

Before we start ...
Prepare to Screenshot.

WELCOME!

Thank you for joining us in these unique circumstances. We hope you're safe and well.

EVENT INFORMATION

Ksqldb - a lightweight stream processing engine for Apache Kafka

Mark Teehan



How to ask questions and interact during and afterwards

Zoom-chat

Any time during the virtual meetup!

Unmute and ask!

(only before and after talks)

Community slack

channel: #events
Throughout and after the event



Confluent Community Catalyst program nominations are open for the class of 2020-2021. Find out more and nominate your community heroes!

[Confluent.io/Nominate](https://confluent.io/Nominate)

THE MEETUP HUB

ALL UPCOMING MEETUPS

NEW MEETUP EMAIL ALERTS

VIDEOS OF PAST MEETUPS

SLIDES FROM THE TALKS

[CNFL.IO/MEETUP-HUB](https://cnfl.io/meetup-hub)



MORE LINKS

developer.confluent.io

All the Apache Kafka® learning resources you'll need

cnfl.io/MU-Try-Cloud

Get \$50 of free usage every month for your first 3 months!



You are being recorded and this footage may be added to public channels

Confluent are giving new users \$50 of free usage per month for their first 3 months

Here's advice on how to use this promotion to try Confluent Cloud for free!

Sign up for a Confluent Cloud account

Please bear in mind that you will be required to enter credit card information but will not be charged unless you go over the \$50 usage in any of the first 3 months or if you don't cancel your subscription before the end of your promotion.

You won't be charged if you don't go over the limit!

Get the benefits of Confluent Cloud, but keep an eye on your account making sure that you have enough remaining free credits available for the rest of your subscription month!!

Cancel before the 3 months end If you don't want to continue past the promotion

If you fail to cancel within your first three months you will start being charged full price. To cancel, immediately stop all streaming and storing data in Confluent Cloud and email cloud-support@confluent.io

bit.ly/TryConfluentCloud

Available on



Google Cloud Platform



K A F K A S U M M I T / 2 0 2 0

Event Streaming Everywhere • Aug 24 - 25

Cnfl.io/Summit-2020-MC



THE VIRTUAL CONFERENCE BY



CONFLUENT



Project Metamorphosis

Unveiling the next-gen event streaming platform

**July 1
(yesterday!)**

cnfl.io/pm

cnfl.io/MU-Try-Cloud



Jay Kreps

Co-founder and CEO
Confluent

Lightweight stream processing engine
For Apache Kafka
ksqlDB

**THE
FOUR-CLUSTER
PROBLEM**

**Mark Teehan
Solution Engineer
Confluent
(Singapore)**

**Translation:
이준호 (Brandon/JUNHO LEE)**

Kafka KRU - Oct 25 2019

kafka KRU





Mark Teehan

Sales Engineer at Confluent

Singapore · [500+ connections](#) · [Contact info](#)

teehan@confluent.io

Confluent in Korea!

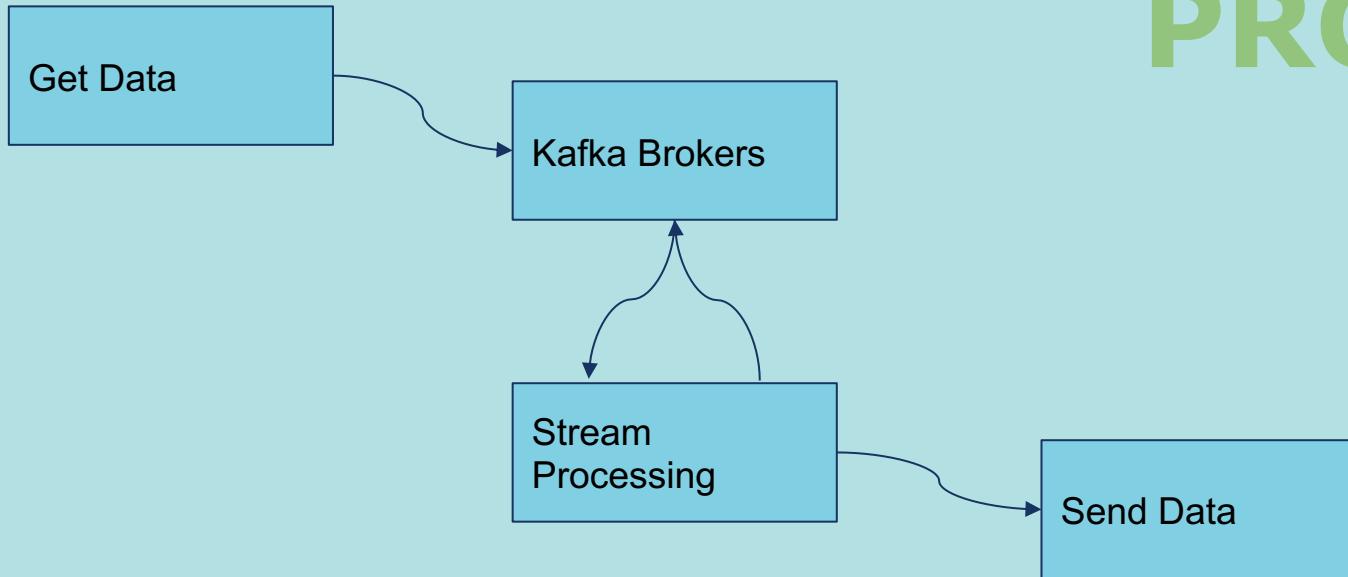
최영주 한국영업대표

dchey@confluent.io

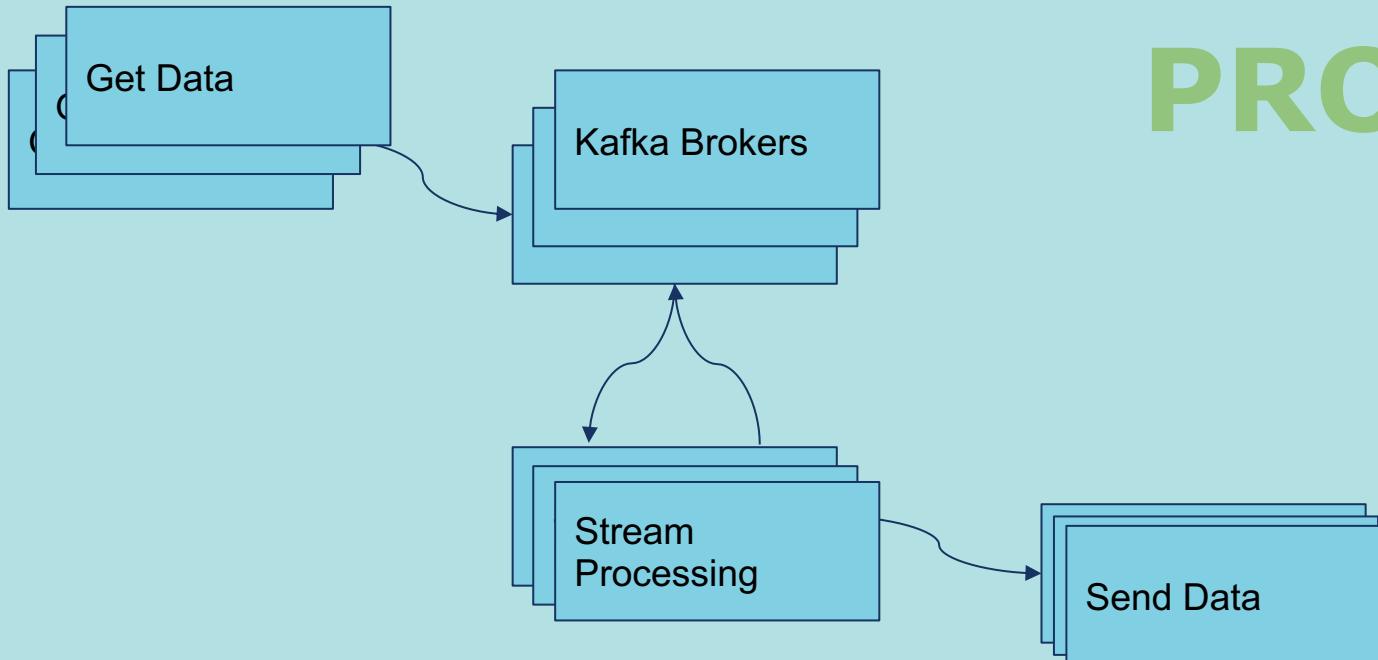


Donna (최영주) Chey · 1st [in](#)
Country Manager at Confluent
Seoul, South Korea · [500+ connections](#)

THE
**FOUR-CLUSTER
PROBLEM**



THE
**FOUR-CLUSTER
PROBLEM**





What is ksqlDB?

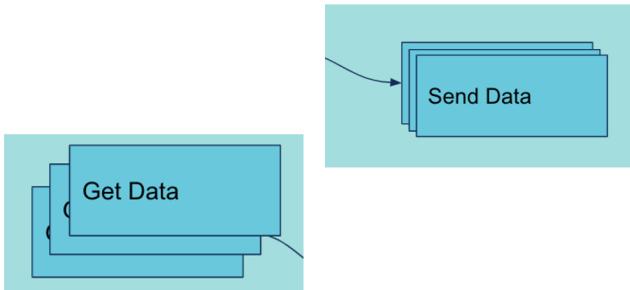
ksqlDB is an **event streaming database** purpose-built to help developers create **stream processing** applications on top of Apache Kafka®.

Is ksqlDB owned by the Apache Software Foundation?

No. This is a community component of Confluent Platform. ksqlDB is owned and maintained by [Confluent Inc.](#) as part of its [Confluent Platform](#) product. However, ksqlDB is licensed under the Confluent Community License.

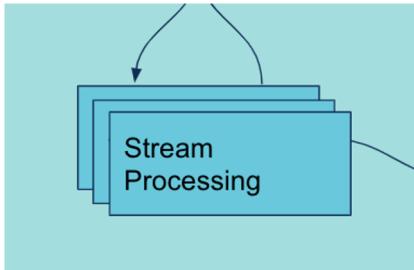


How does ksqlDB solve the four-cluster-problem?



ksqlDB natively integrates with Connect by either

1. communicating with an external Connect cluster or
2. running Connect embedded within the KSQL server process.



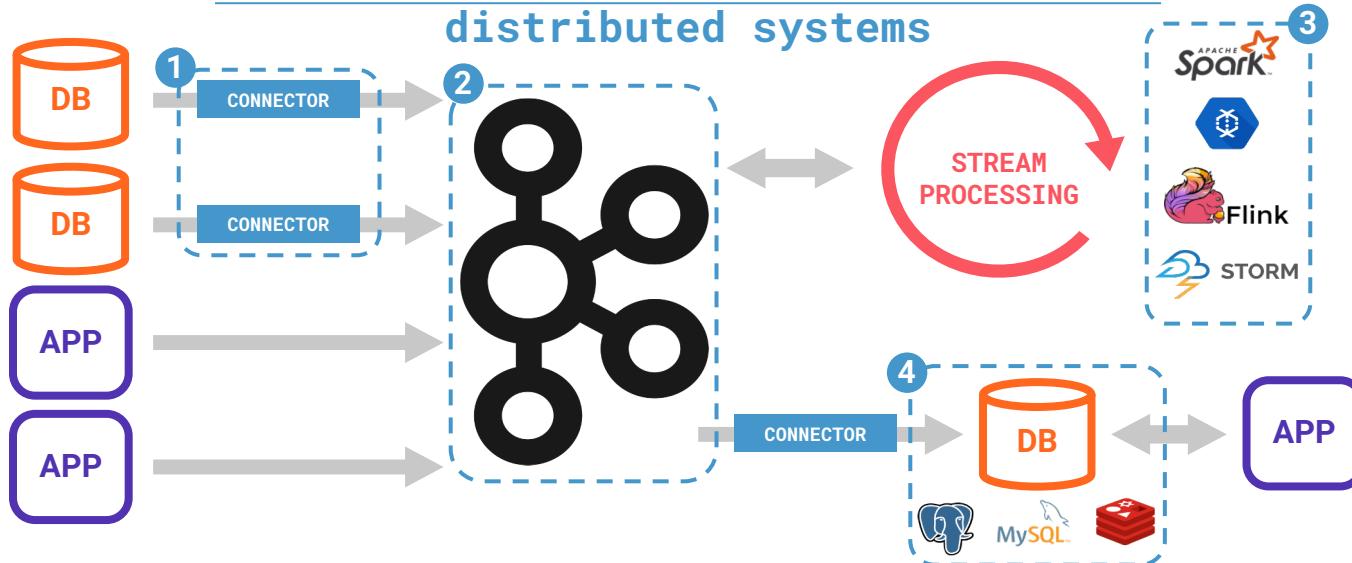
ksqlDB is an event streaming database purpose-built to help developers create stream processing applications on top of Apache Kafka®.

1. Stream Processing with *STREAMs* and *TABLEs*
2. Materialized Views, backed by RocksDB
3. Push and Pull Queries
4. Join Event Streams

Stream processing architectures can be complex

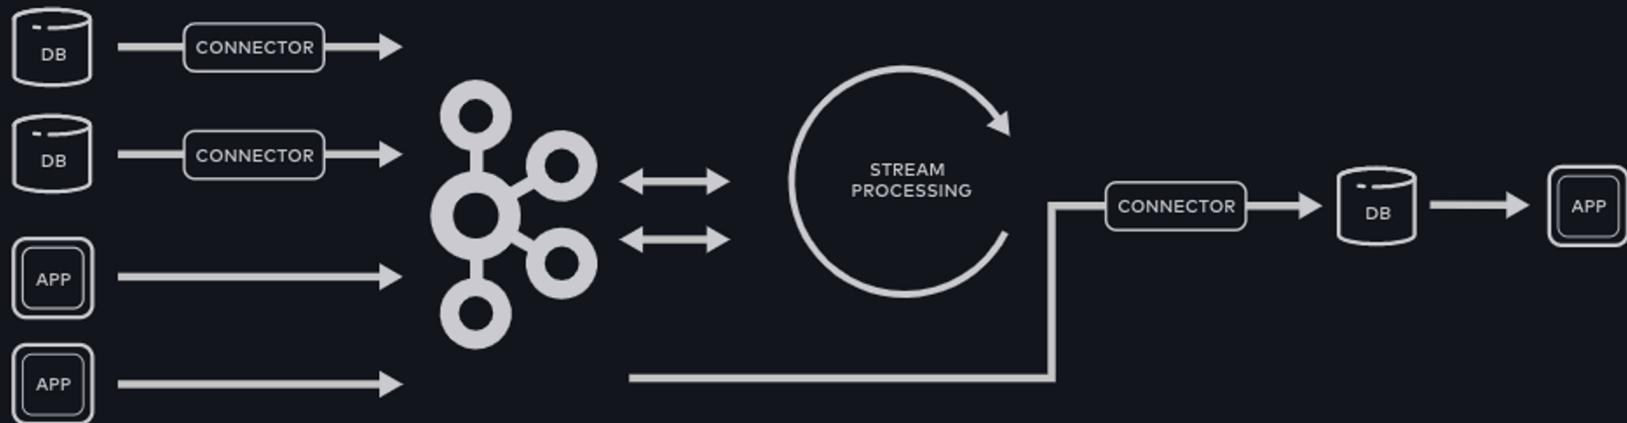


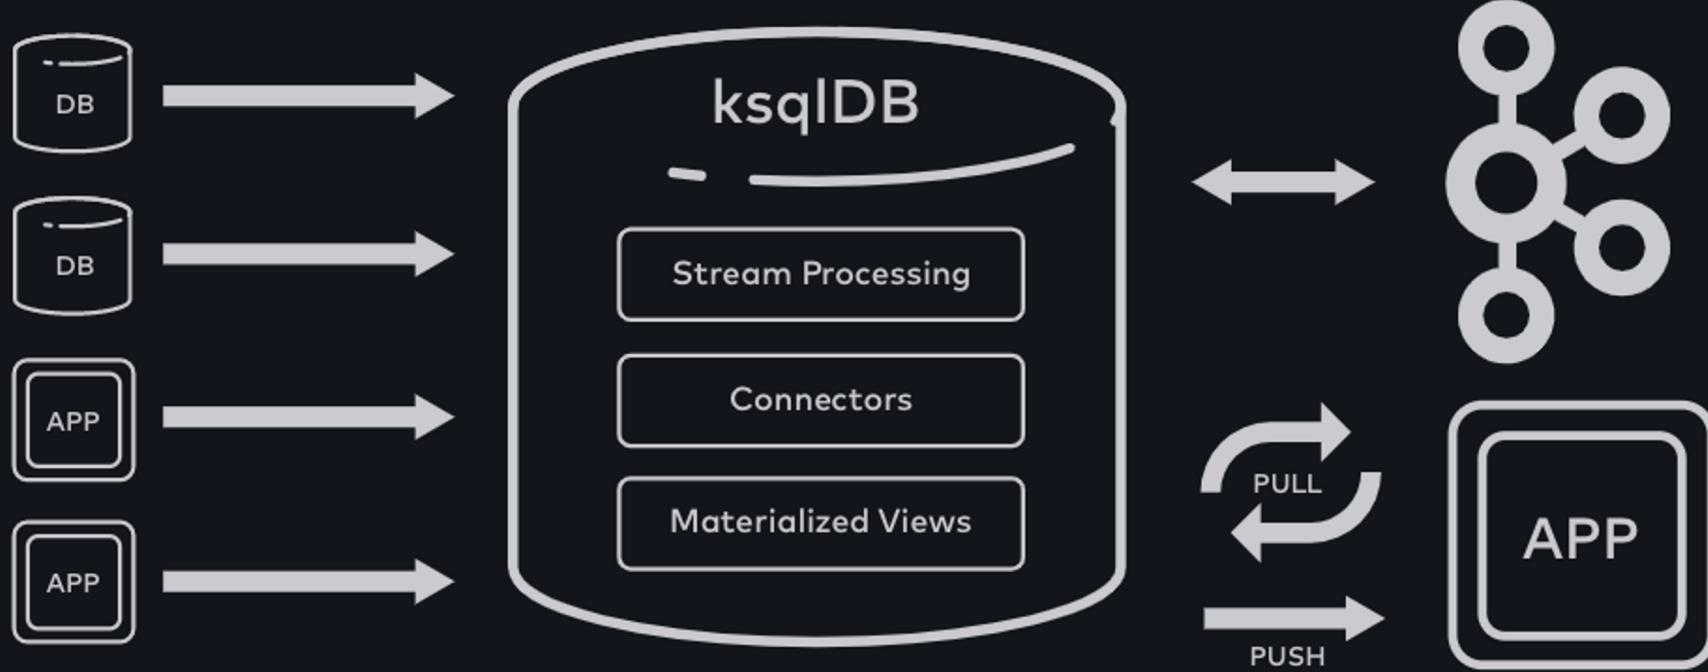
Status quo requires building,
integrating, and managing 3-5
distributed systems



One mental model for the entire stack.

Today, nearly all streaming architectures are complex, piecemeal solutions. They comprise multiple subsystems, each with its own mental model.

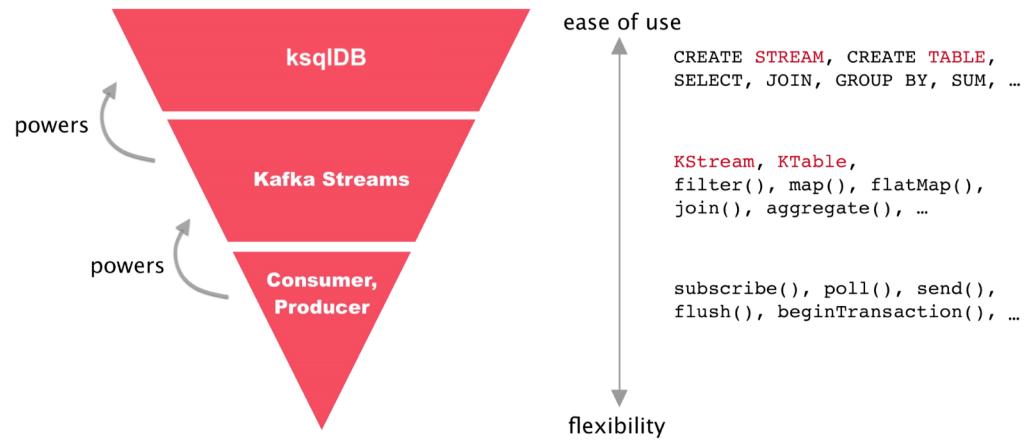




Relationship to Kafka Streams



- ksqlDB is the streaming database for Apache Kafka®.
- With ksqlDB, you can write event streaming applications by using a lightweight SQL syntax.
- Kafka Streams is the Kafka library for writing streaming applications and microservices in Java and Scala.
- ksqlDB is built on Kafka Streams, a robust stream processing framework that is part of Kafka.

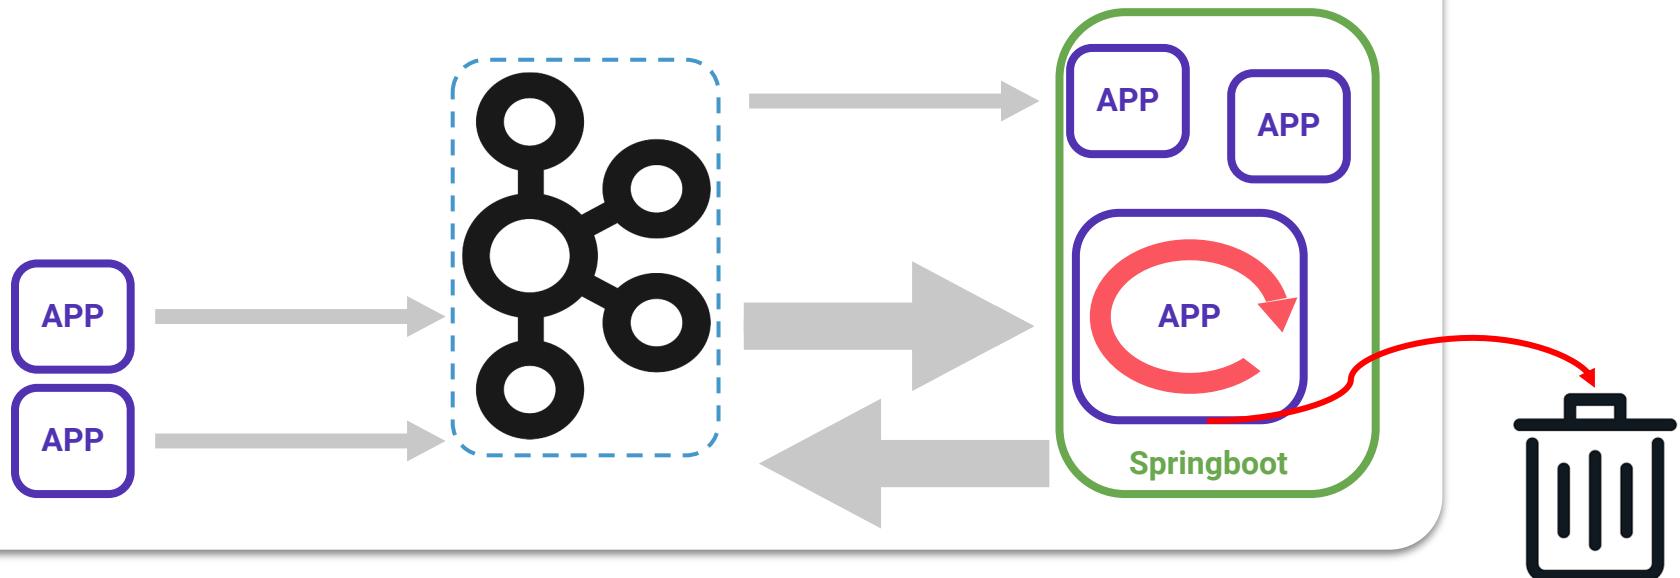


“We don’t do stream processing...”

Are all of your apps equal?



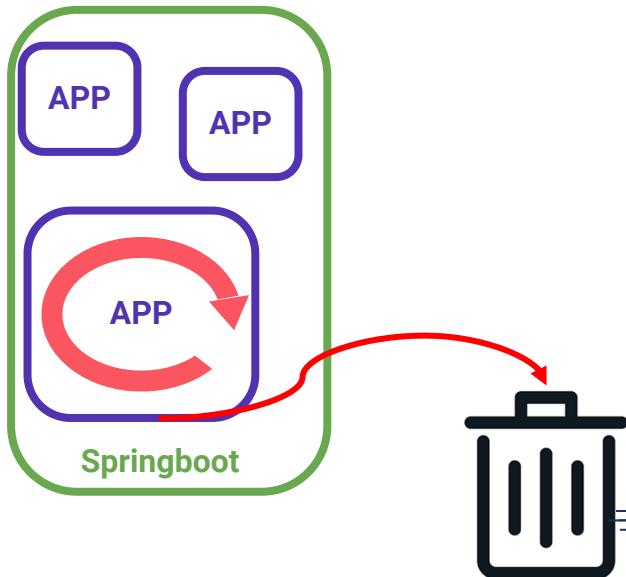
Process messages in the stream, or
in the app



Only consume what you need



It is common to ignore parts of messages



Message: {col1, col2, col3, col4, col5, col6, col7, col8, col9}

Message: {col1, col2, col3, col4, col5, col6, col7, col8, col9}

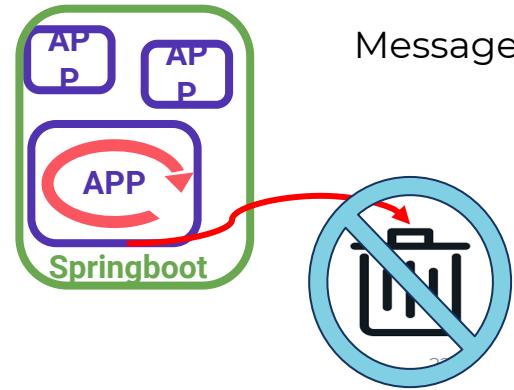
Filter to what the consumer needs:

Message: {col1, col2, col3, col8}

Filtering Columns is stream processing

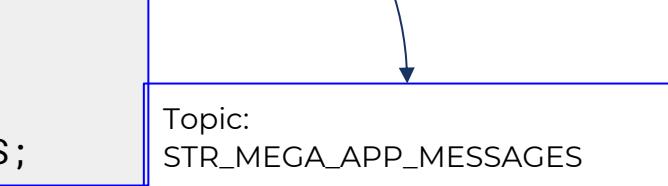


Filter to what the consumer needs



Message: {col1, col2, col3,col8}

```
CREATE STREAM  
STR_MEGA_APP_MESSAGES AS  
SELECT col1  
, col2  
, col3  
, col8  
FROM APP_MESSAGES;
```



Examples of large messages:
ISO Payment Messages (1200 attributes)
Insurance Policies
Syslog entries

Another type of filtering



Filter using a WHERE clause

```
CREATE STREAM  
APP_MESSAGES_KR  
AS SELECT *  
FROM APP_MESSAGES  
WHERE COUNTRY='kr',  
AND...
```

Topic: APP_MESSAGES_KR

KR
APP

8MB/sec

Topic:
APP_MESSAGES

10MB/sec

Topic:
APP_MESSAGES_TW

2MB/sec

TW
APP

```
CREATE STREAM  
APP_MESSAGES_TW  
AS SELECT *  
FROM APP_MESSAGES  
WHERE COUNTRY='tw'
```

The same behaviour as a database **WHERE** clause

Database: views on views

ksqlDB: streams on streams - layers of filters

Useful for “Need to Know” filtering, and replication to cloud

Stream processing:

Filtering (reduce message size)

Filtering (WHERE clause)

JOINS

Aggregation (GROUP BY)

ksqldb Collections: Streams and Tables



STREAM - a regular Kafka Topic - append only

TABLE - a compacted Kafka Topic - uses tombstones

Just normal topics: offsets, partitions, key+value, schemas, consumer groups etc

Use *INSERT*
INTO..VALUES to
Insert/Produce into
a **STREAM** or **TABLE**

You can declare a **STREAM**
or **TABLE** over an
existing topic, or you
can create a new topic
for the collection at
declaration time

Joins



Just like a database join - merge streams from multiple topics

The topics being joined must have the same configuration for partitioning

STREAM+STREAM=STREAM

TABLE+TABLE =TABLE

STREAM+TABLE =STREAM

For STREAM+STREAM joins you must specify a WITHIN clause for matching records within a specified time interval.

A *WITHIN* clause also manages out-of-order and late-arriving data
“...*WITHIN 60 MINUTES...*”



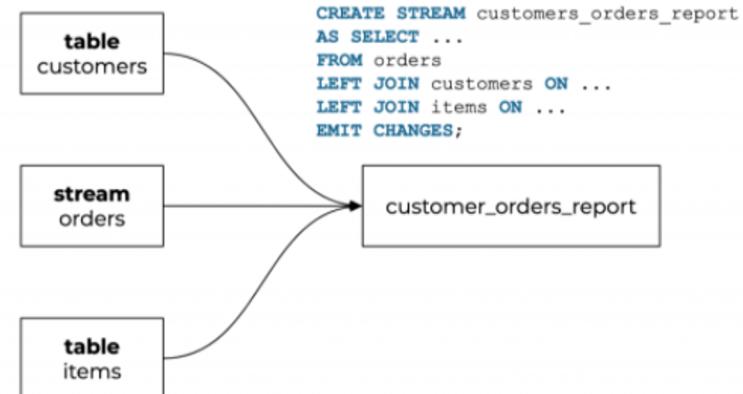
Support for *Multi-Join*

KsqlDB 0.9.0 (19-May-2020): You can do multiple joins into a single statement.

There is no limit to the number of joins in a single statement.

This works with all join types: inner, left, and outer joins.

Normal restrictions apply: you can't repartition tables





ksqldb Change Log

<https://github.com/confluentinc/ksql/blob/master/CHANGELOG.md>

Change Log

0.10.0 (2020-06-25)

Features

- Any key name ([#5093](#)) ([1f0ca3e](#))
- Explicit keys ([#5533](#)) ([d0db0cf](#))
- add extra log messages for pull queries ([#4909](#)) ([d622ecc](#))
- Adds the ability have internal endpoints listen on ksql.internal.listener ([#5212](#)) ([46acb73](#))
- create MetricsContext for ksql metrics reporters ([#5528](#)) ([50561a5](#))
- drop WITH(KEY) syntax ([#5363](#)) ([bb43d23](#))
- expose JMX metric that classifies an error state ([#5374](#)) ([52271bf](#))
- Expose Vert.x metrics ([#5340](#)) ([e82f762](#))
- introduce RegexClassifier for classifying errors via cfg ([#5412](#)) ([b25dd98](#))
- Pull Queries: QPS check utilizes internal API flag to determine if forwarded ([#5392](#)) ([08b428f](#))
- reload TLS certificate without restarting server ([#5516](#)) ([a5920b0](#))
- support TIMESTAMP being a key column ([#5542](#)) ([286ce08](#))
- turn on snappy compression for produced data ([#5495](#)) ([27d8ad5](#))
- client: Java client with push + pull query support ([#5200](#)) ([280ef0c](#))



ksqldb queries

PULL Queries - look up information at a point in time (one result)

PUSH Queries - subscribe to a result as it changes in real-time (a continuous stream)
“push to the user”

PULL: Request/Response

PUSH: Asynchronous flows

PULL :

```
SELECT LAT, LONG  
      FROM LOCATION  
     WHERE ROWKEY='SGX1191'
```

PUSH :

```
SELECT *  
      FROM PAGEVIEWS  
     EMIT CHANGES  
    LIMIT 100;
```

Running ksqlDB

Operate and Deploy

Operations Index

Deploy

Install ksqlDB

Configure ksqlDB CLI

Configure ksqlDB with Docker

Install ksqlDB by using Docker

Check the Health of a ksqlDB Server

Server Configuration

Configure ksqlDB Server

Configure ksqlDB for Avro or Protobuf

Configuration Parameter Reference

Configure Security for ksqlDB

Upgrade ksqlDB

Plan Capacity

KSQL and ksqlDB

Changelog

```
ksqldb:  
  image: confluentinc/ksqldb-server:0.9.0  
  hostname: p35-ksqldb  
  container_name: p35-ksqldb  
  depends_on:  
    - kafka1  
    - schema-registry  
    - kafka-connect  
  ports:  
    - 8088:8088  
  environment:  
    KSQL_BOOTSTRAP_SERVERS: p35-kafka1:29092  
    KSQL_LISTENERS: http://p35-ksqldb:8088  
    KSQL_KSQL_LOGGING_PROCESSING_STREAM_AUTO_CREATE: "true"  
    KSQL_KSQL_LOGGING_PROCESSING_TOPIC_AUTO_CREATE: "true"  
    KSQL_KSQL_CONNECT_URL: http://p35-kc:8083  
    KSQL_KSQL_SCHEMA_REGISTRY_URL: http://p35-sr:8081  
    KSQL_KSQL_SERVICE_ID: p35  
    KSQL_KSQL_INTERNAL_TOPIC_REPLICAS: 1
```

Processing Guarantees



ksqldb supports

- **At-least-once semantics**

Records are never lost but may be redelivered. If your stream processing application fails, no data records are lost and fail to be processed, but some data records may be re-read and therefore re-processed. At-least-once semantics is **enabled by default** in your ksqldb configuration, with `processing.guarantee="at_least_once"`.

- **Exactly-once semantics**

Records are processed once. If a producer within a ksqldb application sends a duplicate record, it's written to the broker exactly once.

Exactly-once stream processing is the ability to execute a read-process-write operation exactly one time.

To enable exactly-once semantics, set `processing.guarantee="exactly_once"` in your ksqldb configuration.

Important: Use the `exactly_once` setting with care. To achieve a true exactly-once system, end consumers and producers must also implement exactly-once semantics.

Q&A

Demo

Confluent in Korea!

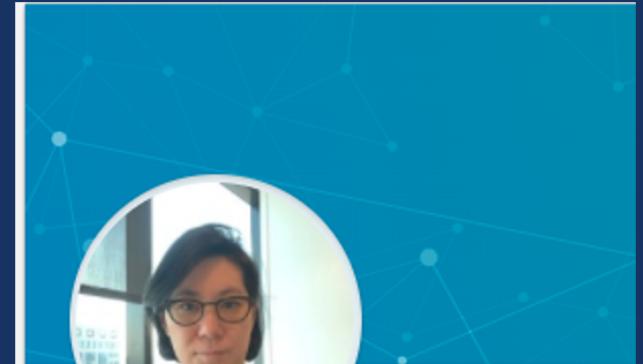
최영주 한국영업대표

**Watch out for a free 2-hour
ksqlDB technical virtual
hands-on Workshop later in
2020.**

**Free of charge, but places are
limited. Please register
interest early.**

Email: dchey@confluent.io

Or contact: 고승범



Donna (최영주) Chey · 1st Country Manager at Confluent
Seoul, South Korea · 500+ connections