

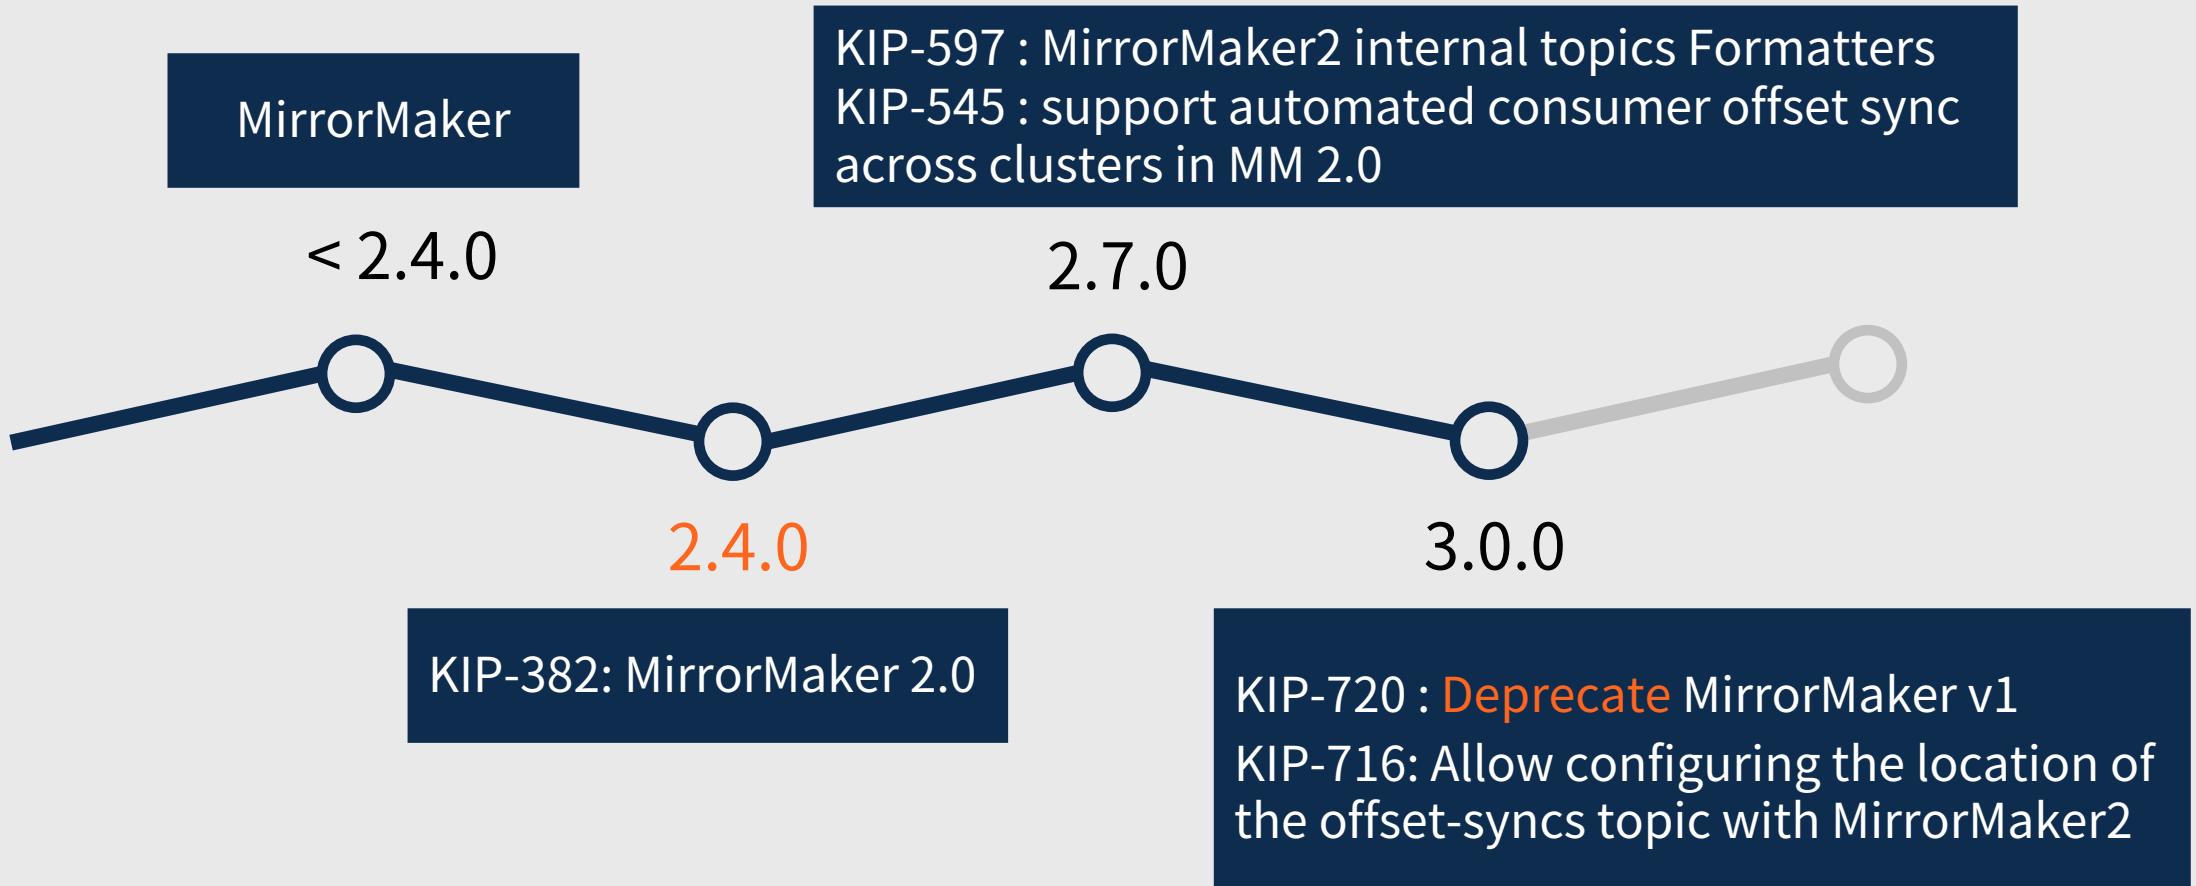
MirrorMaker2 소개

최중호

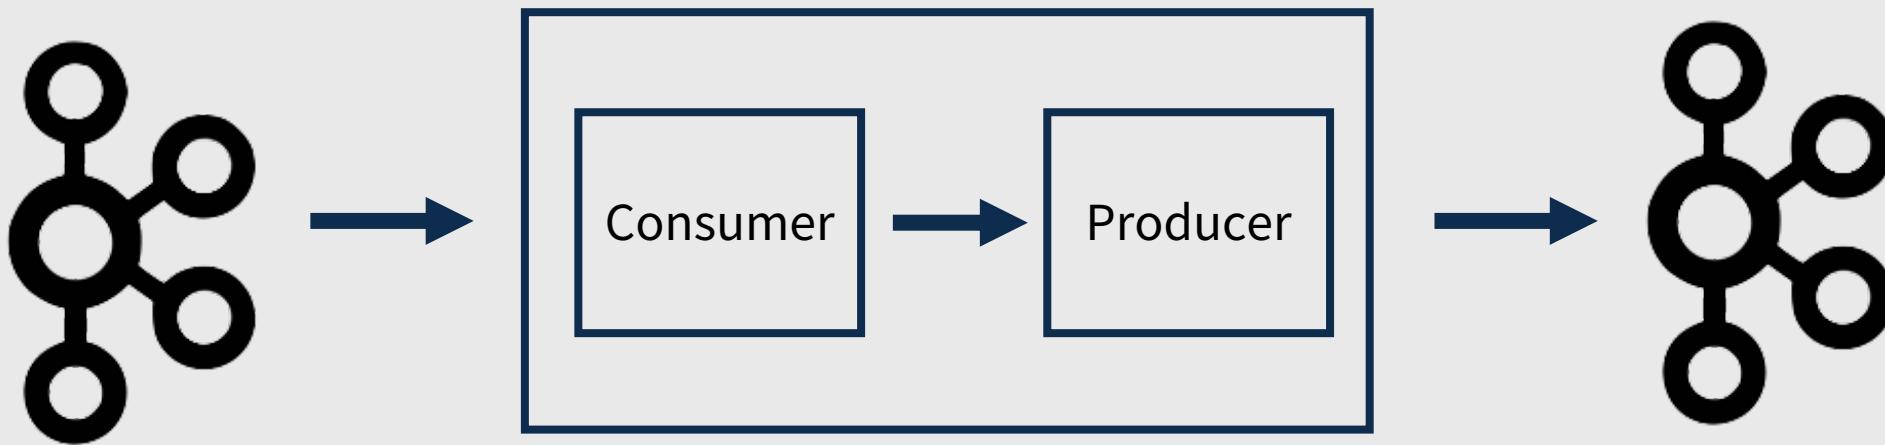
목 차

- MirrorMaker 단점
- MirrorMaker2 아키텍처
- MirrorMaker2 활용
- 모니터링
- 추가 운영 팁

MirrorMaker2 등장



MirrorMaker

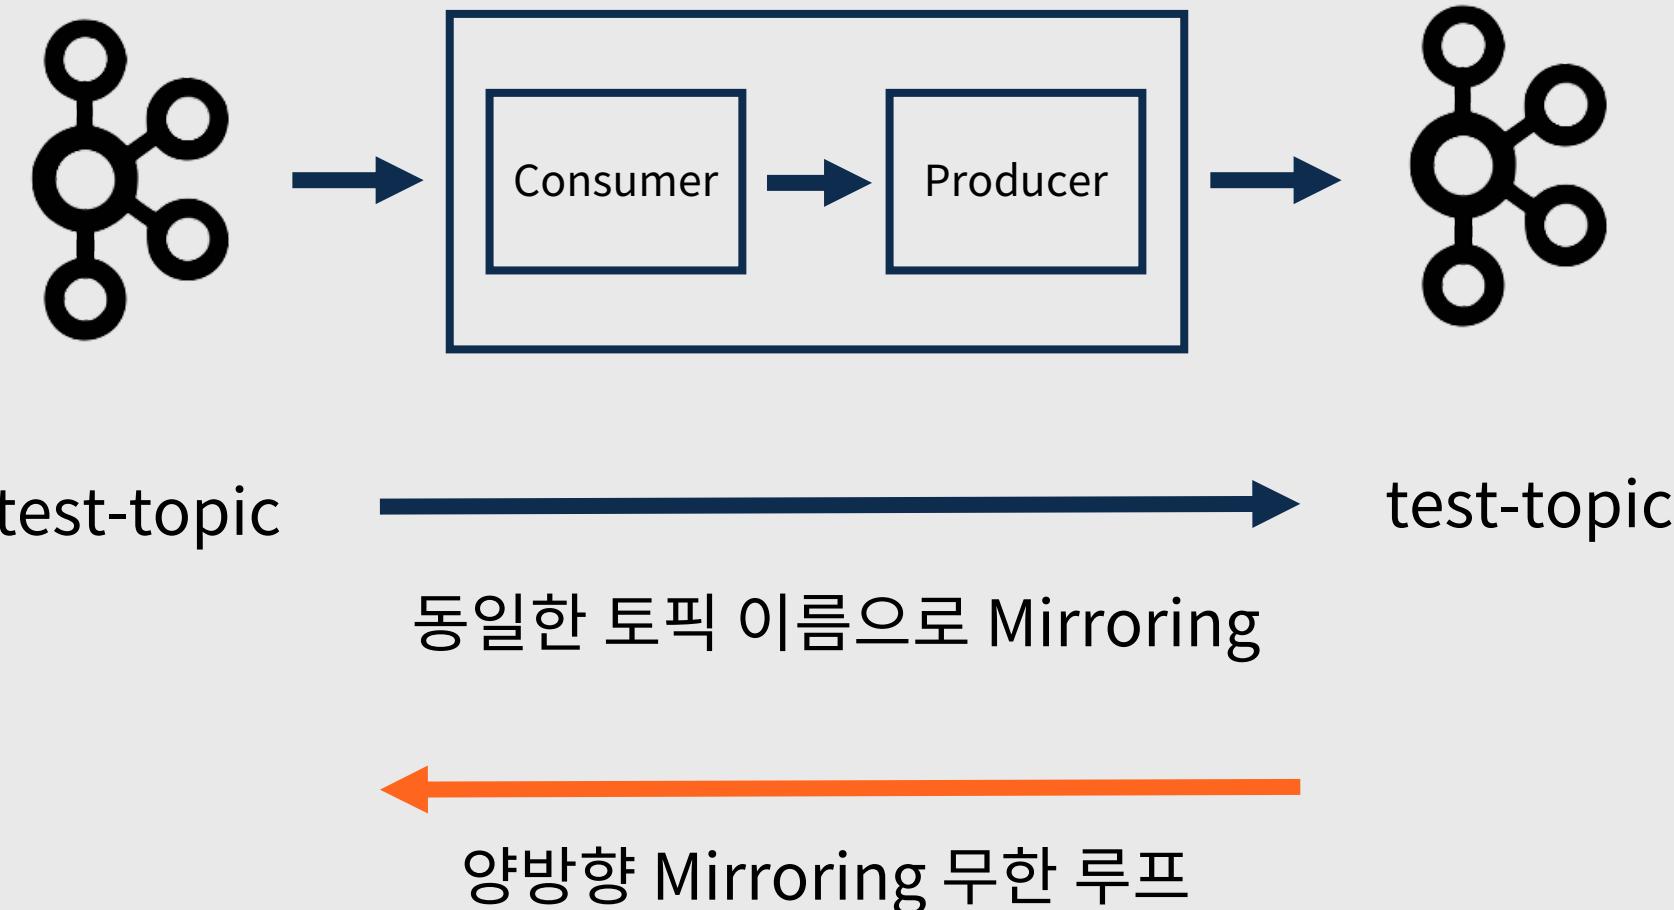


```
./bin/kafka-run-class.sh --daemon --name mirrormaker-v1 kafka.tools.MirrorMaker \
--consumer.config consumer.properties --producer.config producer.properties \
--whitelist= "test-topic"
```

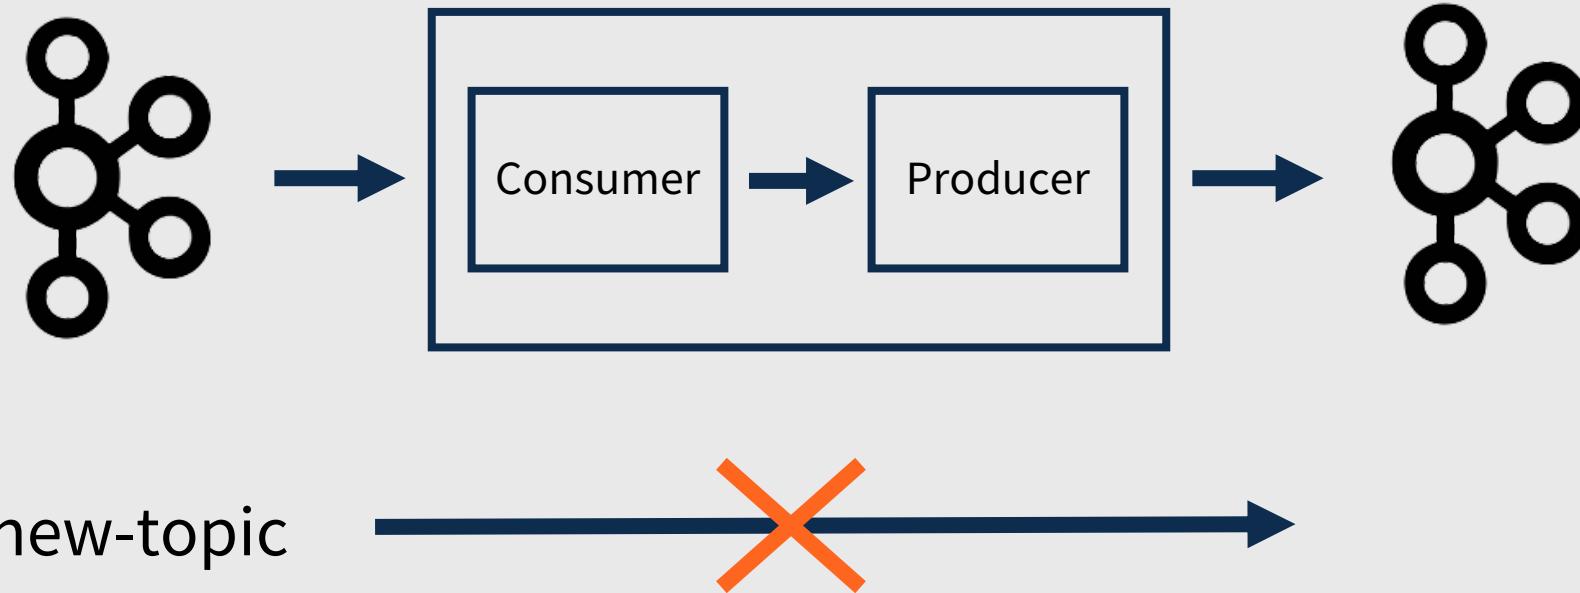
MirrorMaker 단점

- 양방향
- 신규 토픽
- 파티션
- 컨슈머 그룹

MirrorMaker - 양방향

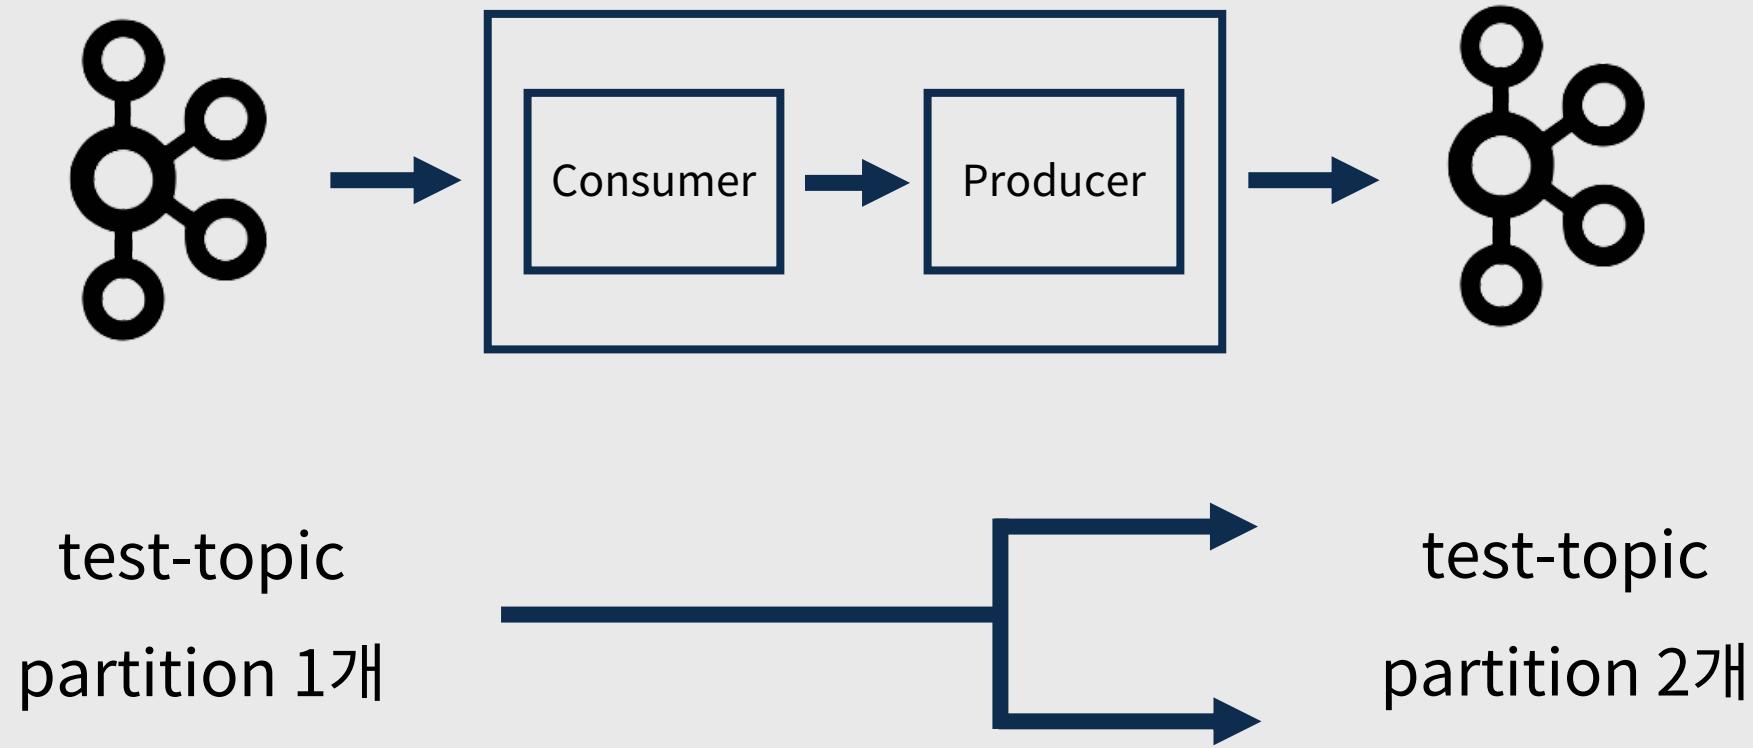


MirrorMaker - 신규 토픽

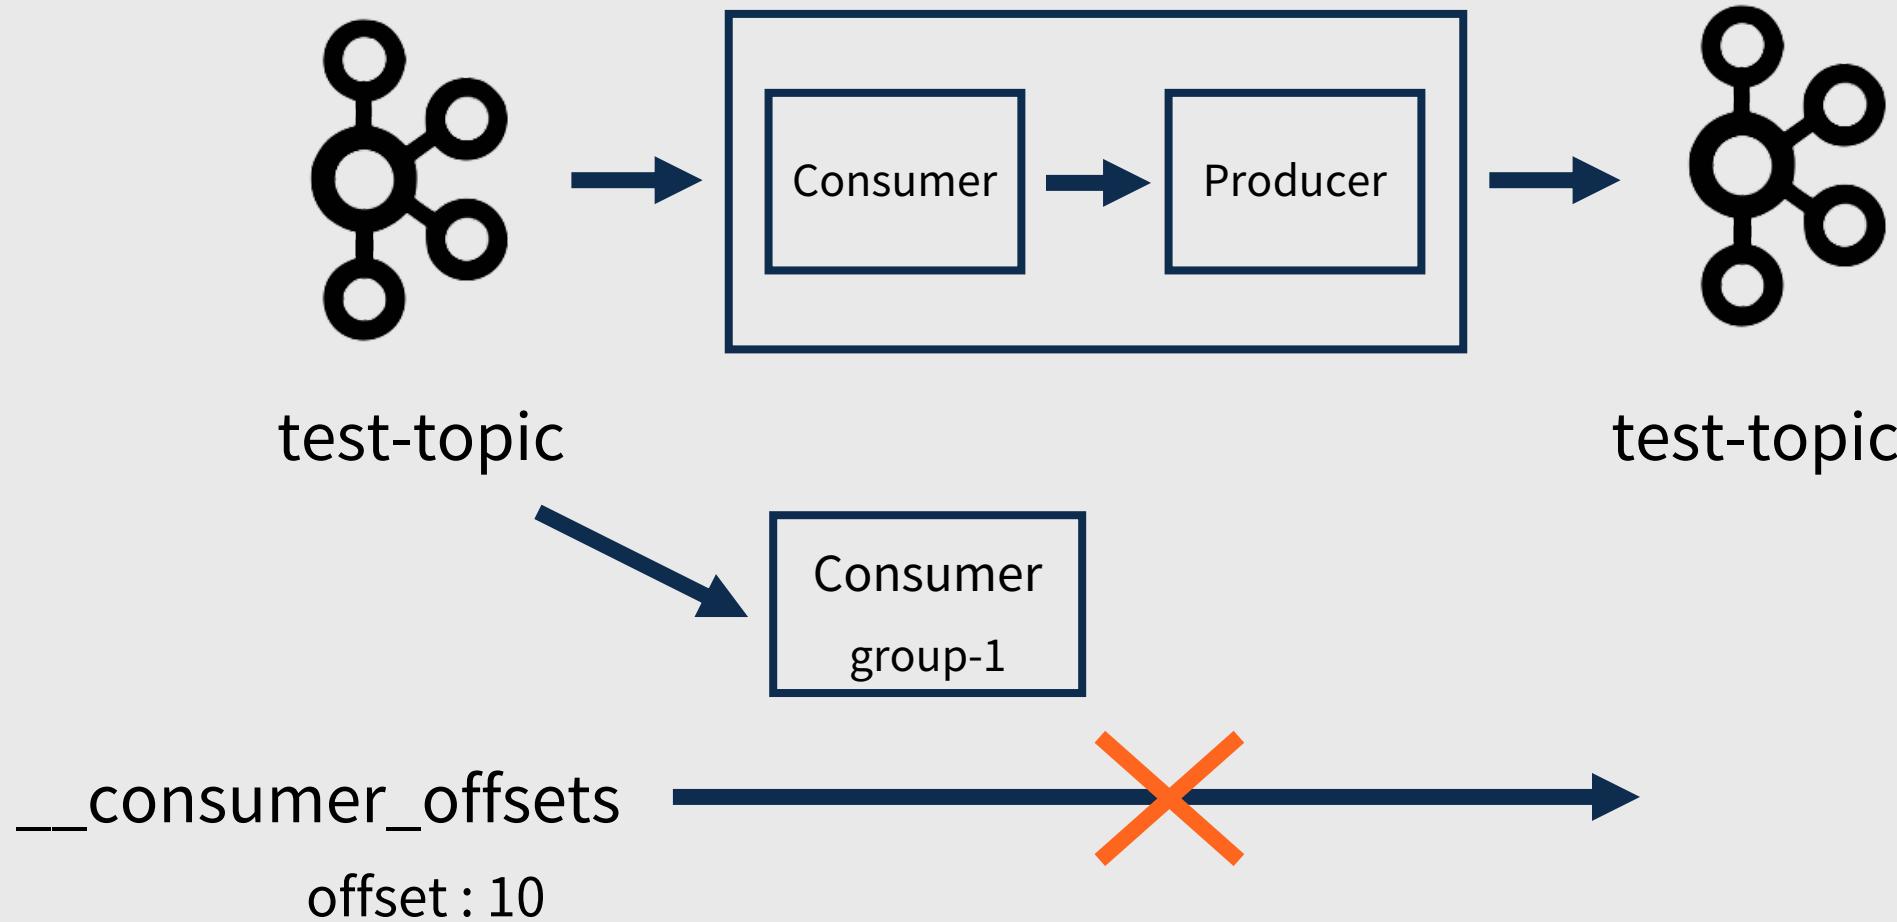


```
./bin/kafka-run-class.sh --daemon --name mirrormaker-v1 kafka.tools.MirrorMaker \  
--consumer.config consumer.properties --producer.config producer.properties \  
--whitelist= "test-topic"
```

MirrorMaker - 패티션

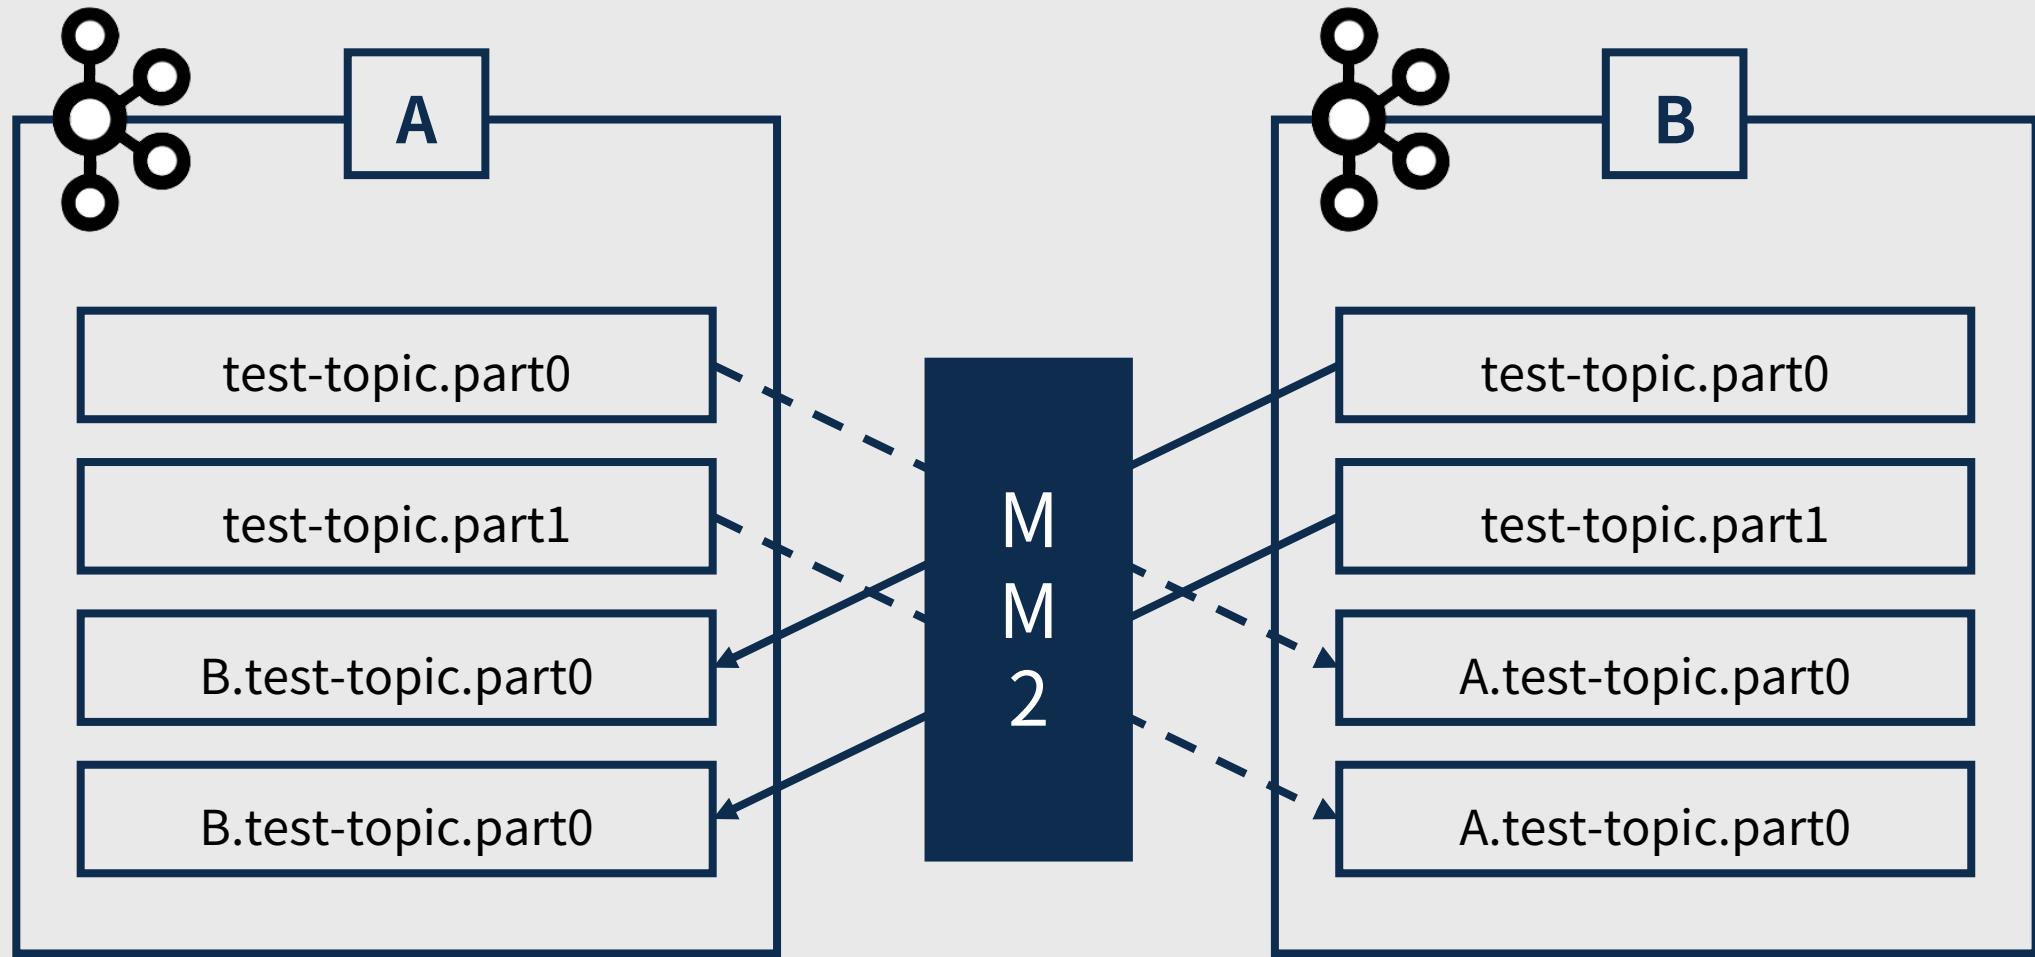


MirrorMaker - 컨슈머 그룹



MirrorMaker2

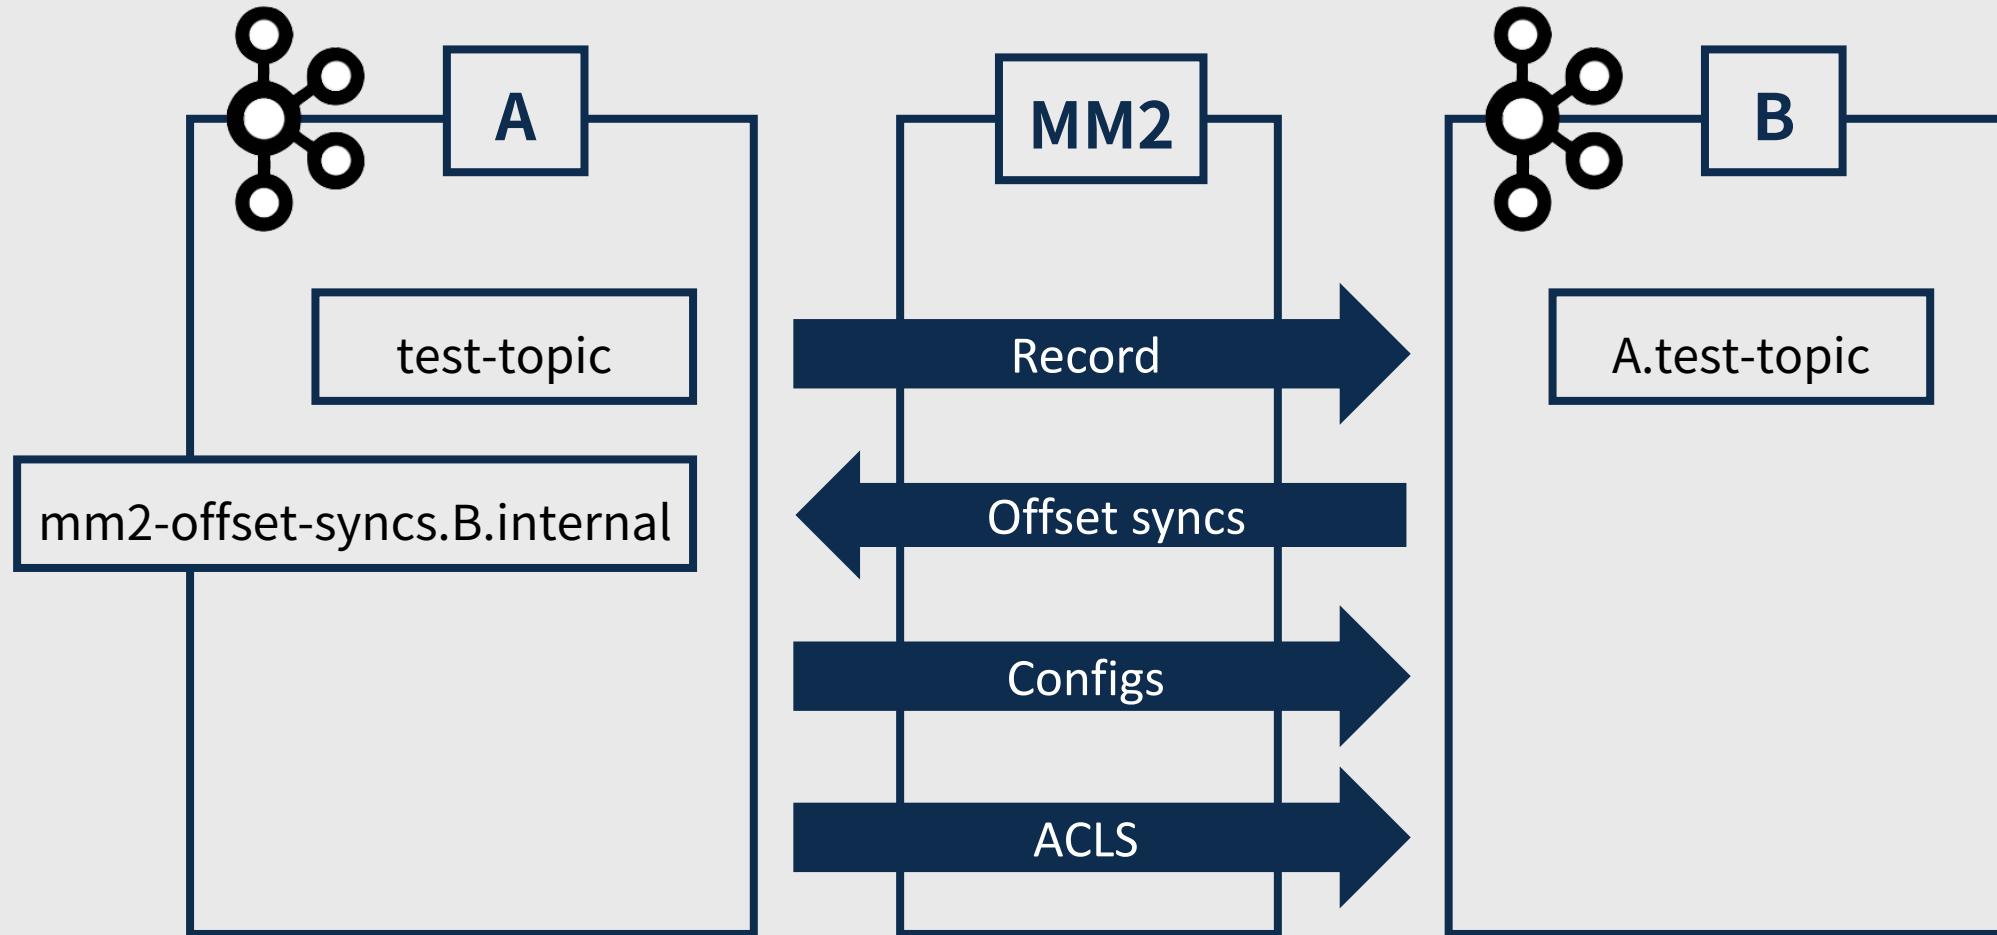
MirrorMaker2 - Active/Active



MirrorMaker2 - 구성

- Kafka Connect 기반
- 3개의 Connector로 구성
 - **MirrorSourceConnector**
 - topic의 데이터를 복제
 - topic의 파티션 정보 및 ACLs 동기화
 - **MirrorCheckpointConnector**
 - consumer group offset 동기화
 - **MirrorHeartbeatConnector**
 - 복제의 모니터링을 위한 Heartbeat 전송

MirrorMaker2 - MirrorSourceConnector



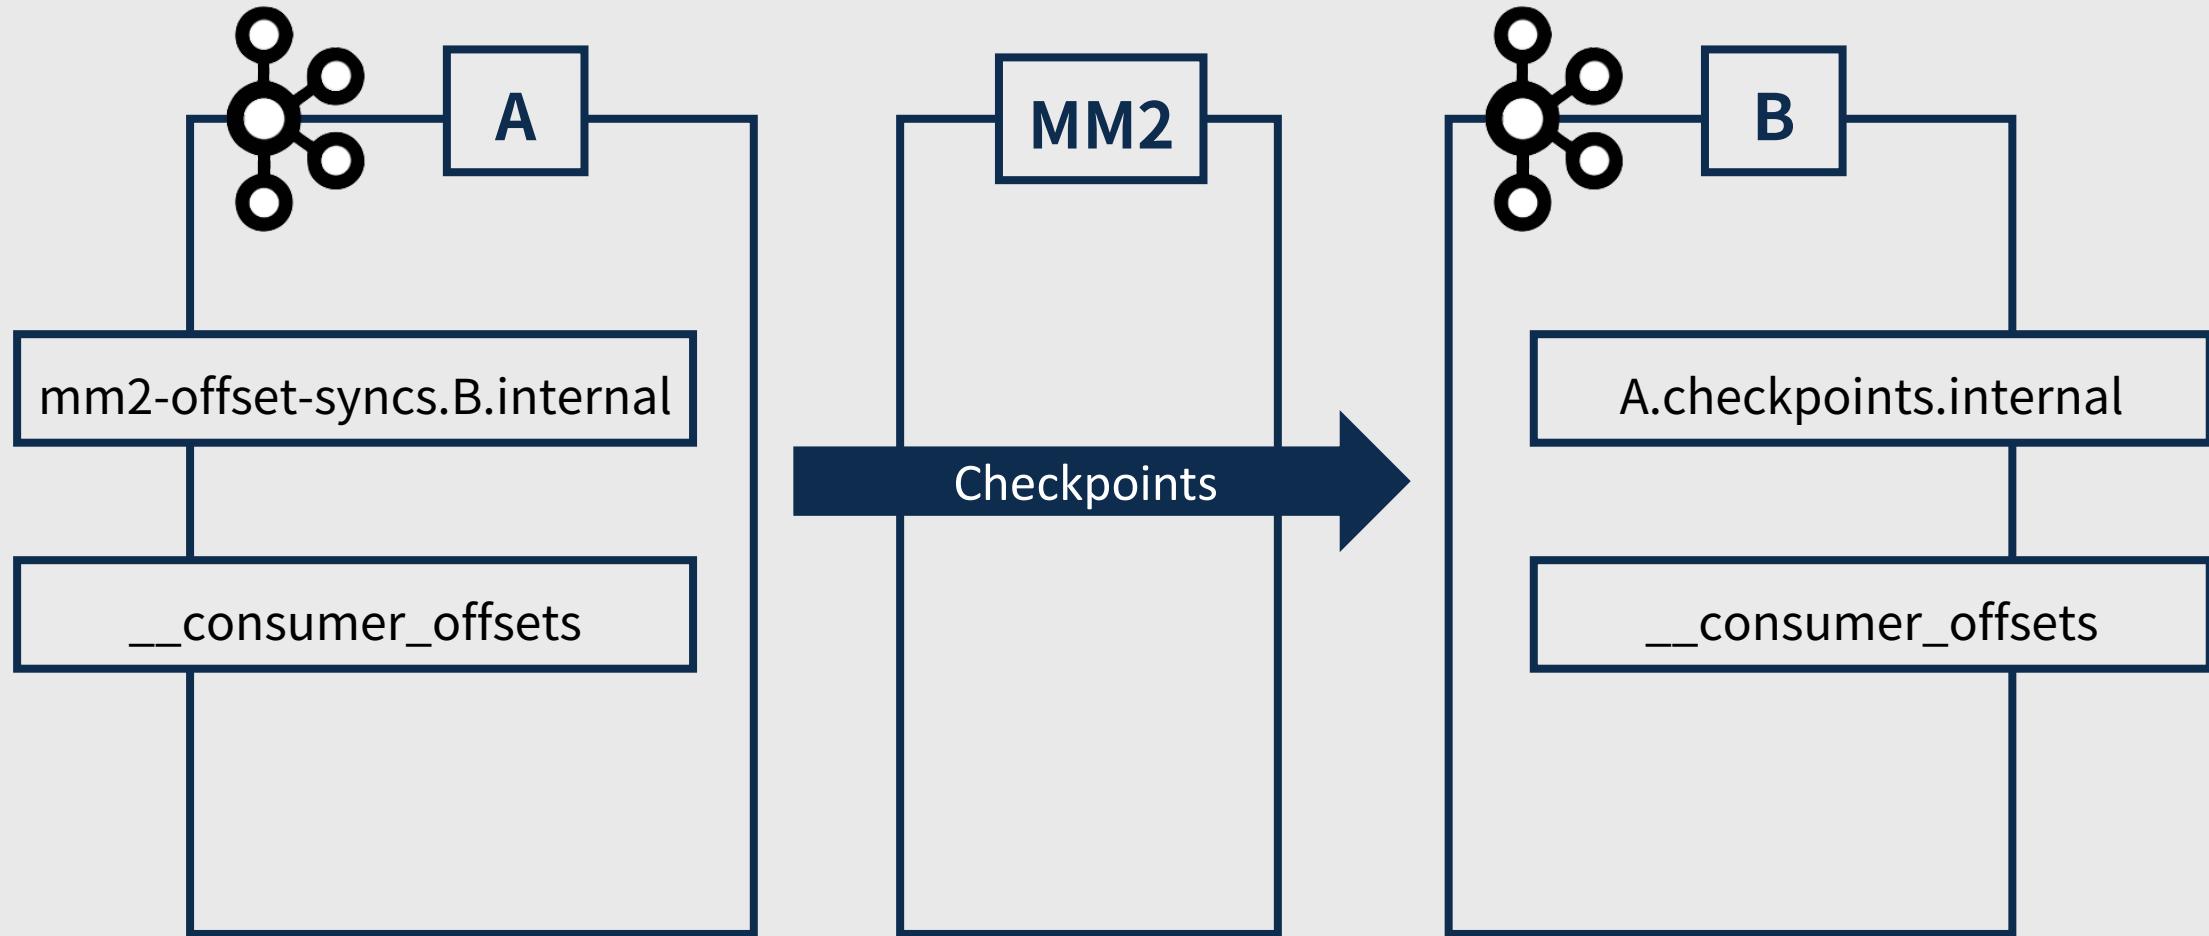
MirrorMaker2 - MirrorSourceConnector

- KIP-597: MirrorMaker2 internal topics Formatters

```
./bin/kafka-console-consumer.sh –bootstrap-server A:9092 \
--topic mm2-offset-syncs.B.internal \
--formatter org.apache.kafka.connect.mirror.formatters.OffsetSyncFormatter
```

```
OffsetSync{topicPartition=test-topic-0,
upstreamOffset=224, downstreamOffset=65}
```

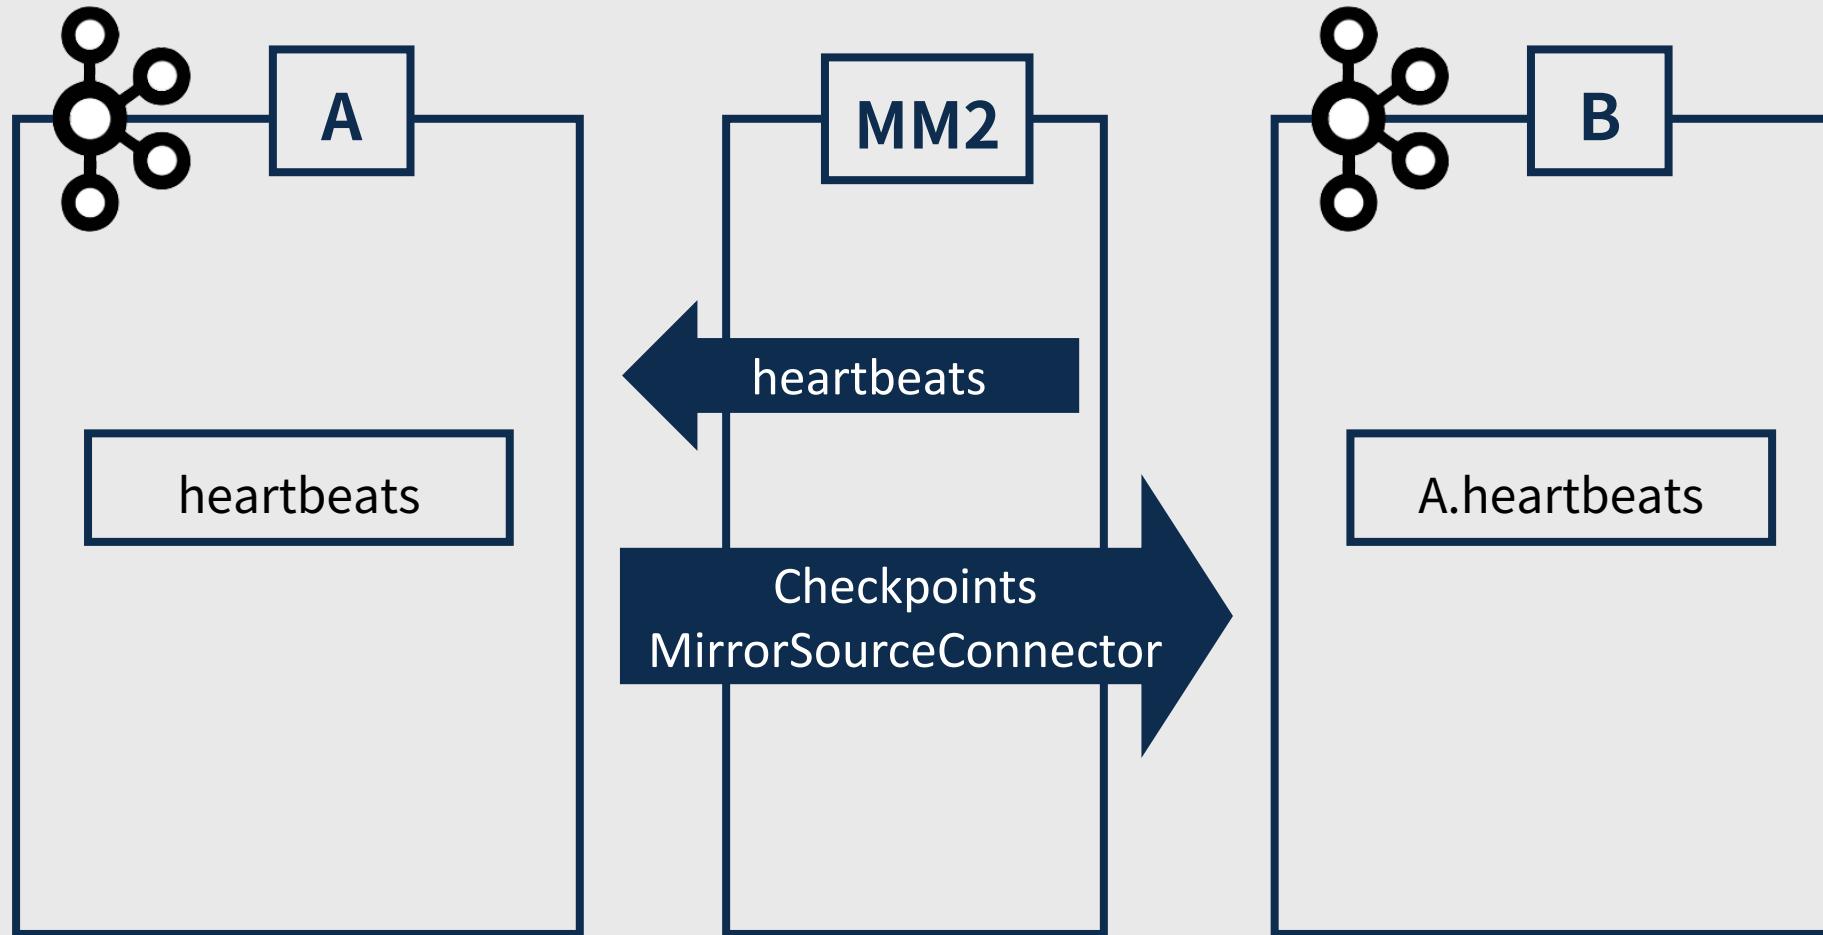
MirrorMaker2 - MirrorCheckpointConnector



MirrorMaker2 - MirrorCheckpointConnector

- KIP-545: support automated consumer offset sync across clusters in MM 2.0
 - A->B.sync.group.offsets.enabled = true
- A 카프카의 컨슈머 오프셋을 B 카프카 오프셋으로 변환한다.
 - 변환할 때 A/B의 오프셋 차이는 mm2-offset-syncs.B.internal로 판별
- 주기적으로 __consumer_offsets에 동기화한다.

MirrorMaker2 - MirrorHeartbeatConnector



MirrorMaker2 - 실행 방법

- Dedicated MirrorMaker cluster
 - MirrorMaker2 전용 클러스터
 - 별도의 properties 구성 후 실행
- Running MirrorMaker in Connect cluster
 - Connect Cluster에 3개의 Connector를 각각 실행

MirrorMaker2 - dedicated cluster (Active/Standby)

mm2.properties

```
clusters = source, target  
source.bootstrap.servers = source:9092  
target.bootstrap.servers = target:9092  
  
source->target.enabled = true  
source->target.topics = .*  
source->target.group = .*  
source->target.sync.group.offsets.enabled = true  
  
target->source.enabled = false
```

`./bin/connect-mirror-maker.sh mm2.properties`

MirrorMaker2 - Connect Cluster (Active/Standy)

MirrorSourceConnector

```
{  
  "name": "mm2-msc",  
  "config": {  
    "name": "mm2-msc",  
    "connector.class": "org.apache.kafka.connect.mirror.MirrorSourceConnector",  
    "clusters": "source, target",  
    "source.cluster.alias": "source",  
    "source.cluster.bootstrap.servers": "source_kafka:9092",  
    "target.cluster.alias": "target",  
    "target.cluster.bootstrap.servers": "target_kafka:9092",  
    "topics": ".*",  
    "tasks.max": "5"  
  }  
}
```

POST http://connect_url/connectors

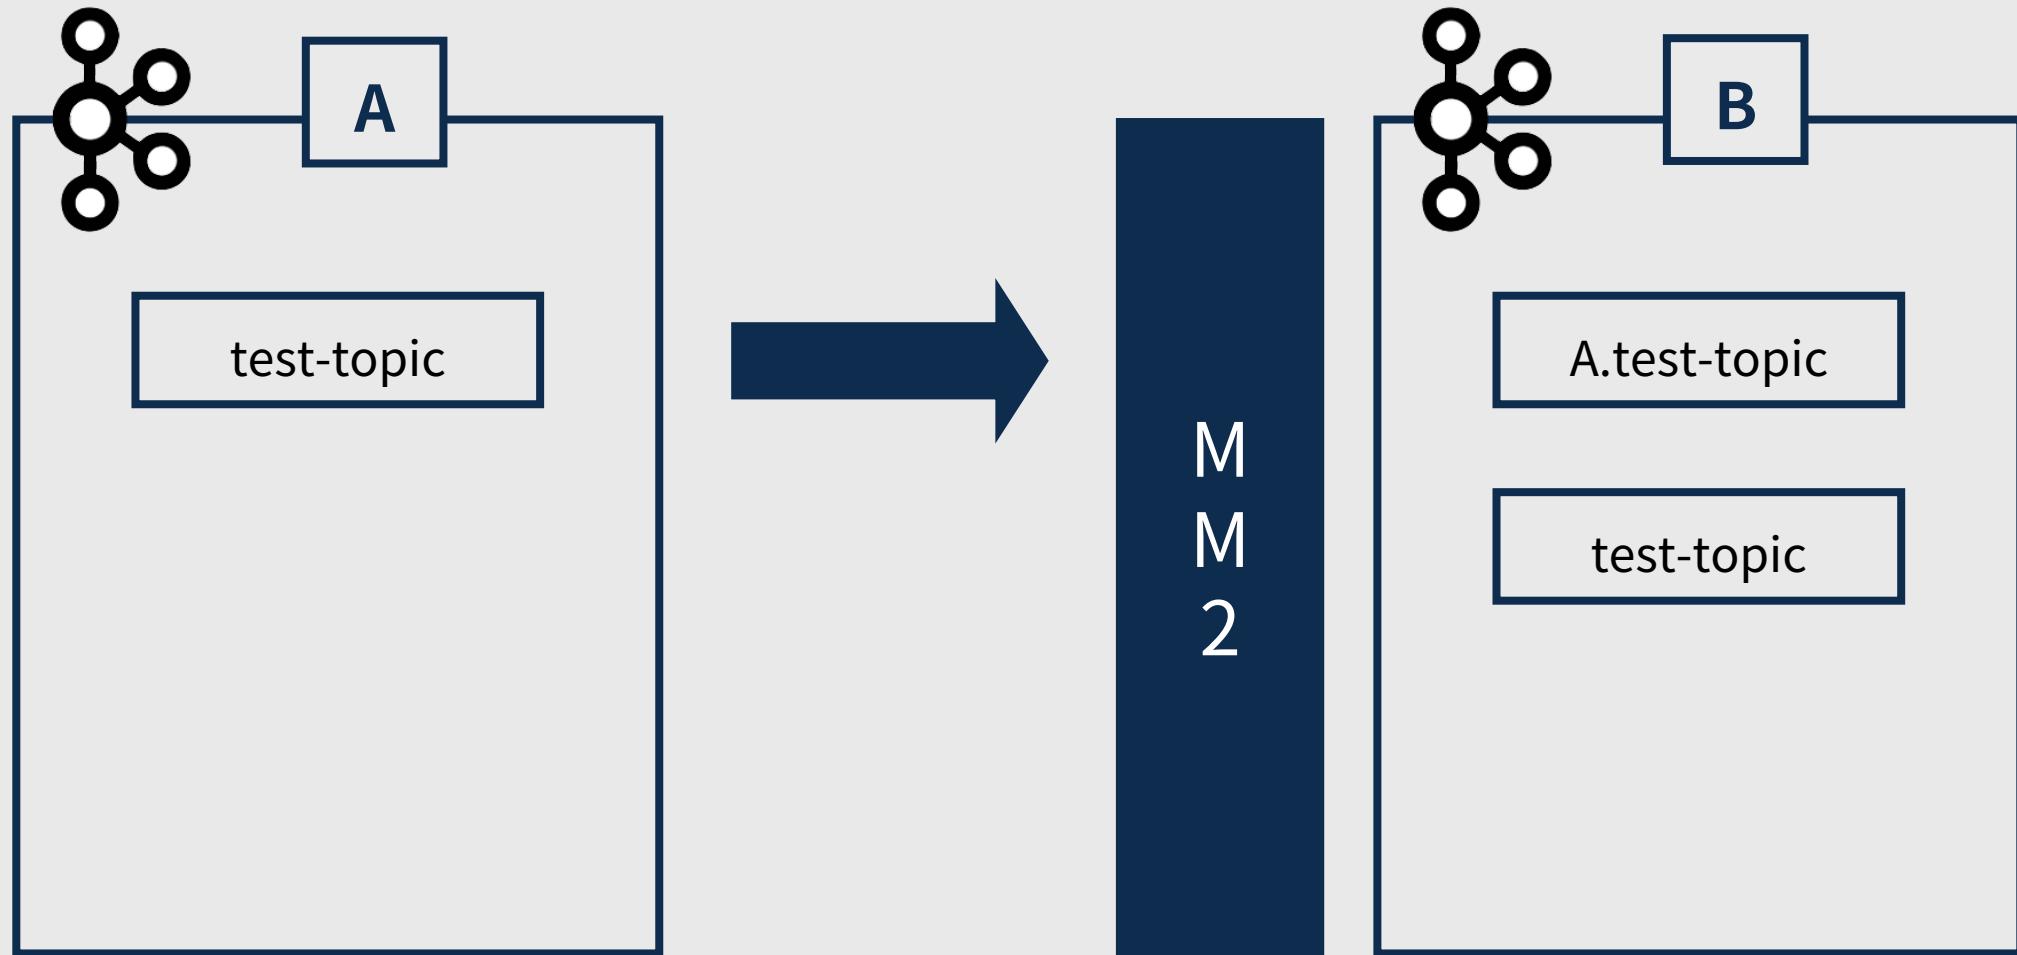
MirrorMaker2 - Connect Cluster (Active/Standy)

MirrorCheckpointConnector

```
{  
  "name": "mm2-cpc",  
  "config":  
  {  
    "name": "mm2-cpc",  
    "connector.class": "org.apache.kafka.connect.mirror.MirrorCheckpointConnector",  
    "clusters": "source, target",  
    "source.cluster.alias": "source",  
    "source.cluster.bootstrap.servers": "source_kafka:9094",  
    "target.cluster.alias": "target",  
    "target.cluster.bootstrap.servers": "target_kafka:9094",  
    "groups": ". *",  
    "tasks.max": "5"  
  }  
}
```

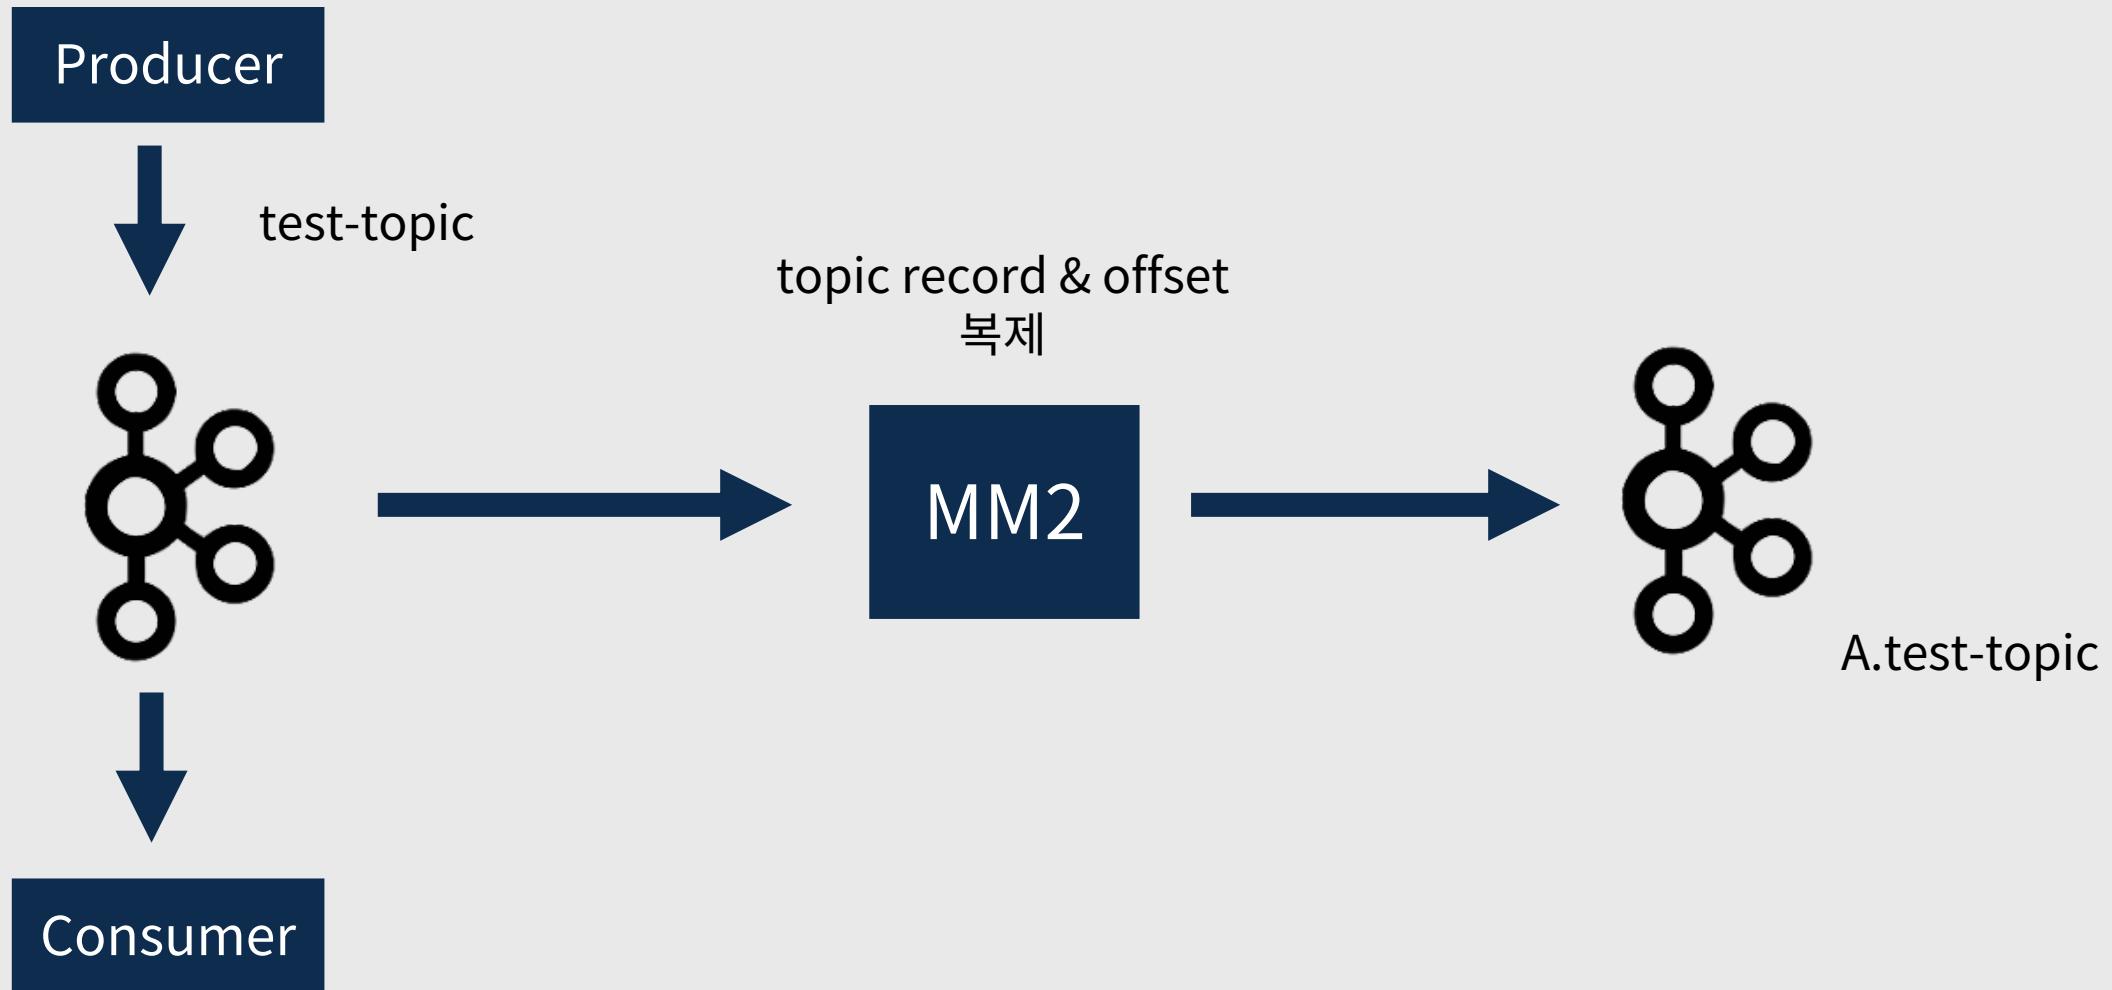
POST http://connect_url/connectors

MirrorMaker2 - Active/ Standby

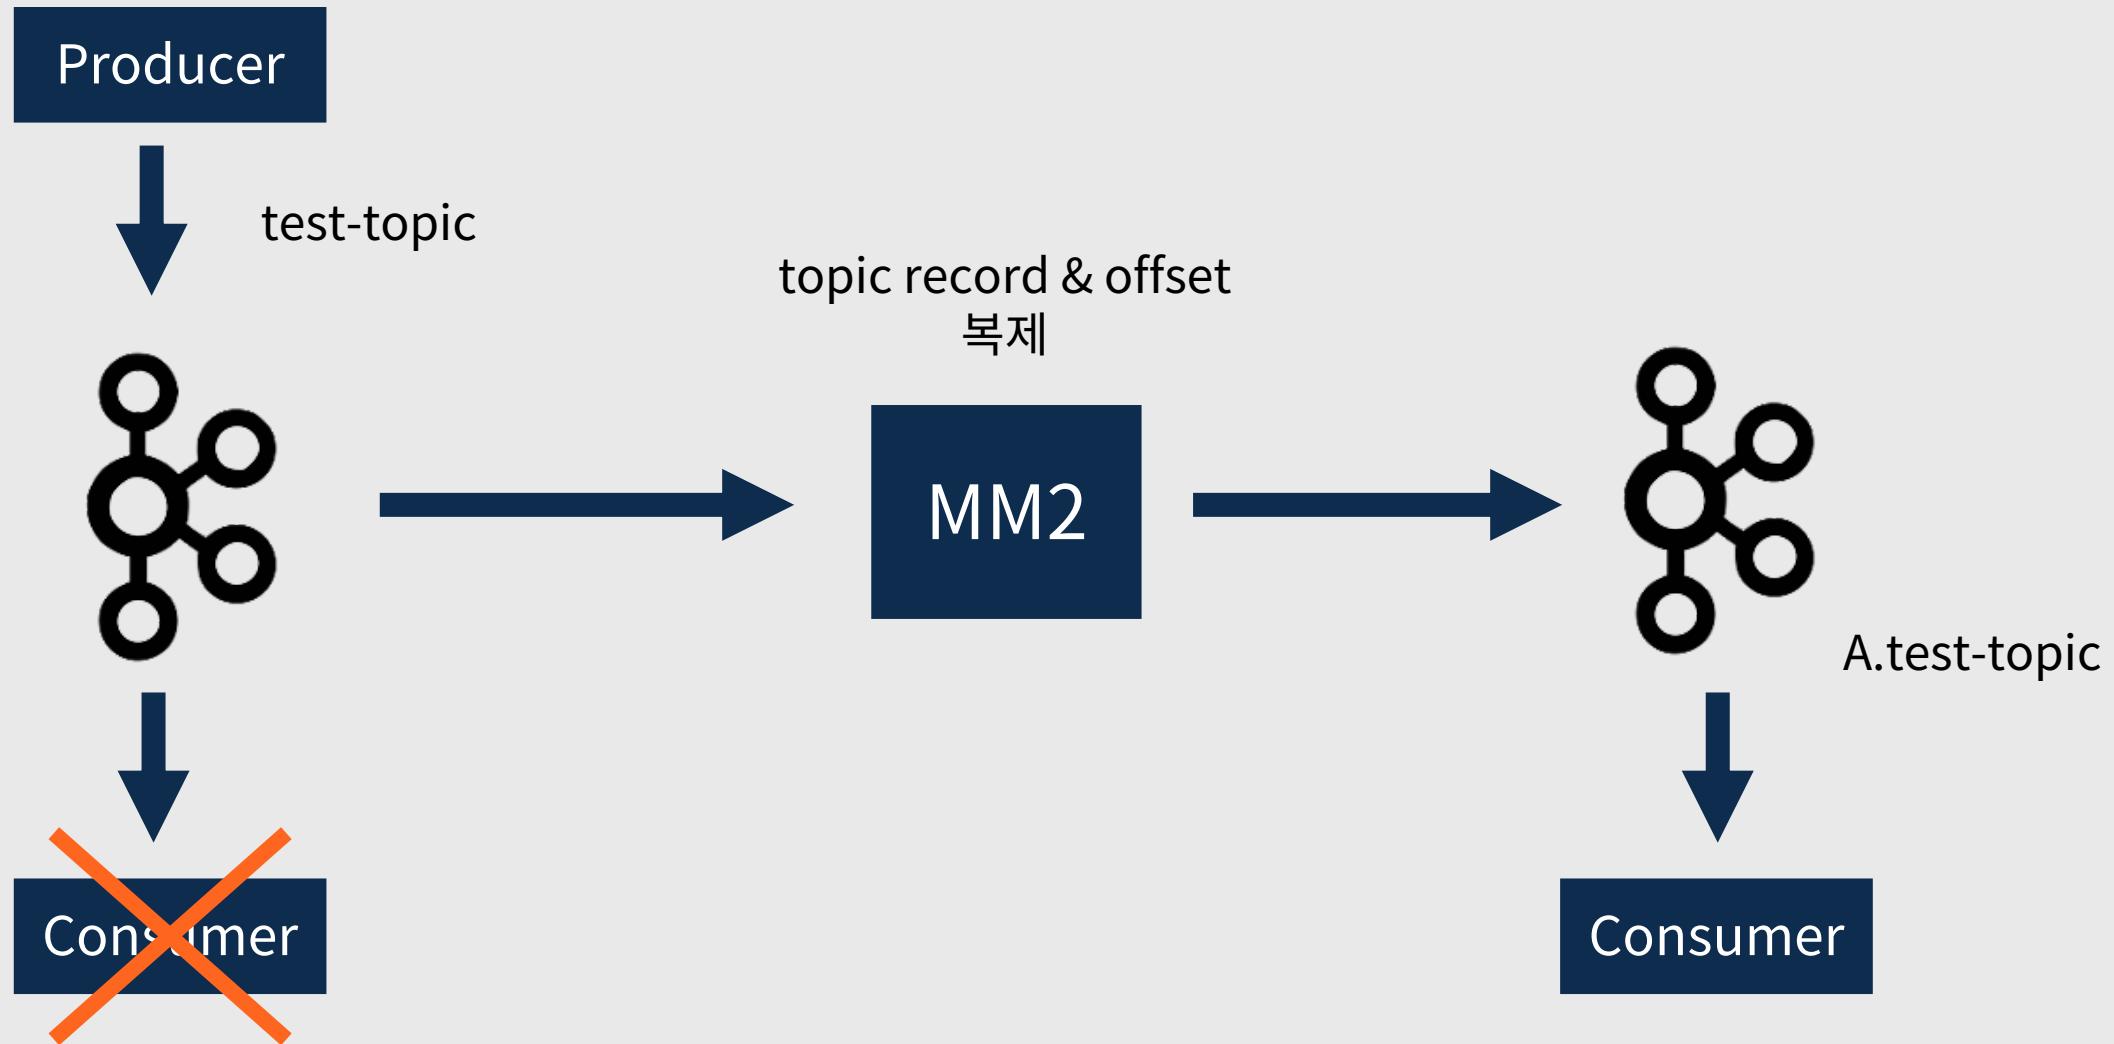


MirrorMaker2 활용

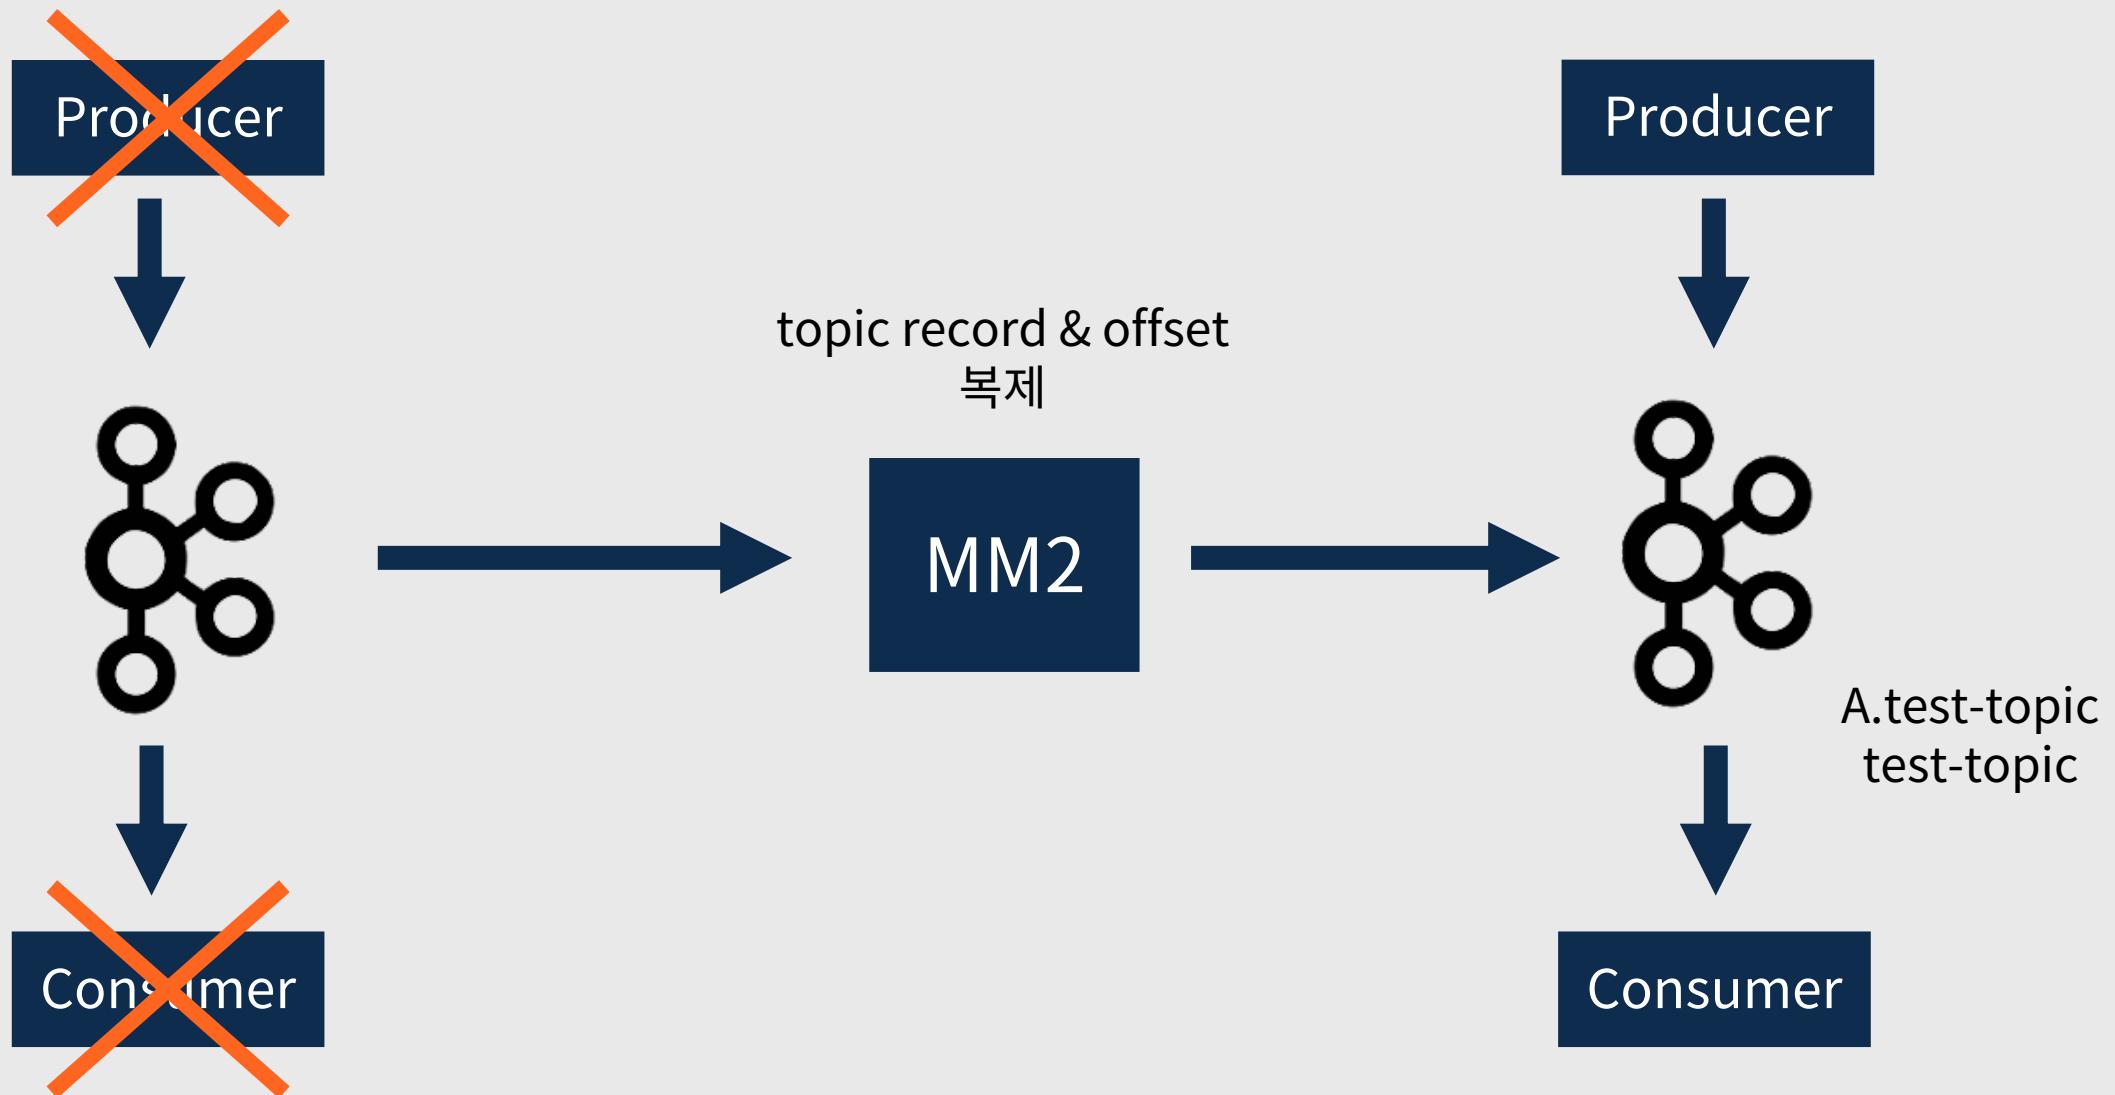
use case 1 - Live Migration



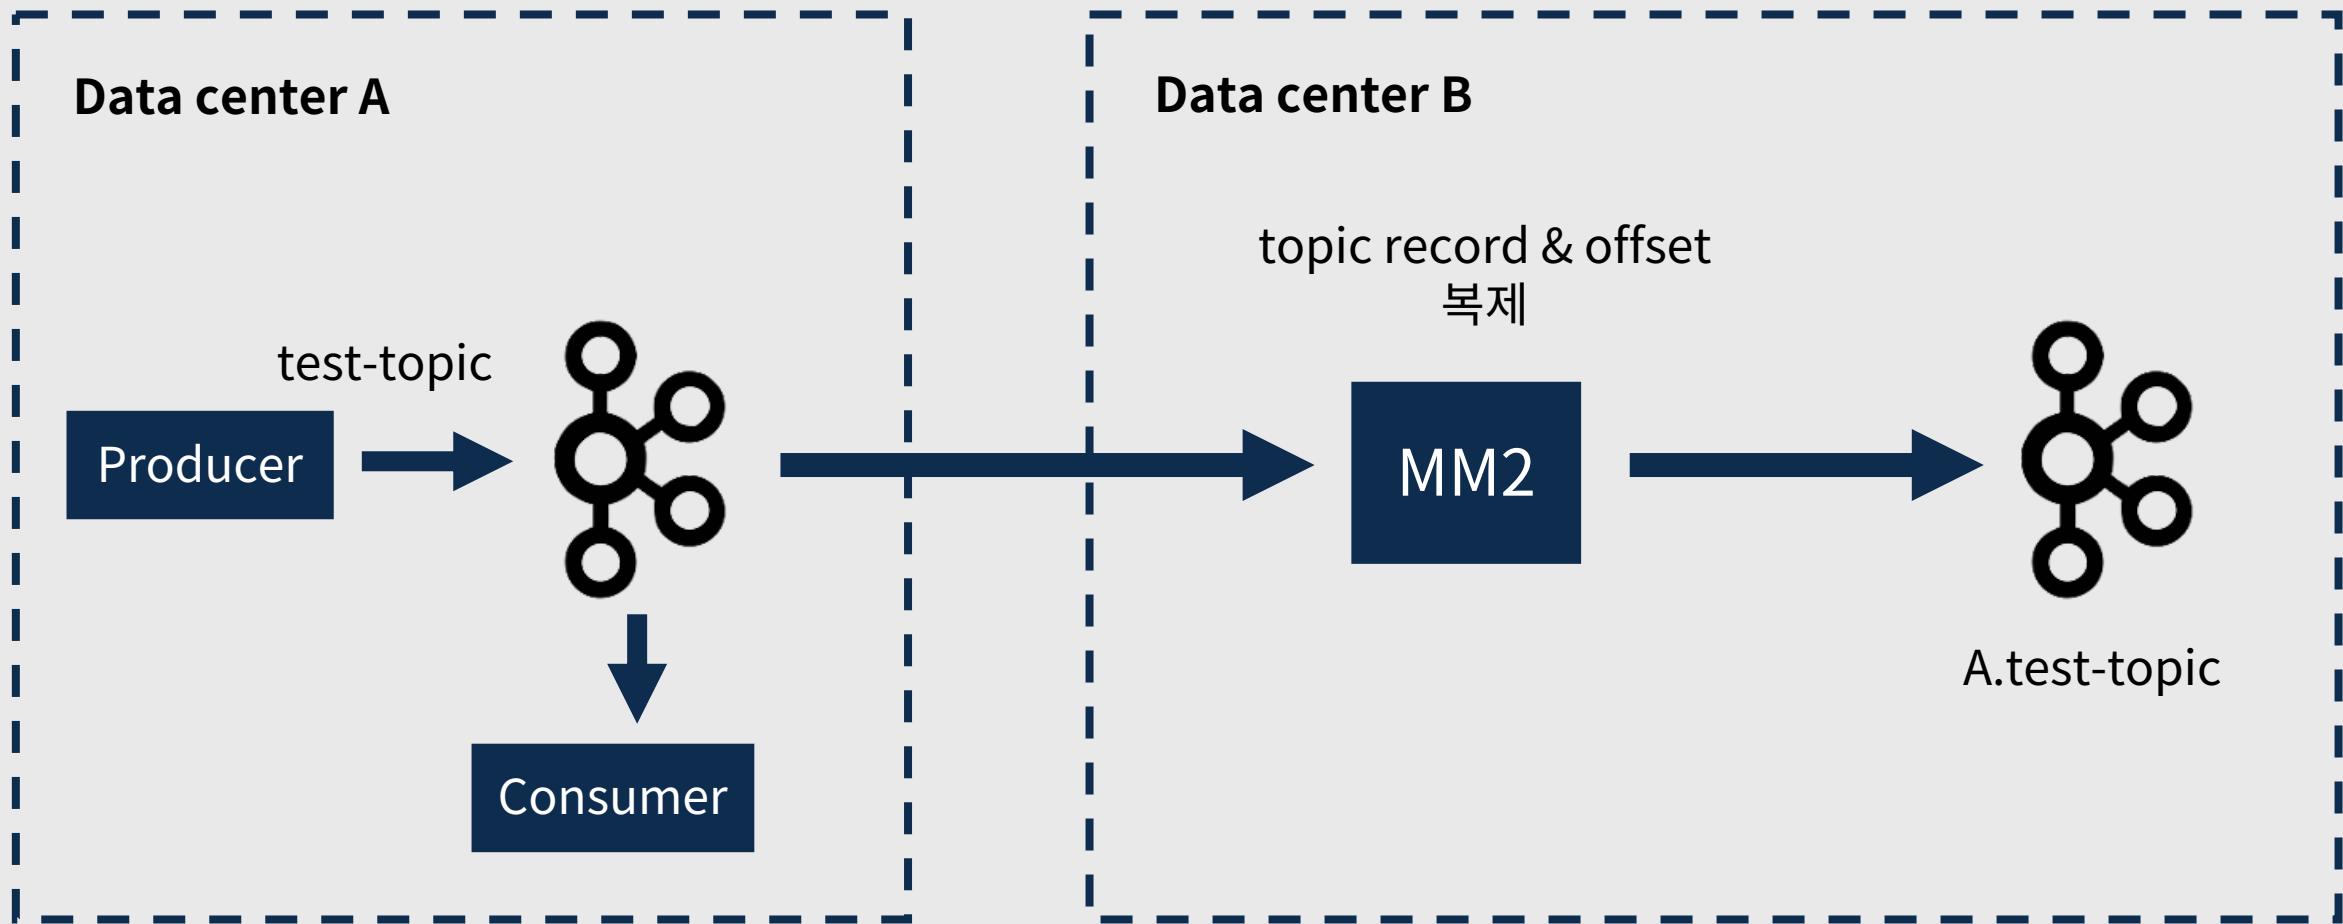
use case 1 - Live Migration



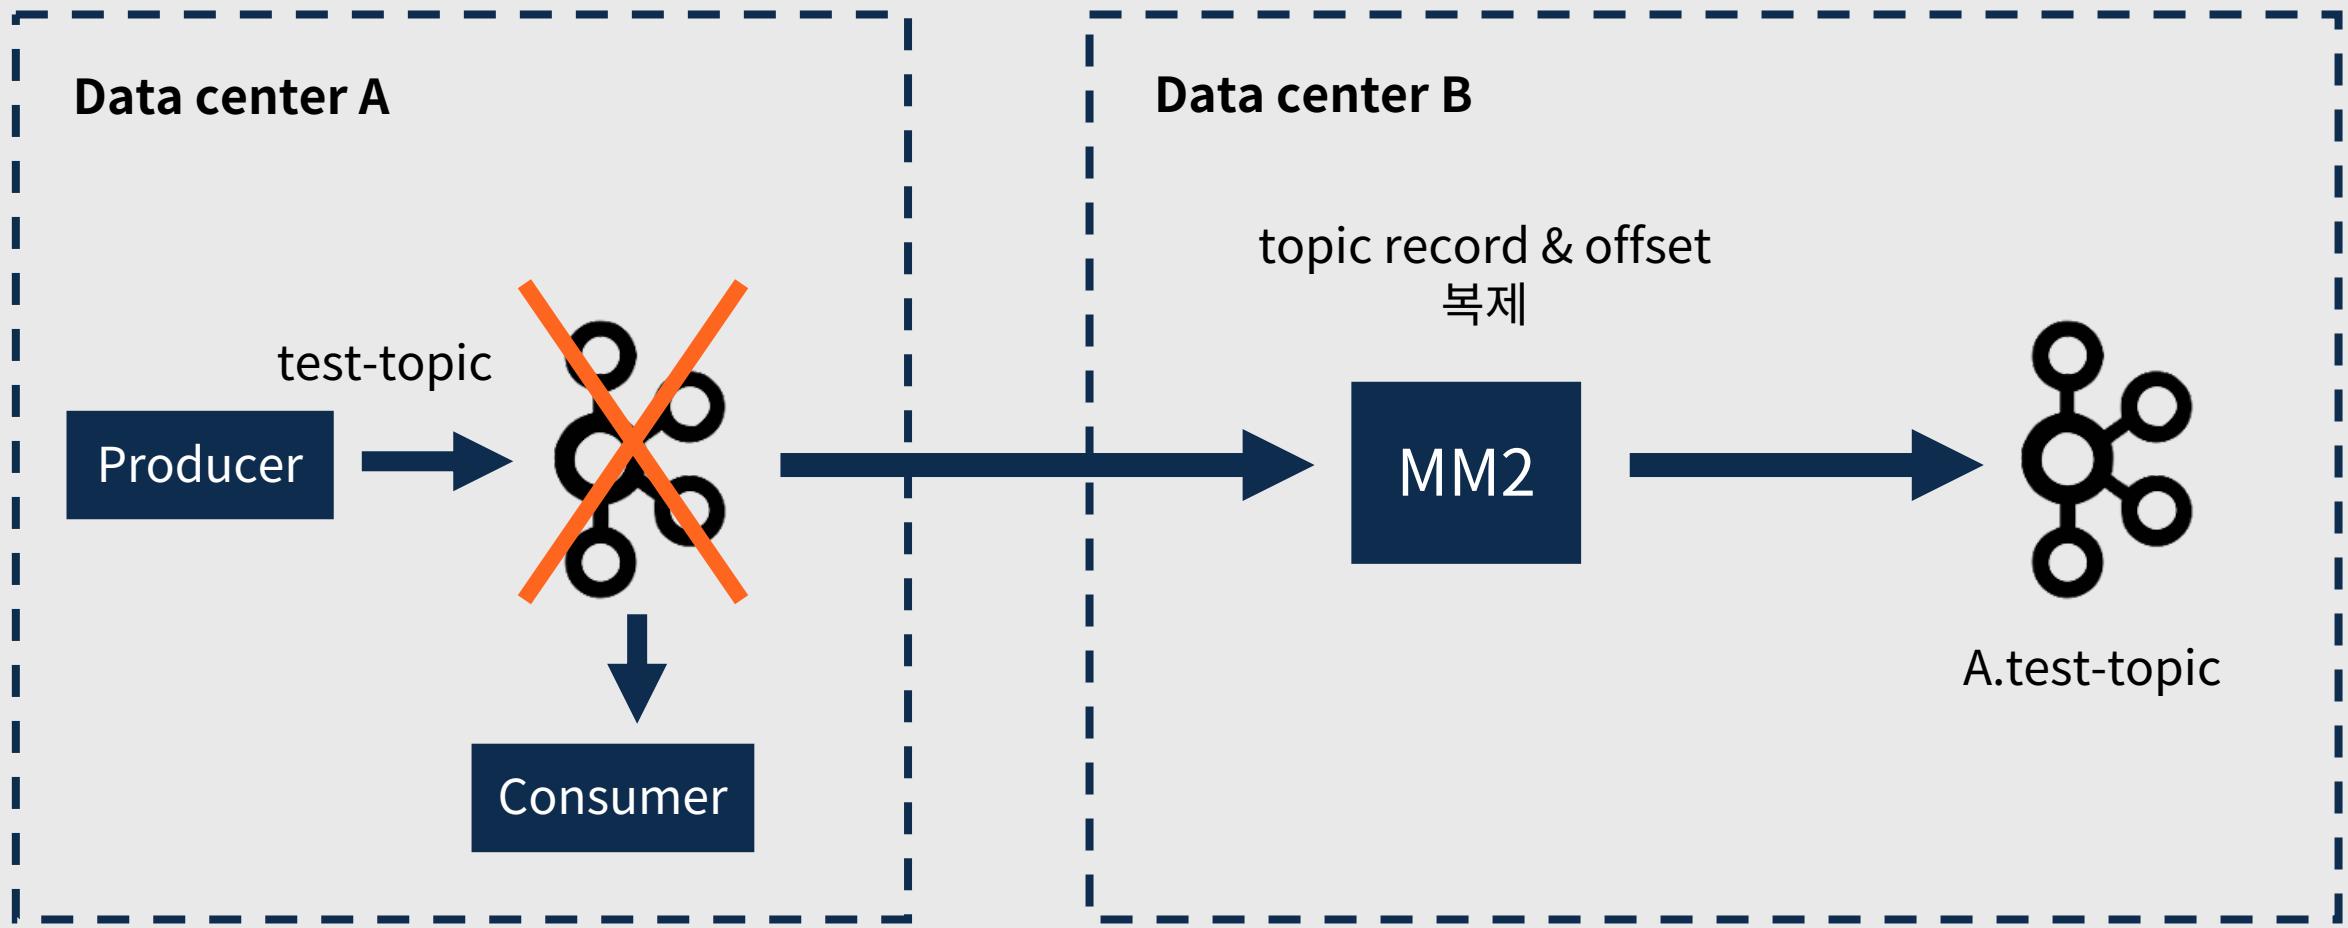
use case 1 - Live Migration



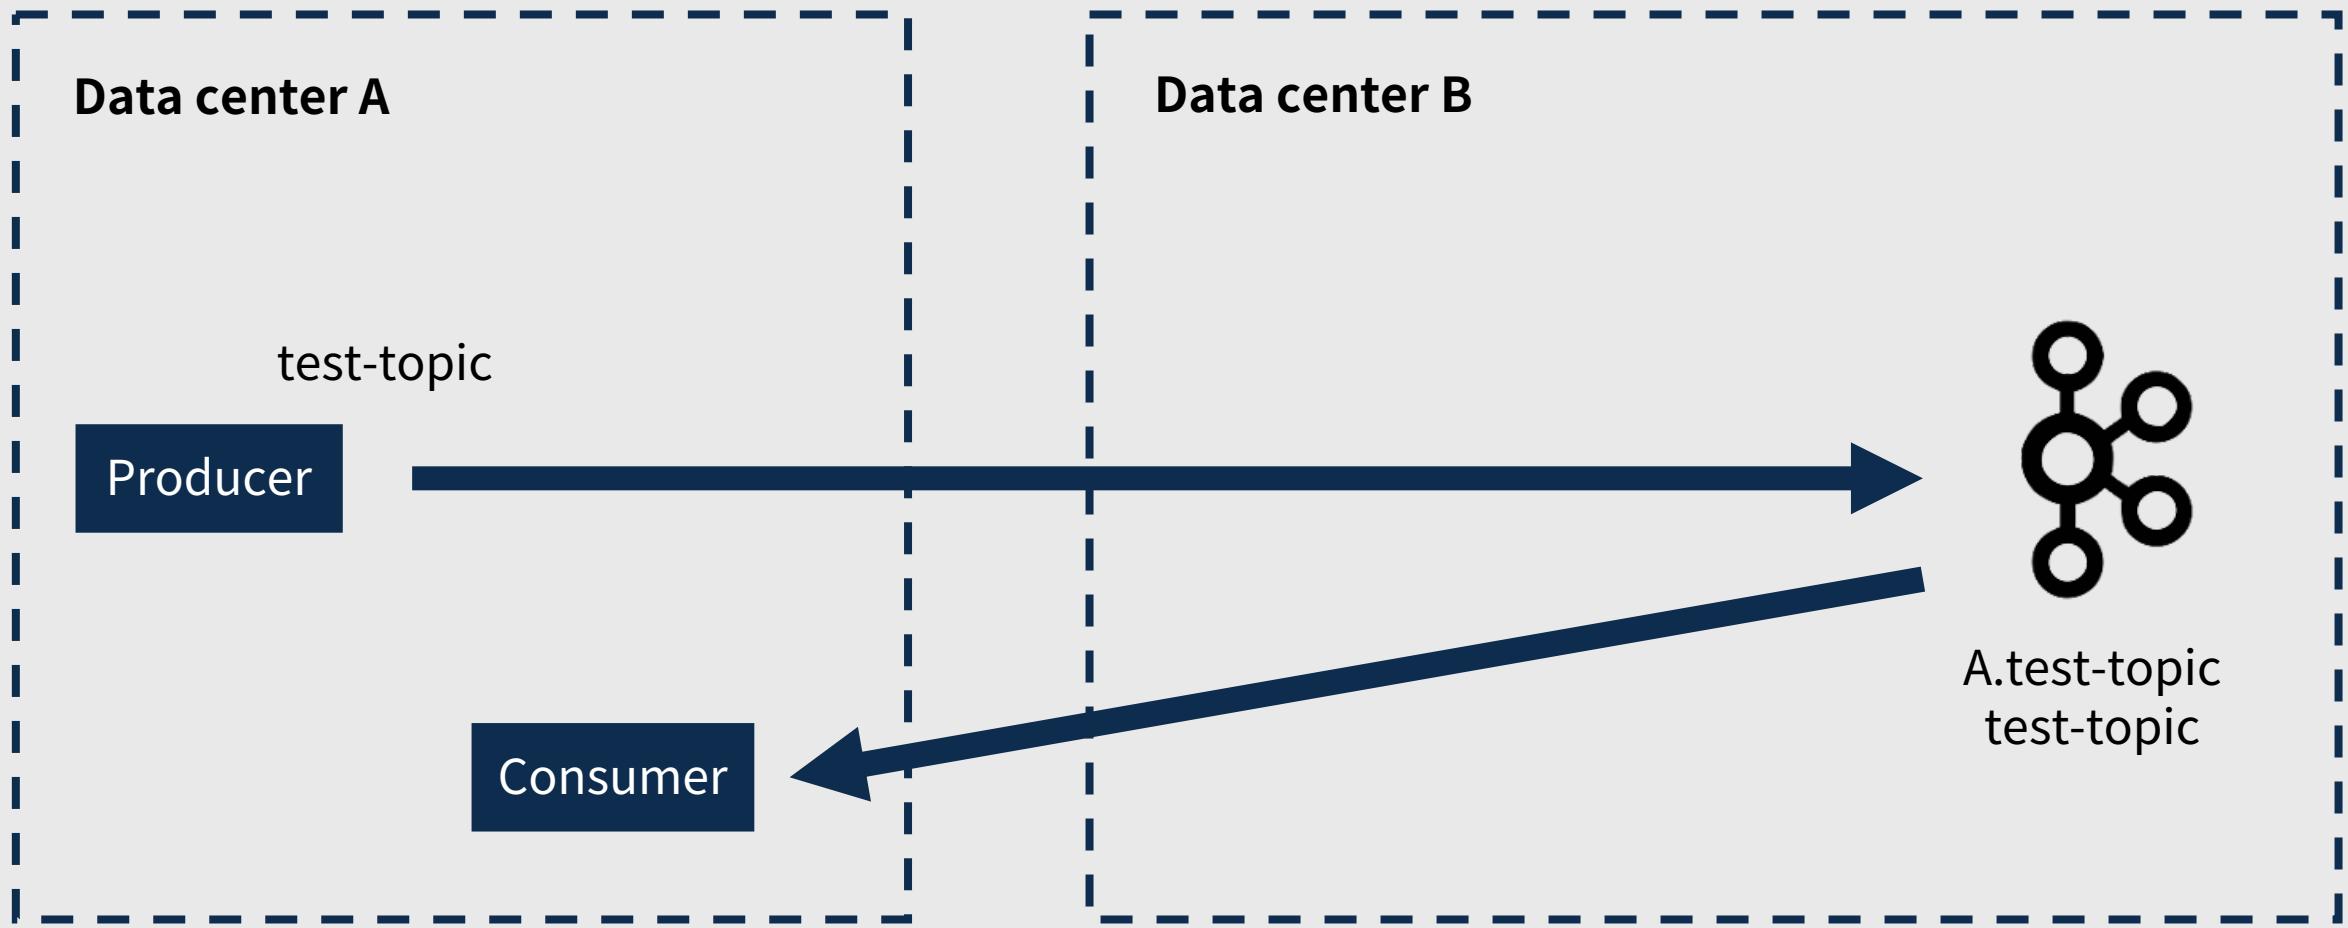
use case 2 - Active/Standy



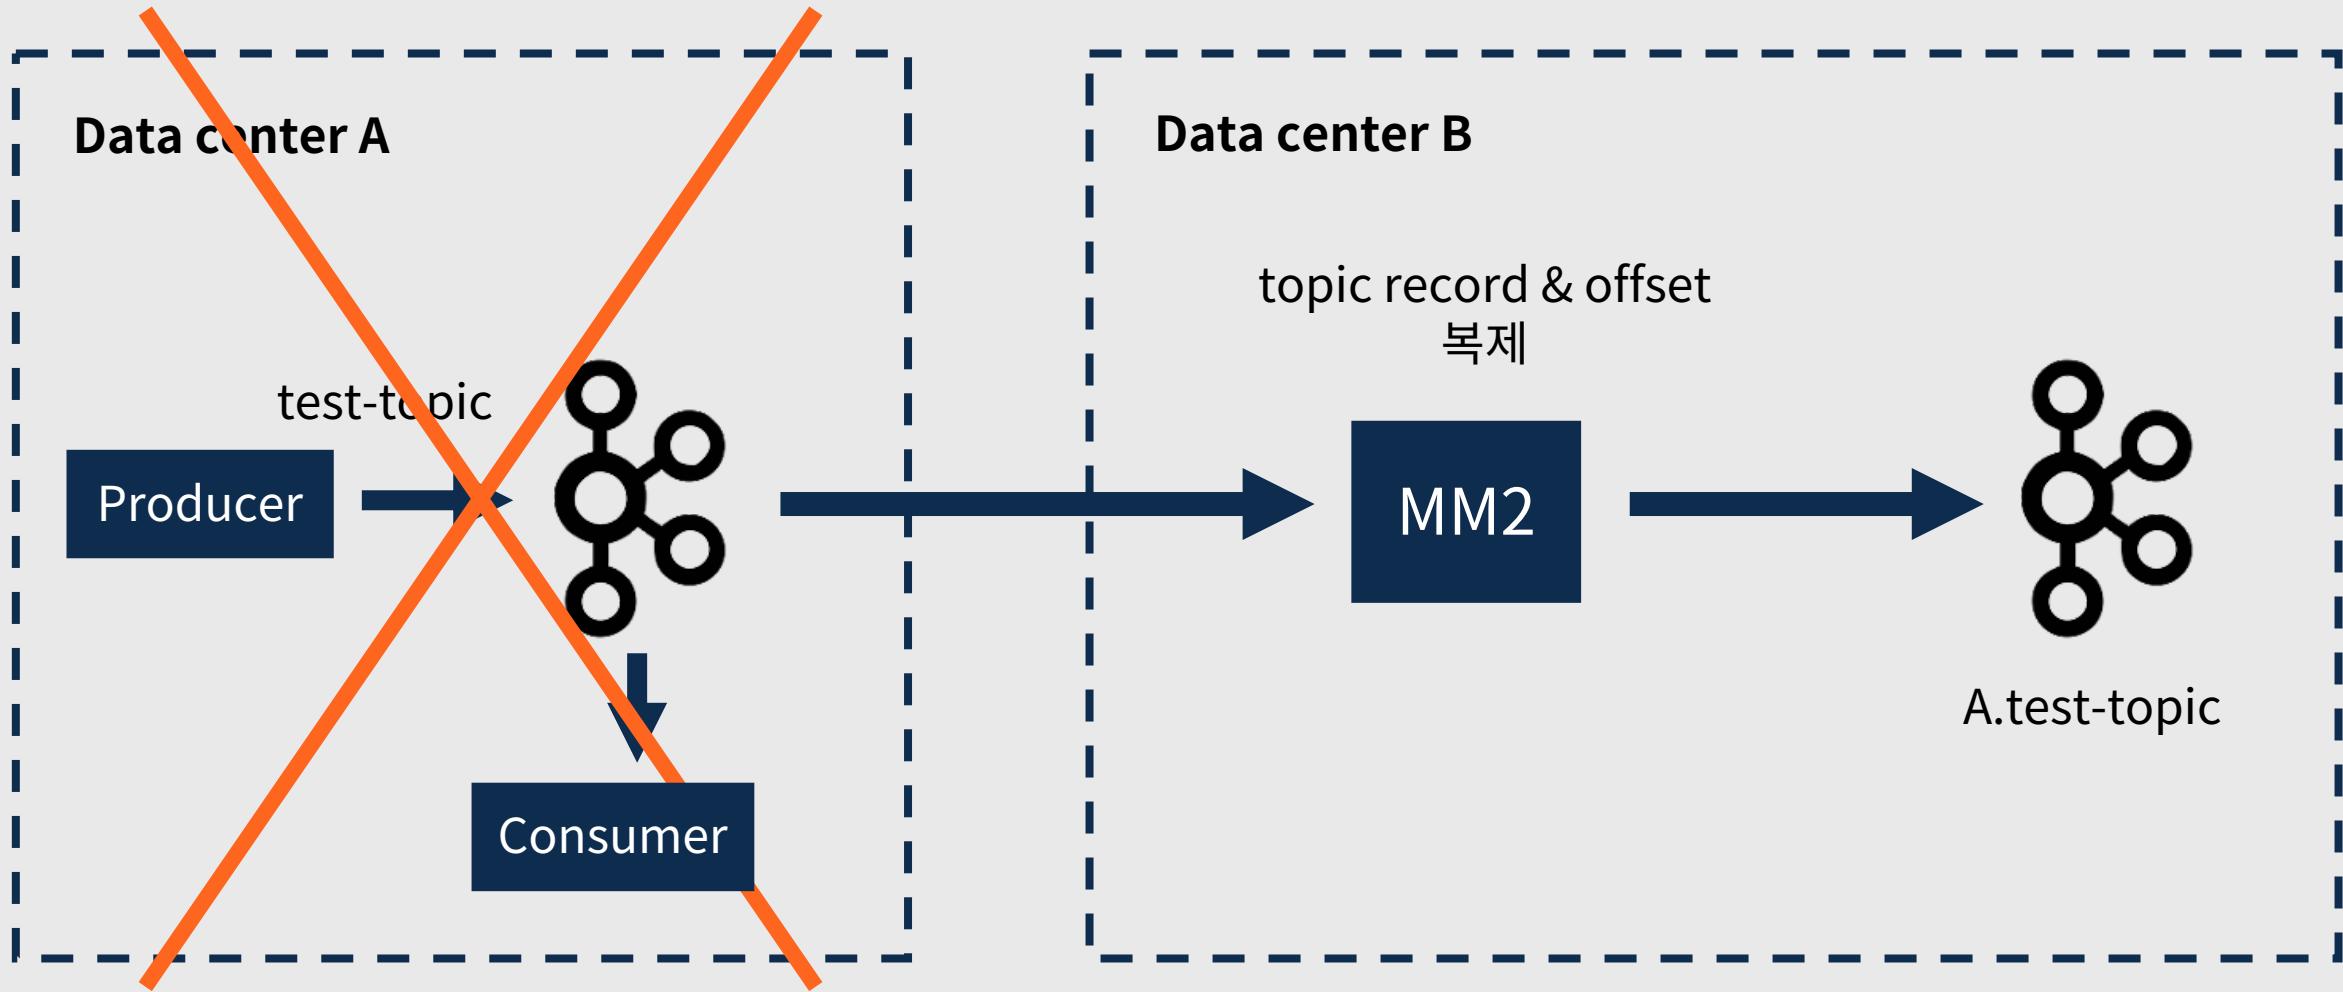
use case 2 - Active/Standy



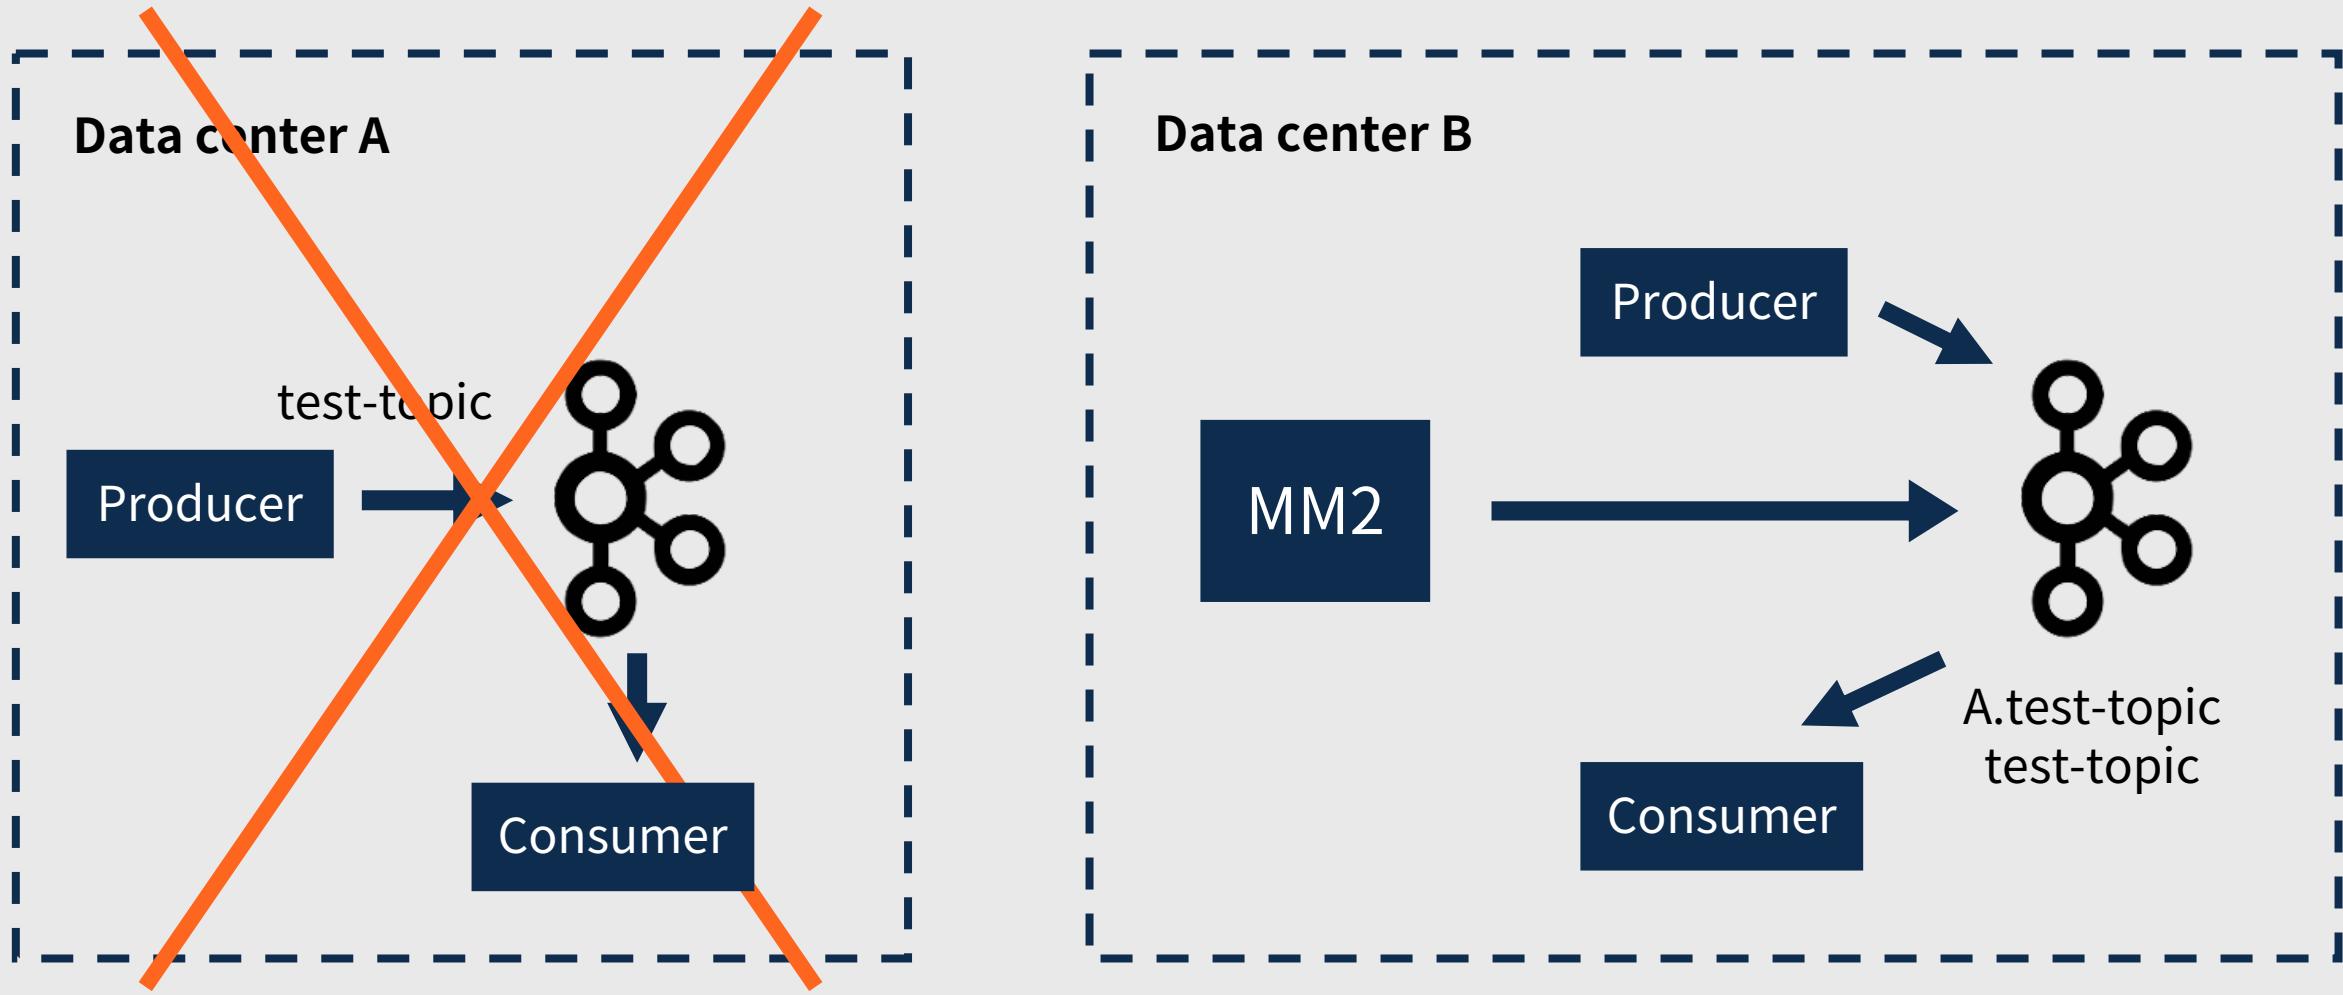
use case 2 - Active/Standby



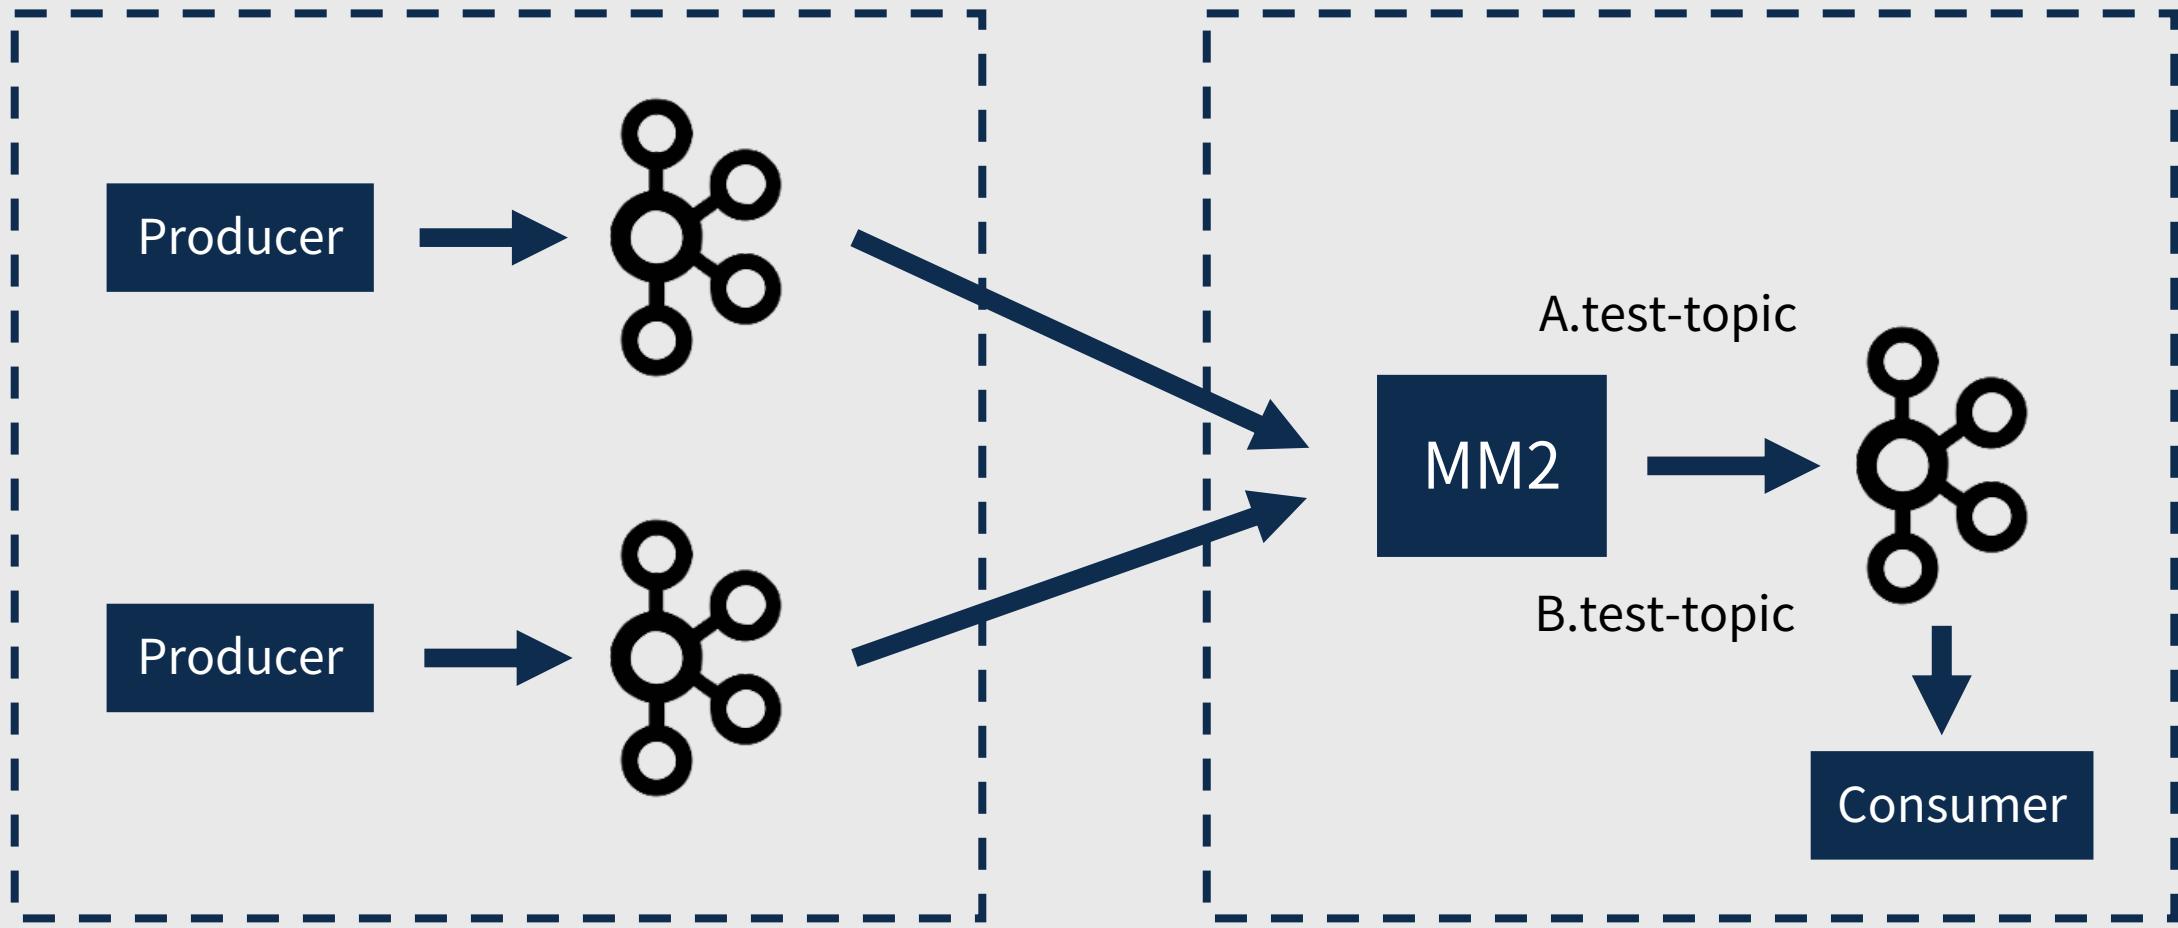
use case 2 - Active/Standy



use case 2 - Active/Standy



use case 3 – Aggregation



MirrorMaker2 모니터링

MirrorMaker2 - 모니터링

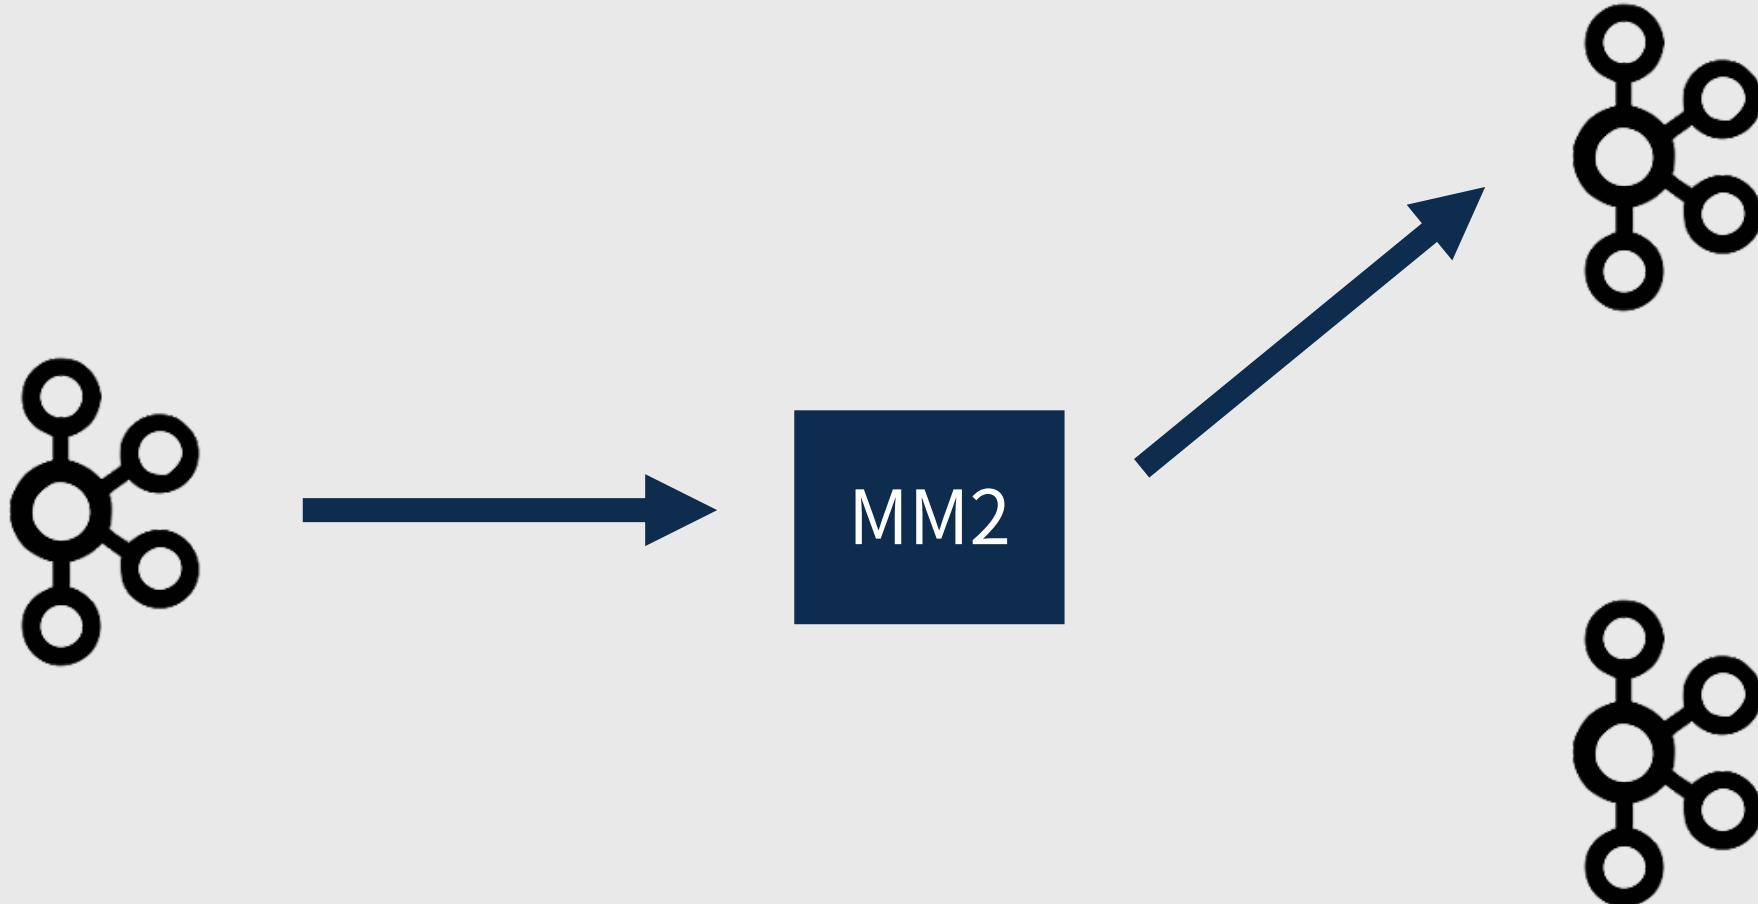
- Throughput/latency per partition
 - kafka.connect.mirror:type=MirrorSourceConnector
 - record-count
 - byte-rate
 - record-age-ms
 - replication-latency-ms
 - offset checkpoint latency
 - kafka.connect.mirror:type=MirrorCheckpointConnector
 - checkpoint-latency-ms

MirrorMaker2 - prometheus & grafana

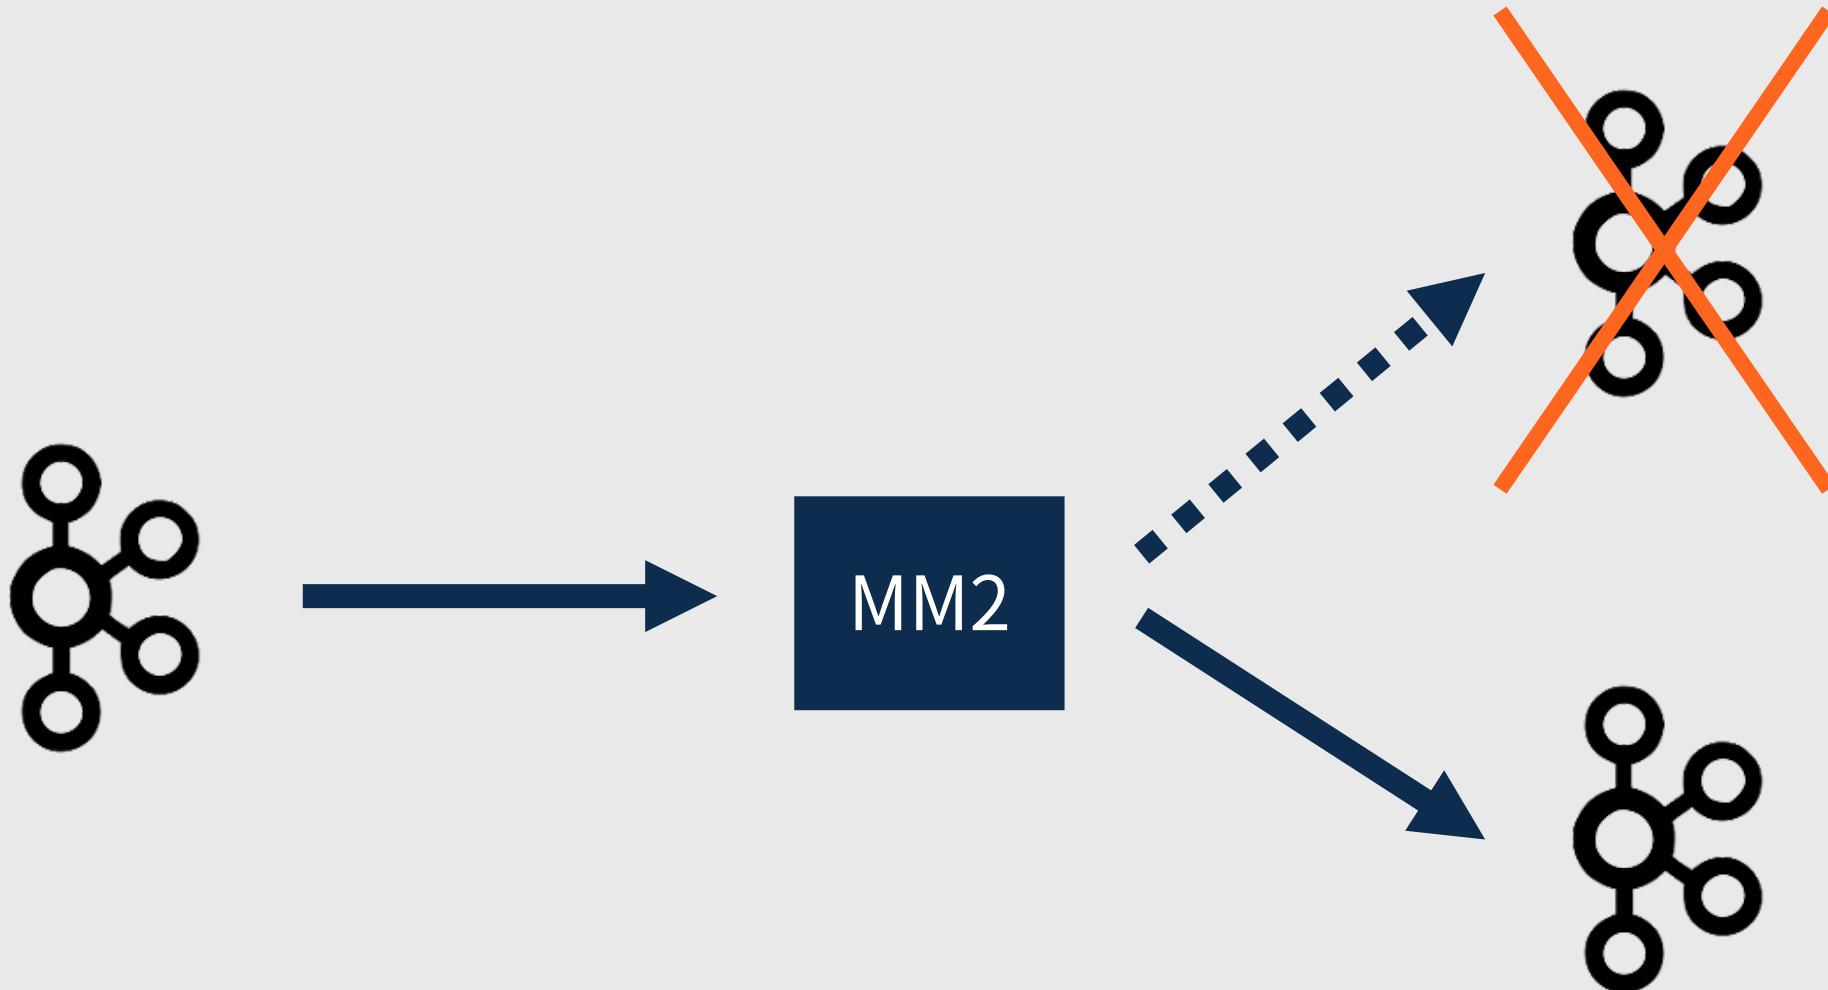


MirrorMaker2 운영 팁

MirrorMaker2 – Migration



MirrorMaker2 - Migration



MirrorMaker2 – Migration

- source connector는 처리한 offset을 저장한다.
 - offset.storage.topic : MirrorMaker2의 설정 (mm2-offset.A.internal)
- MirrorSourceConnector는 재시작할 때 mm2-offset.A.internal 이후 오프셋부터 복제

```
./bin/kafka-console-consumer.sh --bootstrap-server B:9092 \
--topic mm2-offset.A.internal \
--property print.key=true --from-beginning
```

```
header : ["MirrorSourceConnector", {"cluster": "A", "partition": 0, "topic": "test-topic"}]
body : {"offset": 225}
```

MirrorMaker2 – Migration



MM2



mm2-offset.A.internal

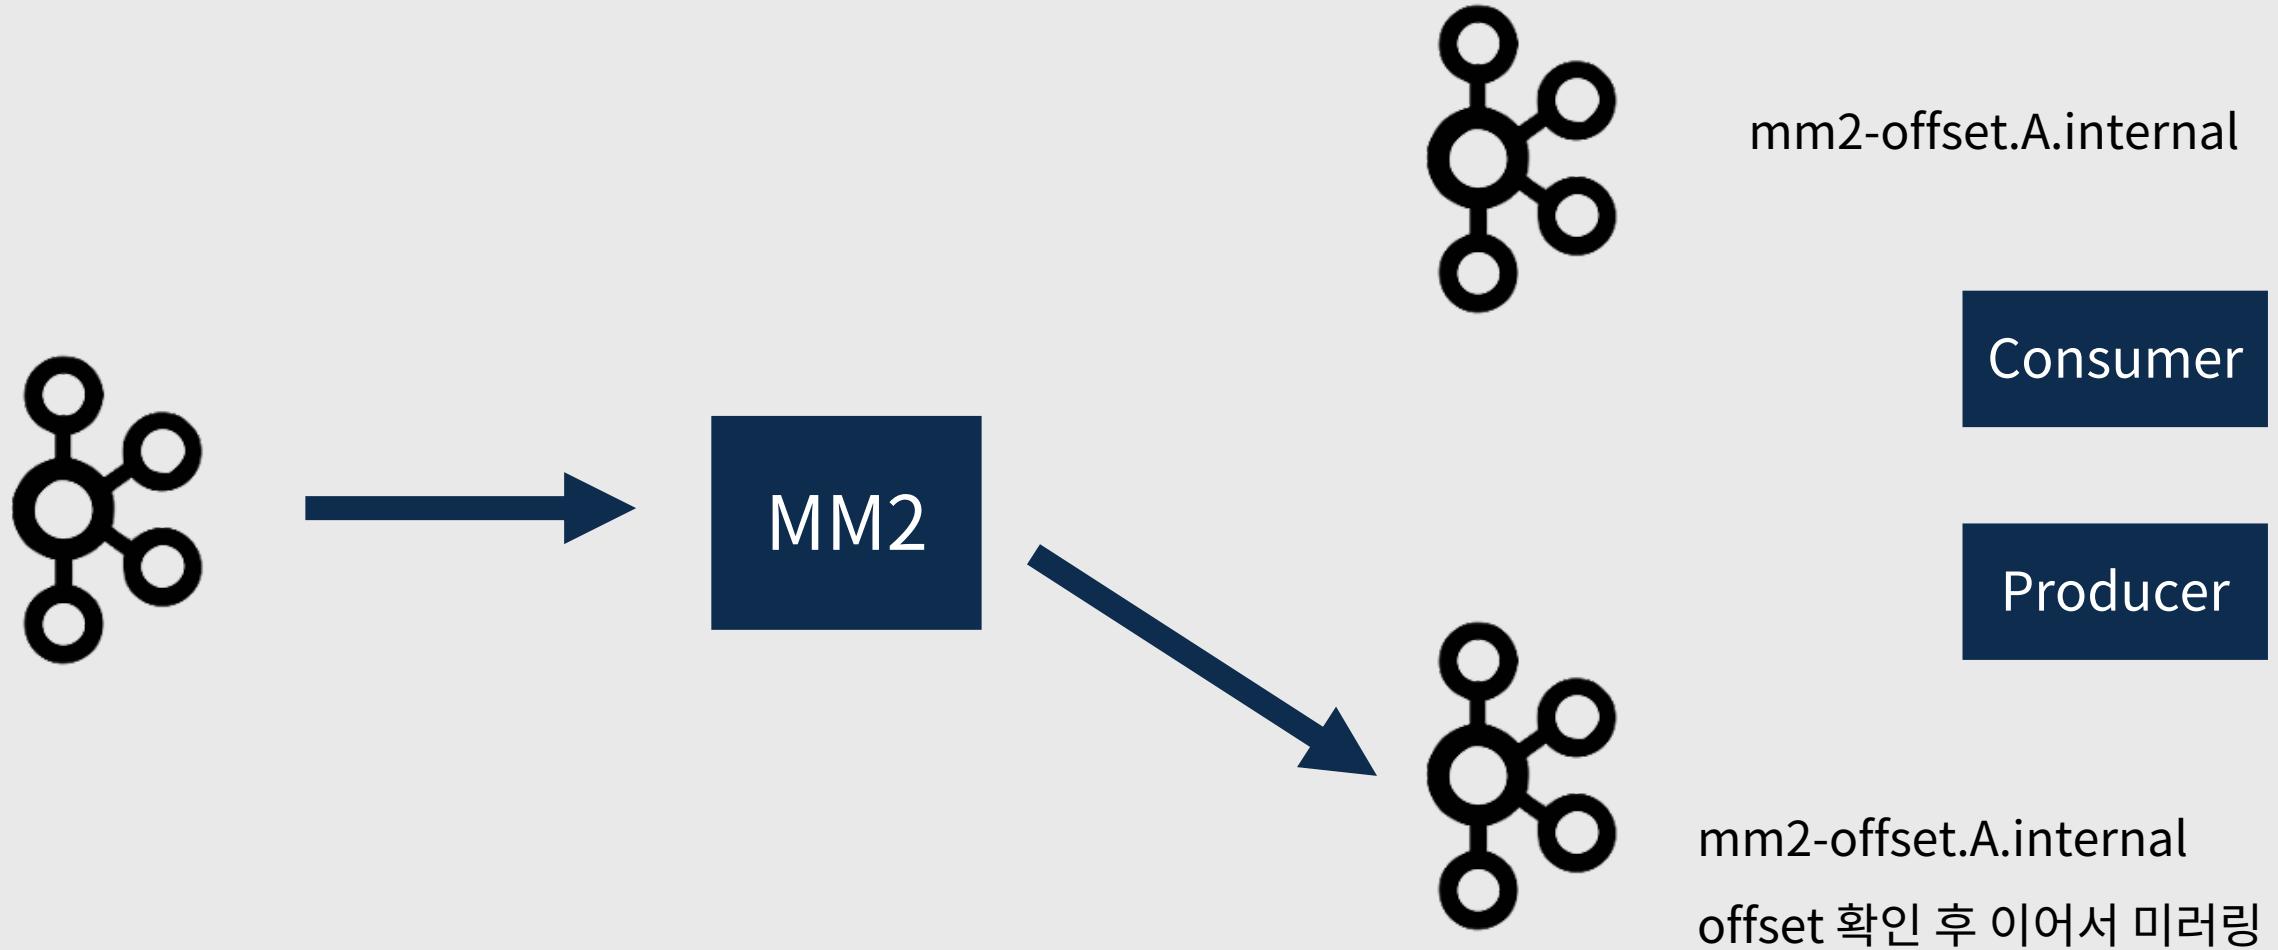
Consumer



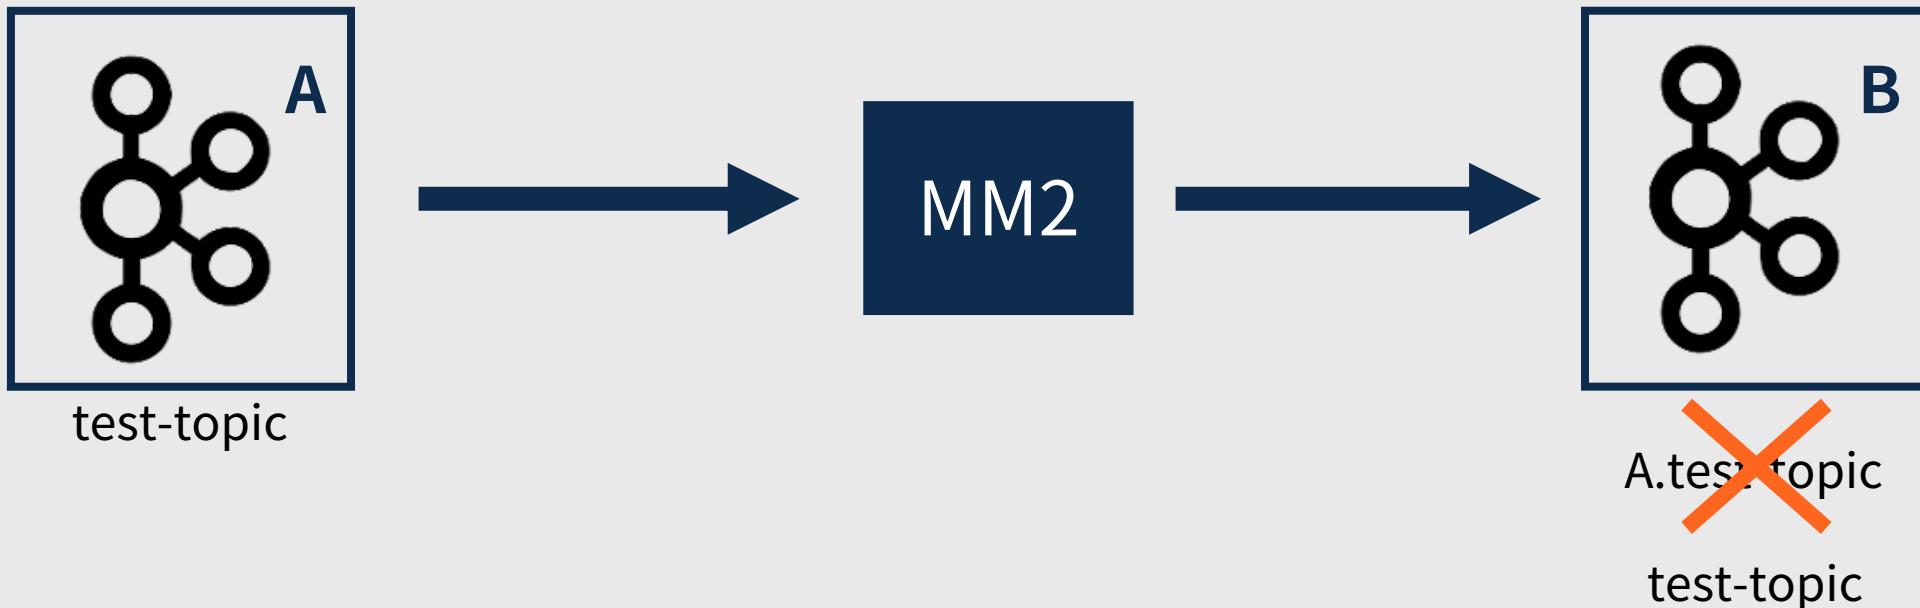
Producer

mm2-offset.A.internal

MirrorMaker2 – Migration



MirrorMaker2 - Prefix 제거



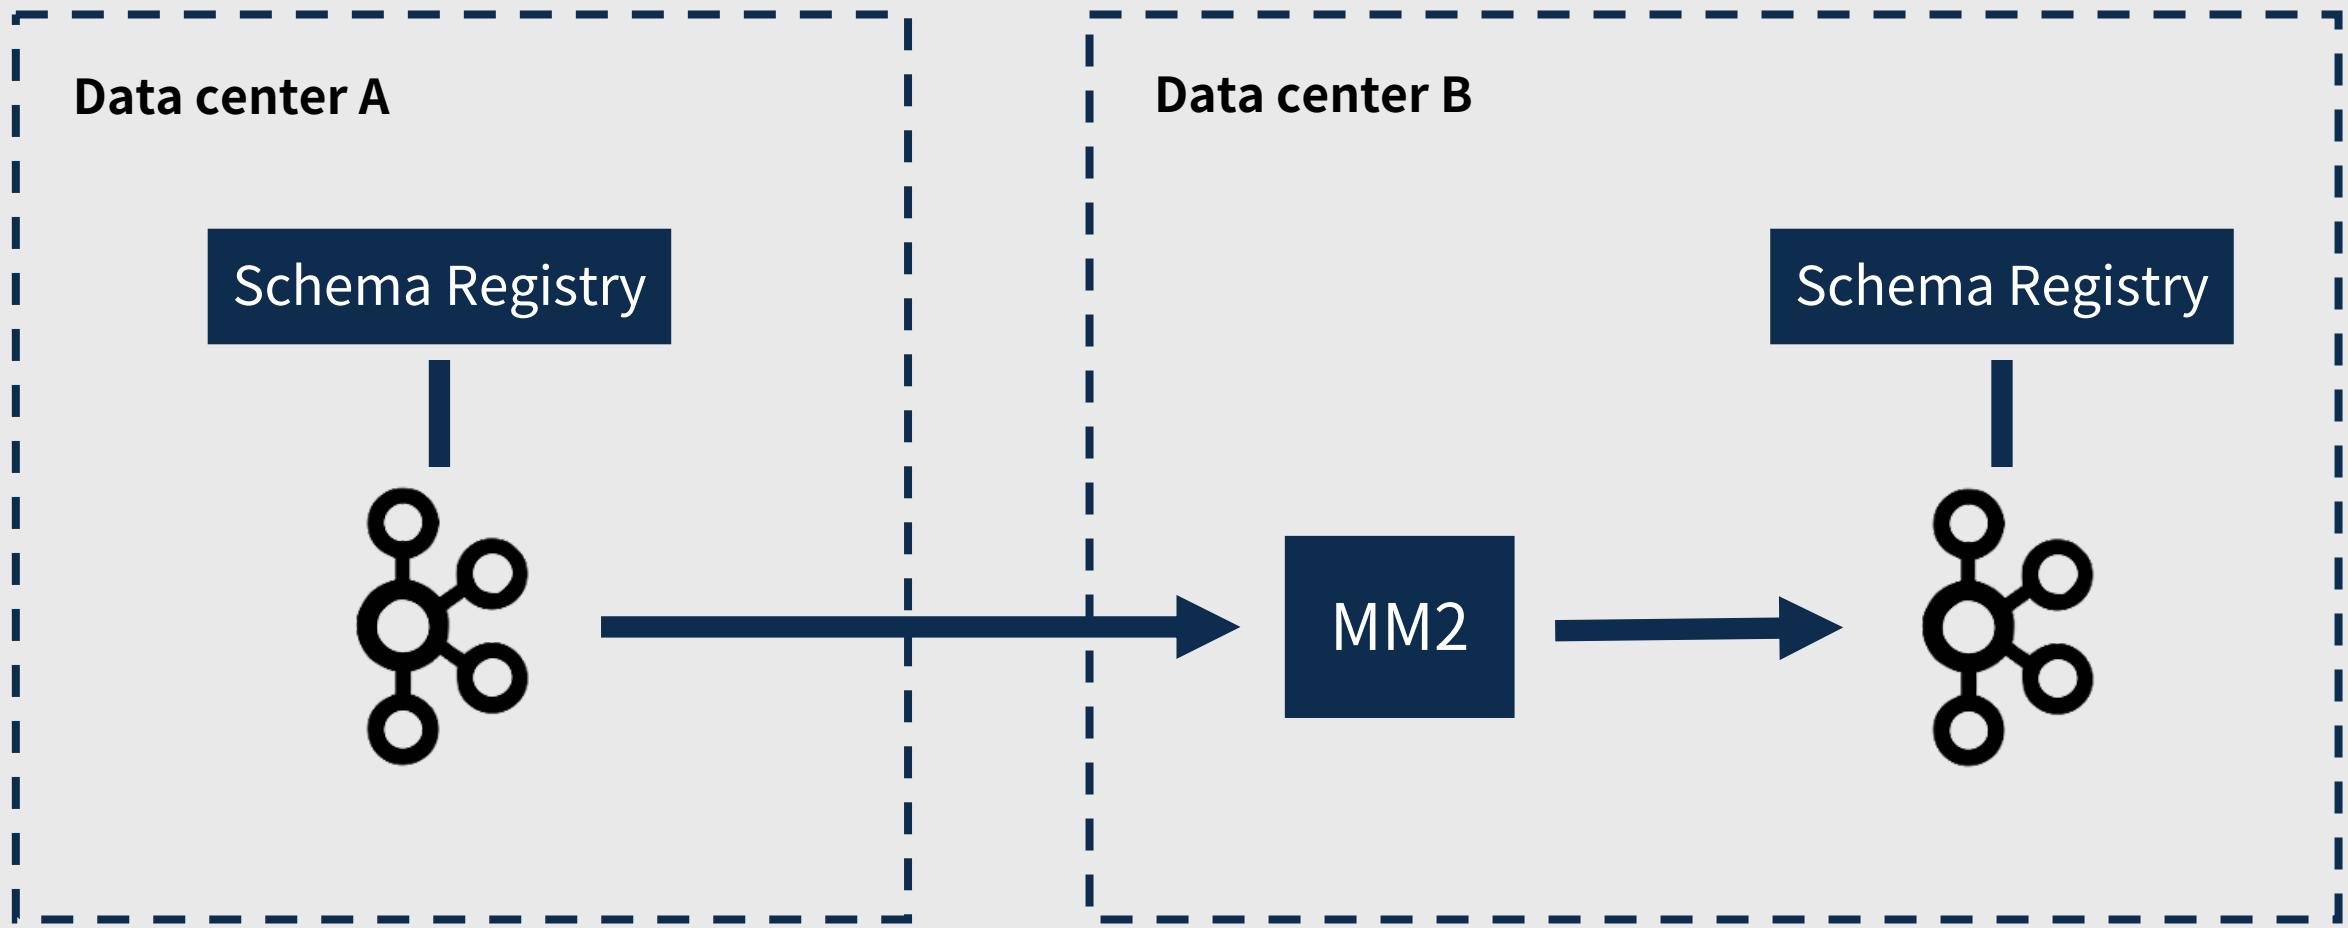
MirrorMaker2 - Prefix 제거

- DefaultReplicationPolicy
 - prefix를 붙여서 토픽 복제하는 클래스
- IdentityReplicationPolicy (KAFKA-9726)
 - prefix 없이 동일 이름으로 토픽 복제

```
replication.policy.class = org.apache.kafka.connect.mirror.IdentityReplicationPolicy
```

- replication.policy.class
 - DefaultReplicationPolicy 클래스 상속 및 ReplicationPolicy, Configurable 구현
 - 다른 복제 정책을 구현하여 적용 가능

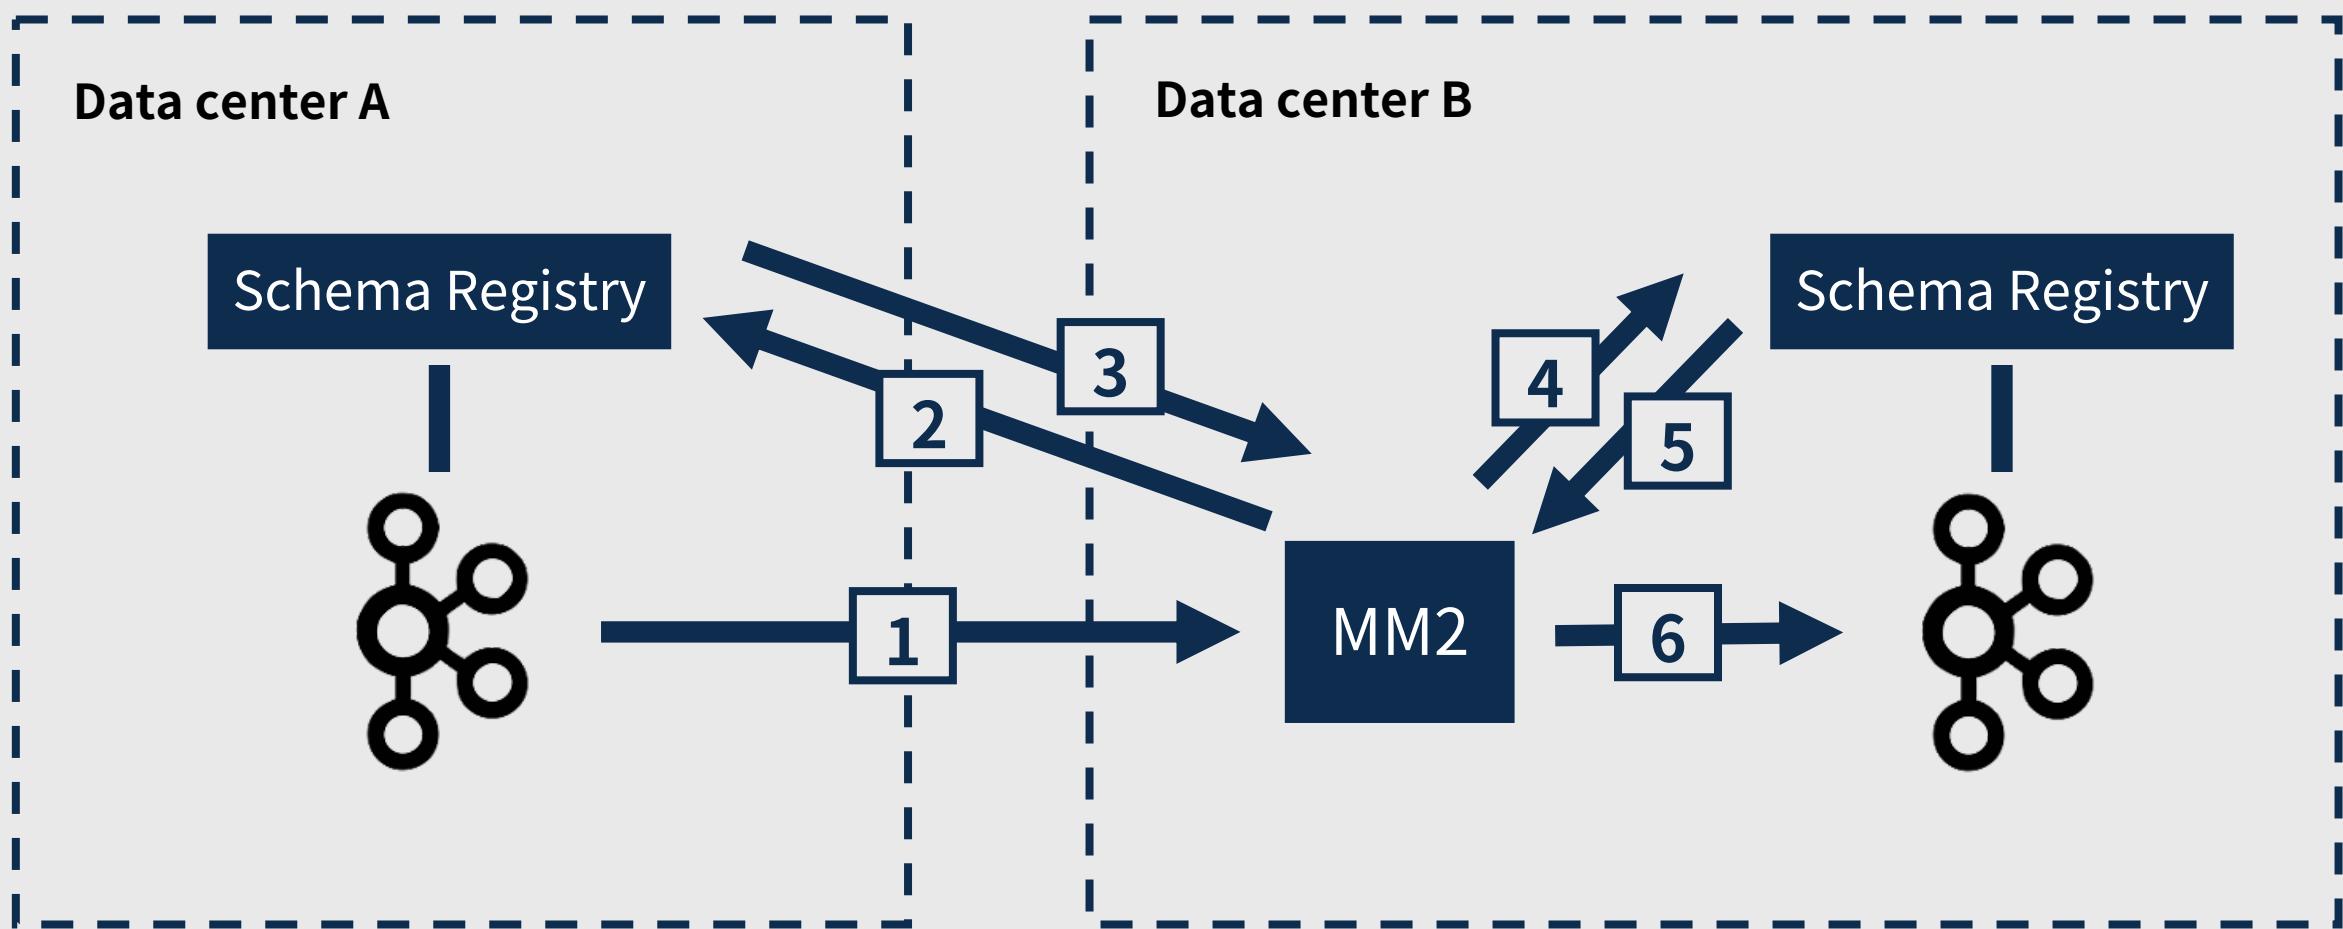
MirrorMaker2 - schema registry



MirrorMaker2 - schema registry

- _schemas 복제가 아닌 schema 신규 등록
- Single Message Transforms (SMT)
 - SMT로 카프카에 데이터 전송 전 변환 가능 (source connector)
- Schema Registry Transfer SMT
 1. source 카프카 레코드의 schema ID 확인
 2. source schema registry에서 schema 확인
 3. target schema registry에 schema 등록
 4. target 카프카에 레코드 전송

MirrorMaker2 - schema registry



감사합니다