

Monitoring Kafka Cluster /w Prometheus

by swoogi@areum.kr

JW Song / Aug. 30, 2022

Something worth What to talk about

- Why important?
- Various Solutions
- Why Prometheus?
- Examples
- Need more metrics
- I can't watch 24/7
- Our implementations

About Speaker

CCAAK & CCDAK since 2020

- 카프카 클러스터 관리 운영 도구 개발 중
 - 2020년 ~
 - 2021년 ~ 다수 고객사 v1 구축 및 운영 중
 - 연내 시즌2 및 Community Edition 출시 목표
- My Role
 - 기술 탐색 및 아키텍처
 - Kafka 관련 SW 패키징
 - zookeeper, kafka, kafka connect, kafka mirror maker2, etc.
 - Config file generator
 - Installation Packages
 - Dockerize
 - Helm Charts

Cluster Health

Why monitoring is so important?

- 카프카 사용 범위 확대
 - 예전에는 로그 수집 파이프라인 등의 2차 시스템으로 도입
 - 카프카가 멈추더라도 애플리케이션에 영향이 거의 없음
 - 카프카 클러스터 복구에 충분한 시간을 확보할 수 있음
 - 중요한 이벤트 메시지를 중계하는 1차 시스템으로 도입이 증가
 - 카프카가 멈추면 실시간으로 이루어지는 애플리케이션 서비스에서 척추 신경에 장애가 생기는 것과 같은 상황 발생할 수 있음
 - 최대한 빠른 시간안에 카프카 클러스터를 복구해야 함
 - 능동적인 모니터링을 통해 장애를 방지하는 것이 최선

Monitoring Kafka

<https://docs.confluent.io/platform/current/kafka/monitoring.html>

- 카프카는 내부 상태를 확인할 수 있는 여러 가지 측정 지표(metrics) 제공
 - broker metrics
 - kafka.server:type=ReplicaManager,name=UnderMinIsrPartitionCount
 - kafka.server:type=ReplicaManager,name=UnderReplicatedPartitions
 - kafka.controller:type=KafkaController,name=OfflinePartitionsCount
 - kafka.controller:type=KafkaController,name=ActiveControllerCount
 - kafka.server:type=FetcherLagMetrics,name=ConsumerLag,clientId=([-.\w]+),topic=([-.\w]+),partition=([0-9]+)
 - zookeeper metrics
 - kafka.server:type=SessionExpireListener,name=SessionState
 - producer metrics
 - Starting with 0.8.2
 - consumer metrics
 - Starting with 0.9.0.0

Critical Cases

Broker or cluster stopped

- 데이터 디스크 full에 의한 브로커 중지
 - log.dirs로 지정된 디스크 볼륨에 남은 공간이 '0'이 되면서 브로커 프로세스는 살아 있지만 기능이 멈춤
 - log size per partition 등의 정보를 기본으로 제공하고 있지만, 브로커 서버 데이터 볼륨의 전체 크기 정보가 없어 부족함
- 주키퍼 네트워크 단절에 의한 클러스터 중단
 - 주키퍼 앙상블이 다른 네트워크 존에 구성됨
 - 네트워크 작업에 의해 카프카 브로커와 주키퍼 사이 네트워크가 단절됨
 - 계속된 메타데이터 업데이트 실패로 인해 브로커 프로세스 기능이 멈춤

Various Solutions

Monitoring kafka cluster

1. metric.reporters

- 브로커 설정에 metric reporter를 등록
 - org.apache.kafka.common.metrics.MetricsReporter 인터페이스 구현
 - 대표적으로 org.apache.kafka.common.metrics.JmxReporter 클래스가 있음
 - JmxReporter는 metrics.reporters 설정에 항상 자동으로 포함됨
 - confluent는 io.confluent.metrics.reporter.ConfluentMetricsReporter를 통해 카프카 토픽으로 메트릭을 수집

2. JConsole via JMX

- JConsole은 JMX를 통해 JVM의 상태를 바로 확인할 수 있는 도구

3. Prometheus /w JMX exporter

4. Elastic Stack /w Jolokia + Metricbeat

JMX?

Java Management eXtensions

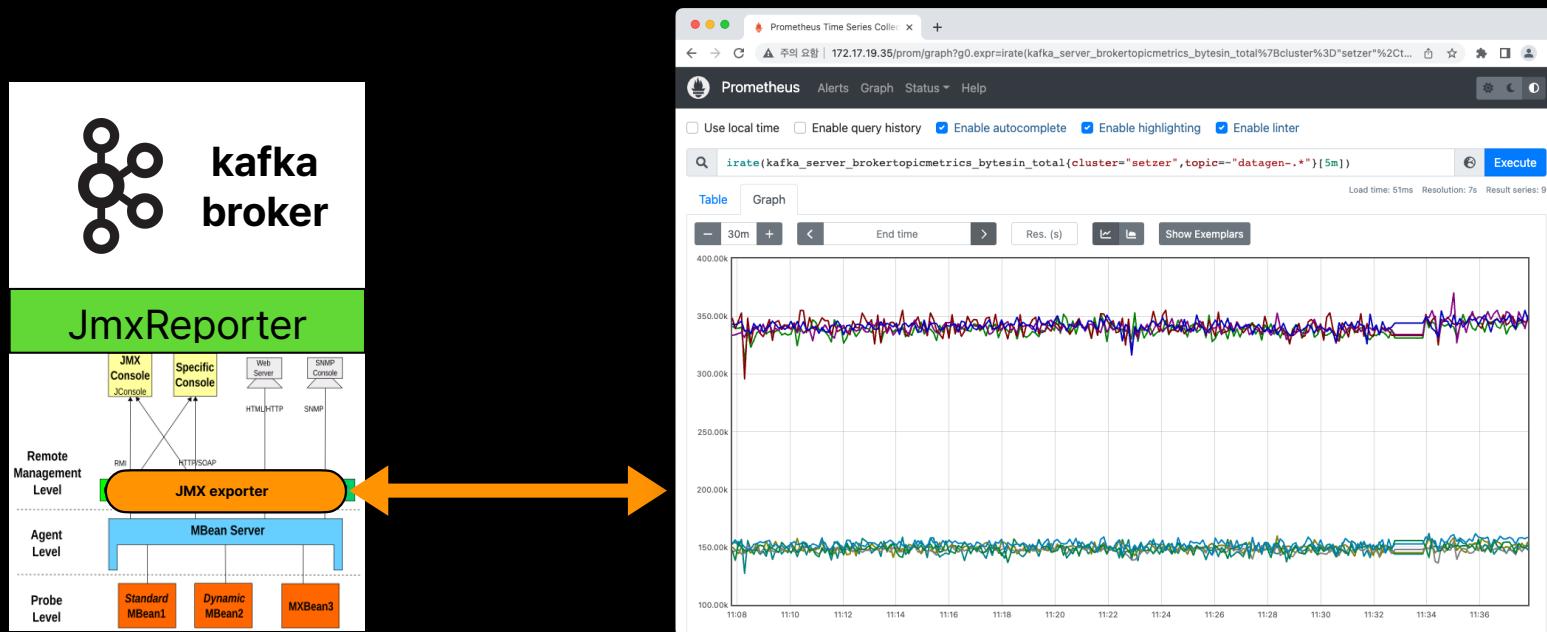
- 자바 애플리케이션 관리와 모니터링을 위한 기술
 - 애플리케이션, 시스템 객체, 장치 등을 모니터링 및 관리
 - MBean(Managed Bean)에 의해 리소스가 제공됨
 - JDK와 함께 배포되는 JConsole을 통해 JVM에 접속 가능



Scraping Metrics

Prometheus /w JMX exporter

- JMX exporter
 - JMX 메트릭을 HTTP(S)로 노출해주는 javaagent
 - javaagent는 주 애플리케이션과 동일한 JVM에서 실행되며 최초 진입점 `main()`보다 먼저 실행될 수 있는 특별한 애플리케이션
 - Jolokia + Metricbeat도 비슷한 형태



Expose Metrics

Starting kafka brokers /w JMX exporter

- 카프카 시작 시 특별한 환경 변수
 - JMX_PORT, KAFKA_JMX_OPTS
 - LOG_DIR, KAFKA_LOG4J_OPTS
 - KAFKA_OPTS
 - KAFKA_HEAP_OPTS, KAFKA_JVM_PERFORMANCE_OPTS
- javaagent : JMX exporter
 - https://github.com/prometheus/jmx_exporter
 - https://github.com/prometheus/jmx_exporter/tree/master/example_configs

```
# START KAFKA BROKER
KAFKA_OPTS=-javaagent:/path/to/jmx_prometheus_javaagent-0.17.0.jar=9404:/path/to/kafka-2_0_0.yml \
bin/kafka-server-start.sh -daemon ./config/server.properties
```

```
# TEST JMX exporter
curl localhost:9404/metrics
```

Exporter Rules

Configuring JMX exporter

```
1 lowercaseOutputName: true
2
3 rules:
4 # very specific rule for cluster id
5 - pattern: 'kafka.server<type=KafkaServer, name=ClusterId>>Value: (.+)'
6   name: kafka_server_kafkaserver_clusterid
7   value: 1
8   labels:
9     clusterId: "$1"
10    help: "kafka server JMX metric cluster id"
11    type: GAUGE
12 ## very specific rule for app-info
13 - pattern: kafka.(\w+)<type=app-info>>start-time-ms
14   name: kafka_$1_start_time_seconds
15   help: "Kafka $1 JMX metric start time seconds"
16   type: GAUGE
17   valueFactor: 0.001
18 - pattern: 'kafka.(\w+)<type=app-info>>(commit-id|version): (.+)'
19   name: kafka_$1_info
20   value: 1
21   labels:
22     $2: "$3"
23   help: "Kafka $1 JMX metric info version and commit-id"
24   type: GAUGE
25 # Special cases and very specific rules
26 - pattern : kafka.server<type=(.+), name=(.+), clientId=(.+), topic=(.+), partition=(.*)>>Value
27   name: kafka_server_$1_$2
28   type: GAUGE
29   labels:
30     clientId: "$3"
31     topic: "$4"
32     partition: "$5"
33 - pattern : kafka.server<type=(.+), name=(.+), clientId=(.+), brokerHost=(.+), brokerPort=(.+)>>Value
34   name: kafka_server_$1_$2
35   type: GAUGE
36   labels:
37     clientId: "$3"
38     broker: "$4:$5"
39 - pattern : kafka.coordinator.(\w+)<type=(.+), name=(.+)>>Value
40   name: kafka_coordinator_$1_$2_$3
41   type: GAUGE
42
43 # Generic per-second counters with 0-2 key/value pairs
44 - pattern: kafka.(\w+)<type=(.+), name=(.+)>>Value
45   name: kafka_$1_$2_$3_total
```

Extra Rules

Example Config

Scraping kafka metrics from brokers

```
1 ## prometheus.yml ##
2
3 ##### job : kafka #####
4 - job_name: kafka
5   scrape_interval: 5s
6   file_sd_configs:
7     - files:
8       - targets/kafka.yaml
9       refresh_interval: 5m
10
11 ##### job : kafka_node #####
12 - job_name: kafka_node
13   scrape_interval: 5s
14   file_sd_configs:
15     - files:
16       - targets/kafka_node.yaml
17       refresh_interval: 5m
```

```
1 ## kafka.yaml ##
2
3 - targets:
4   - "172.17.25.20:9404"
5   labels:
6     cluster: "ccs"
7     brokerId: "0"
8 - targets:
9   - "172.17.25.21:9404"
10  labels:
11    cluster: "ccs"
12    brokerId: "1"
13 - targets:
14   - "172.17.25.22:9404"
15  labels:
16    cluster: "ccs"
17    brokerId: "2"
18 - targets:
19   - "172.17.25.40:9404"
20  labels:
21    cluster: "setzer"
22    brokerId: "0"
23 - targets:
24   - "172.17.25.41:9404"
25  labels:
26    cluster: "setzer"
27    brokerId: "1"
28 - targets:
29   - "172.17.25.42:9404"
30  labels:
31    cluster: "setzer"
32    brokerId: "2"
```

Metrics Example

kafka.server:type=ReplicaManager,name=UnderMinIsrPartitionCount

The screenshot shows the Prometheus Time Series Collector interface with two panels of metrics results.

Panel 1 (Top):

- Query: `sum(kafka_cluster_partition_underminisr{cluster="setzer"})`
- Load time: 39ms Resolution: 345s Result series: 1
- Table view: Shows a single row with Evaluation time, value 0, and a timestamp 0.

Panel 2 (Bottom):

- Query: `sum(kafka_cluster_partition_underminisr{cluster="ccs"})`
- Load time: 16ms Resolution: 14s Result series: 1
- Table view: Shows a single row with Evaluation time, value 12, and a timestamp 0.

A red box highlights Panel 2, and a blue button at the bottom left says "Add Panel".

Metrics Example

kafka.server:type=ReplicaManager,name=UnderReplicatedPartitions

The screenshot shows the Prometheus Time Series Collector interface with two panels of metrics.

Top Panel:

- Query: `sum(kafka_cluster_partition_underreplicated{cluster="setzer"})`
- Evaluation time: {}
- Result: 0
- Load time: 42ms Resolution: 345s Result series: 1

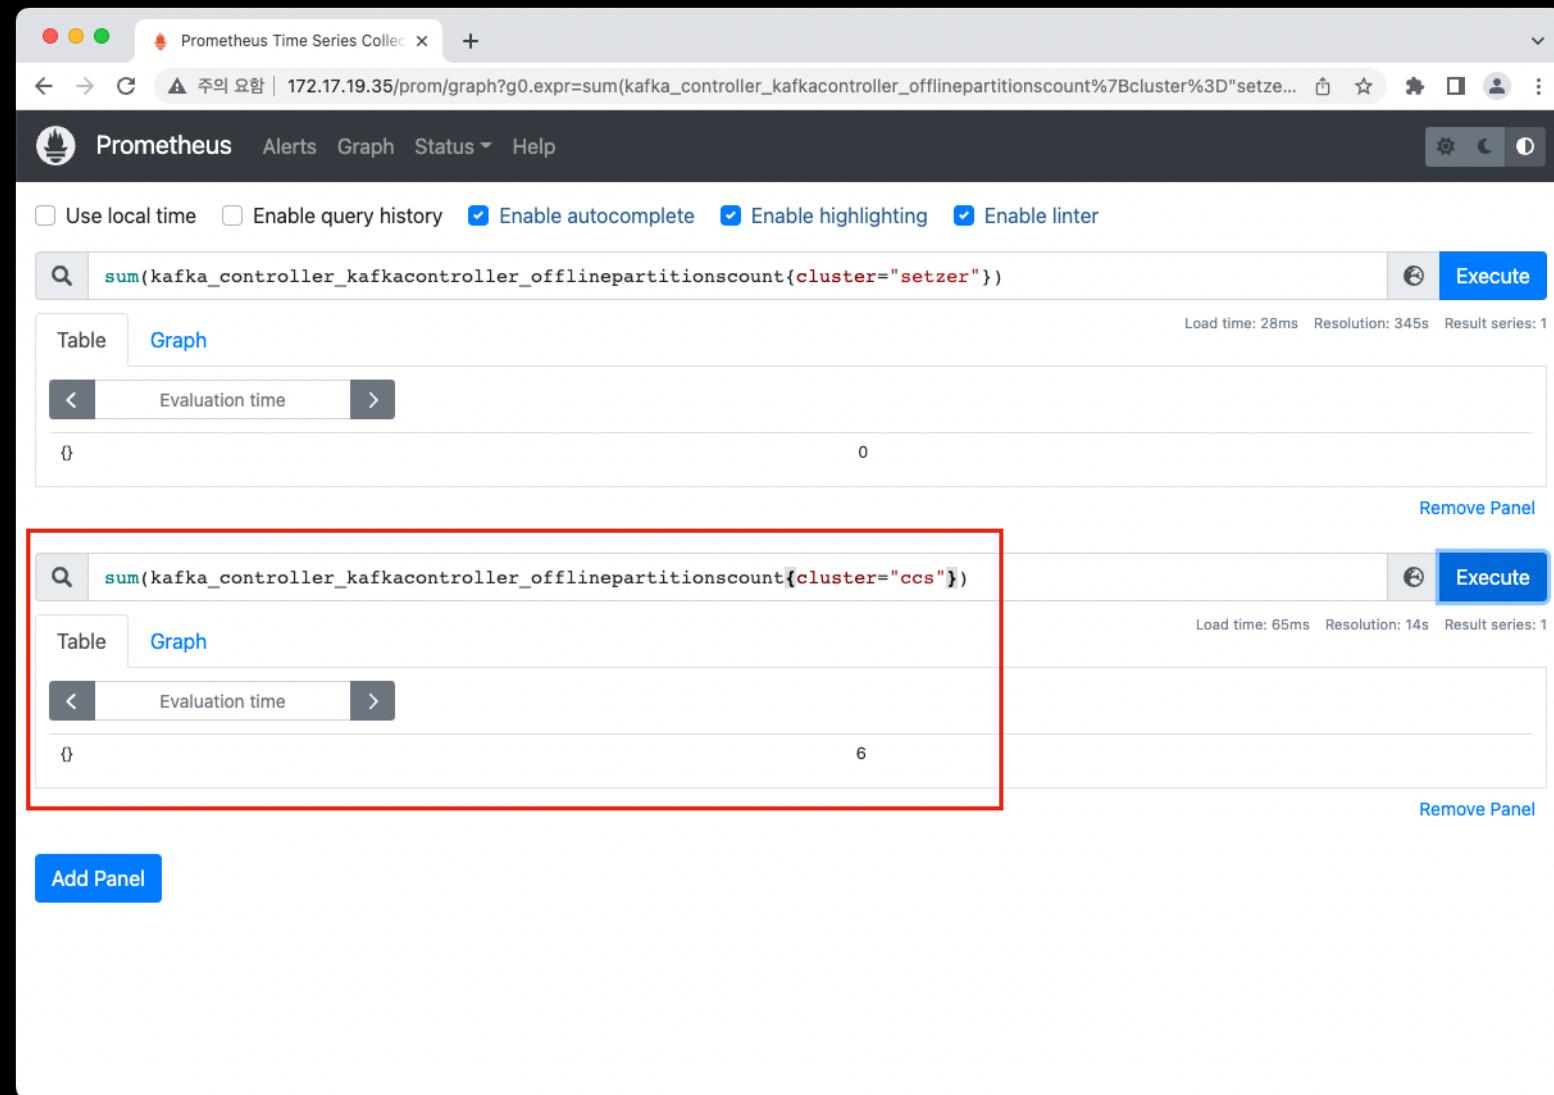
Bottom Panel (Highlighted by a Red Box):

- Query: `sum(kafka_cluster_partition_underreplicated{cluster="ccs"})`
- Evaluation time: {}
- Result: 51
- Load time: 112ms Resolution: 14s Result series: 1

Buttons at the bottom left include "Add Panel" and "Remove Panel".

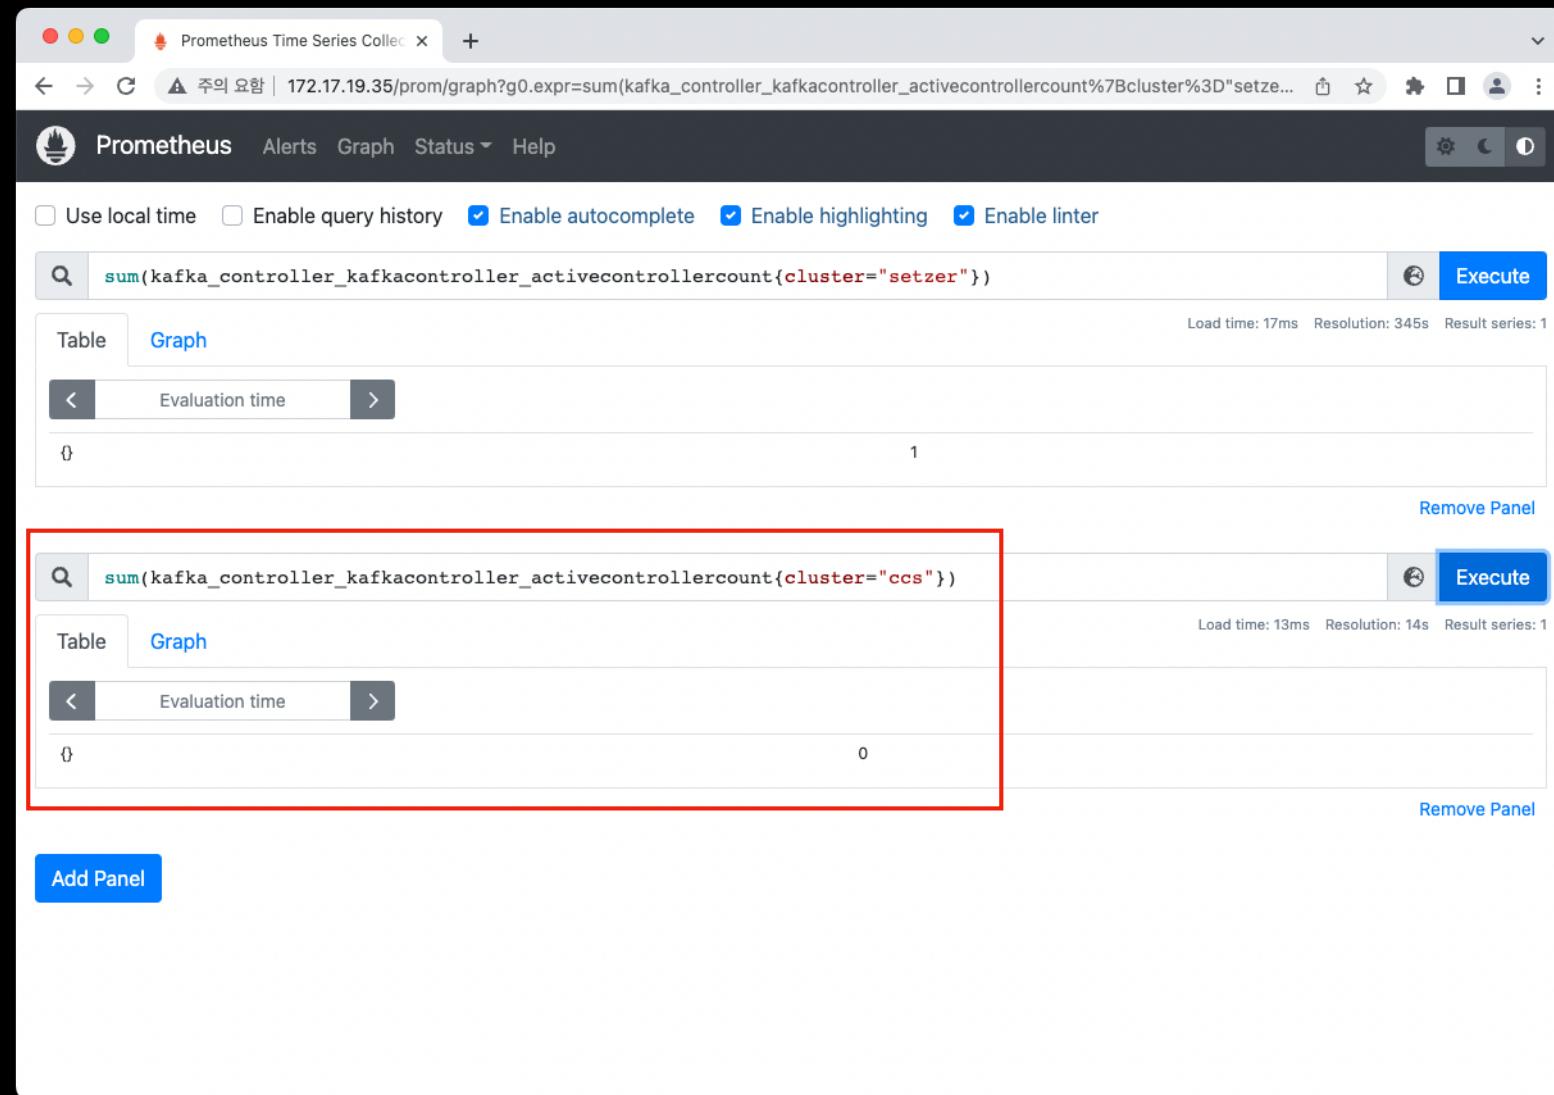
Metrics Example

kafka.controller:type=KafkaController,name=OfflinePartitionsCount



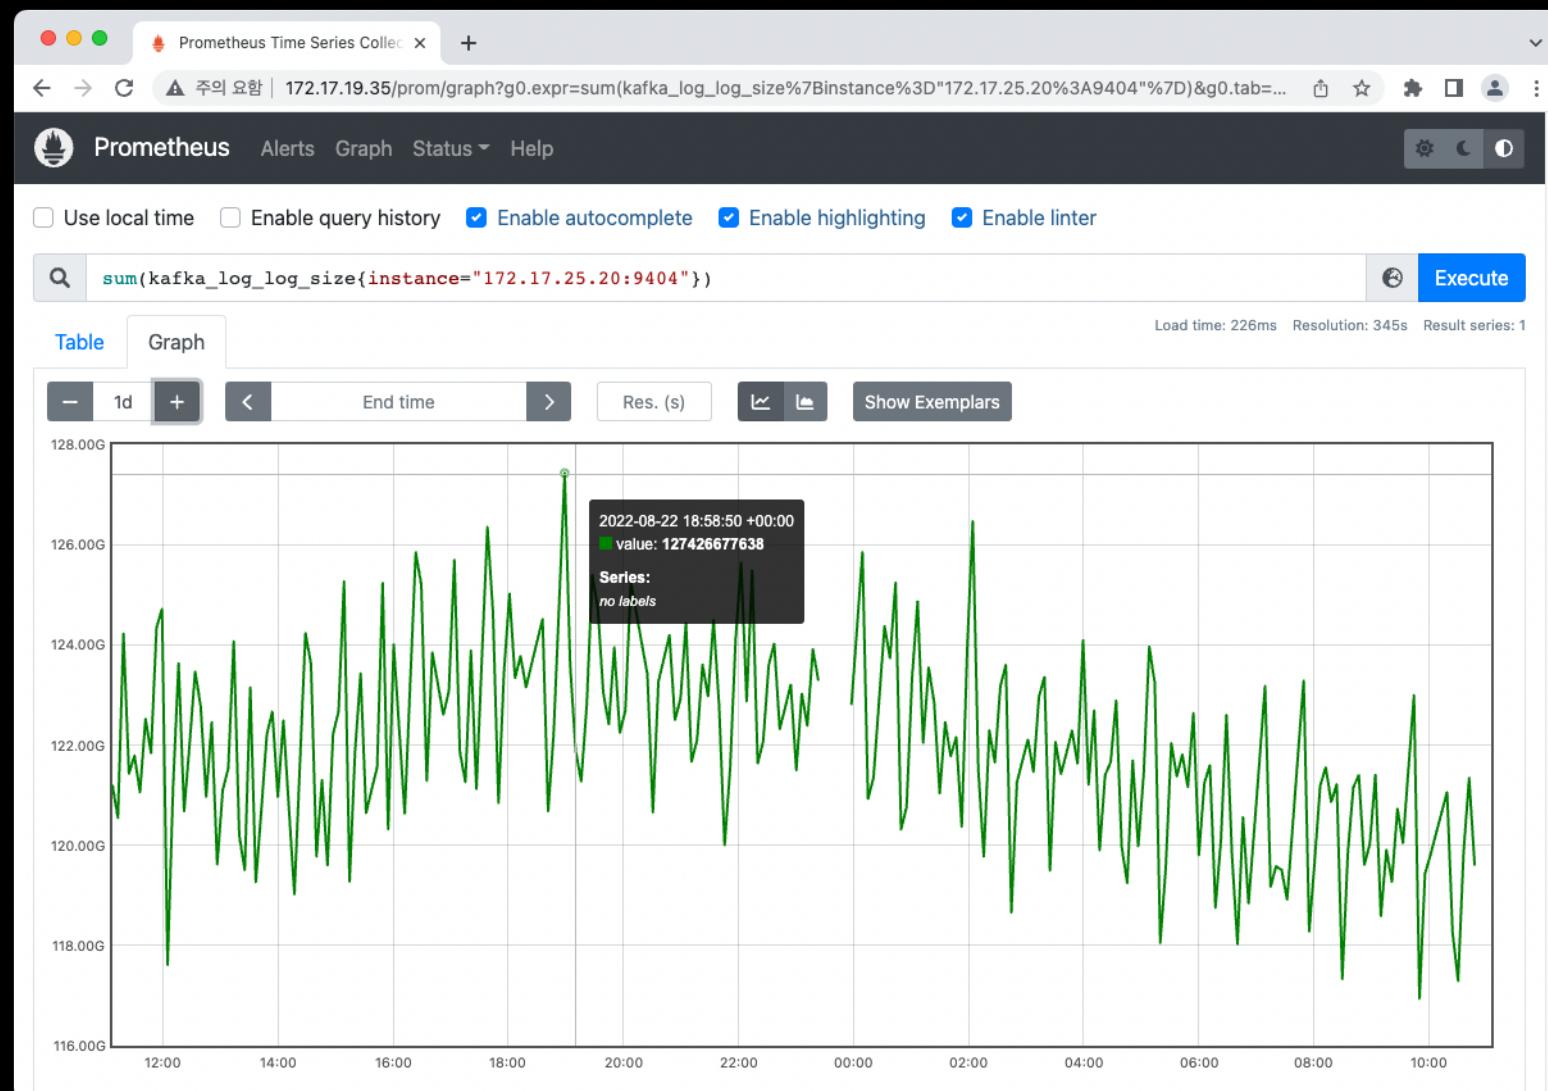
Metrics Example

kafka.controller:type=KafkaController,name=ActiveControllerCount



Metrics Example

kafka.log:type=Log,name=Size



Advantages

Why prometheus is better

- Apache ZooKeeper
 - 3.6.0 버전부터 Metrics Provider로 Prometheus를 번들로 포함
 - 간단한 설정만으로 주키퍼 및 JVM 메트릭을 HTTP를 통해 노출
 - 한 가지 도구를 이용하여 카프카 클러스터 모니터링 시스템 구축

```
# ENABLE METRICS PROVIDER @zoo.cfg
metricsProvider.className=org.apache.zookeeper.metrics.prometheus.PrometheusMetricsProvide
metricsProvider.httpPort=7000
metricsProvider.exportJvmInfo=true
```

What about Server

More metrics are needed

- Node Exporter
 - 노드의 각종 메트릭을 HTTP로 노출해 주는 서버 애플리케이션
 - CPU, Memory, Disk, 네트워크 및 다수의 메트릭 제공
 - https://github.com/prometheus/node_exporter

Example Config

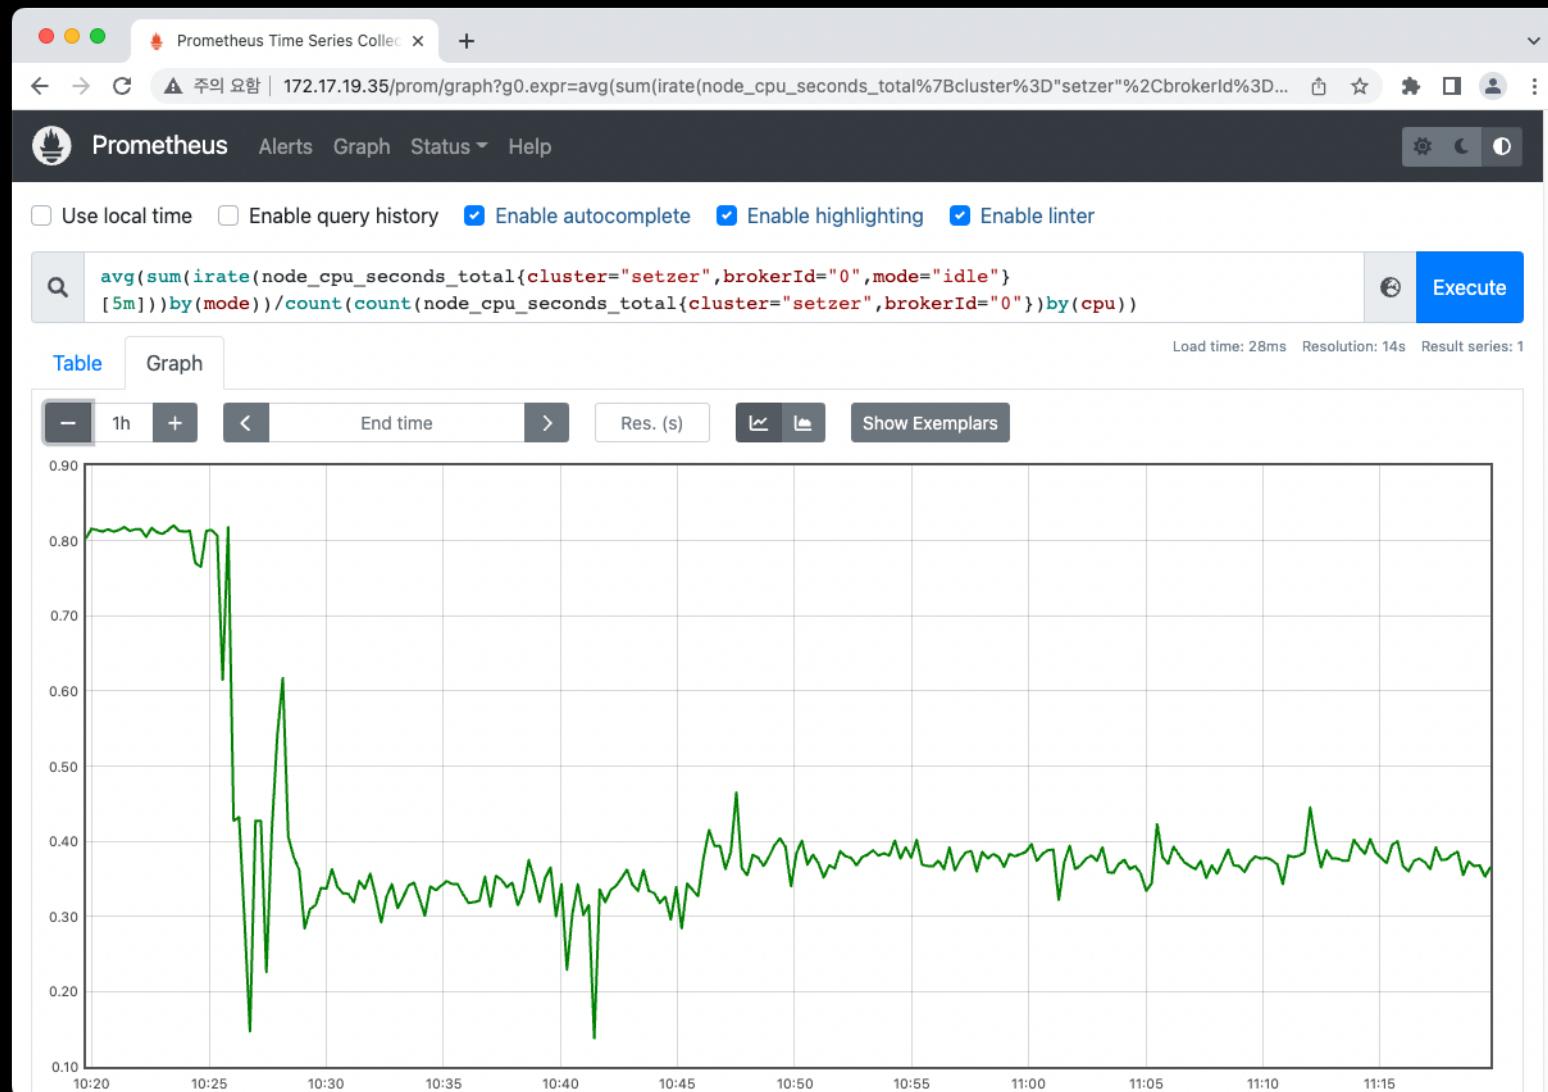
Scraping node metrics from brokers

```
1 ## prometheus.yml ##
2
3 ##### job : kafka #####
4 - job_name: kafka
5   scrape_interval: 5s
6   file_sd_configs:
7     - files:
8       - targets/kafka.yaml
9       refresh_interval: 5m
10
11 ##### job : kafka_node #####
12 - job_name: kafka_node
13   scrape_interval: 5s
14   file_sd_configs:
15     - files:
16       - targets/kafka_node.yaml
17       refresh_interval: 5m
```

```
1 ## kafka_node.yaml ##
2
3 - targets:
4   - "172.17.25.20:9100"
5   labels:
6     cluster: "ccs"
7     brokerId: "0"
8 - targets:
9   - "172.17.25.21:9100"
10  labels:
11    cluster: "ccs"
12    brokerId: "1"
13 - targets:
14   - "172.17.25.22:9100"
15  labels:
16    cluster: "ccs"
17    brokerId: "2"
18 - targets:
19   - "172.17.25.40:9100"
20  labels:
21    cluster: "setzer"
22    brokerId: "0"
23 - targets:
24   - "172.17.25.41:9100"
25  labels:
26    cluster: "setzer"
27    brokerId: "1"
28 - targets:
29   - "172.17.25.42:9100"
30  labels:
31    cluster: "setzer"
32    brokerId: "2"
```

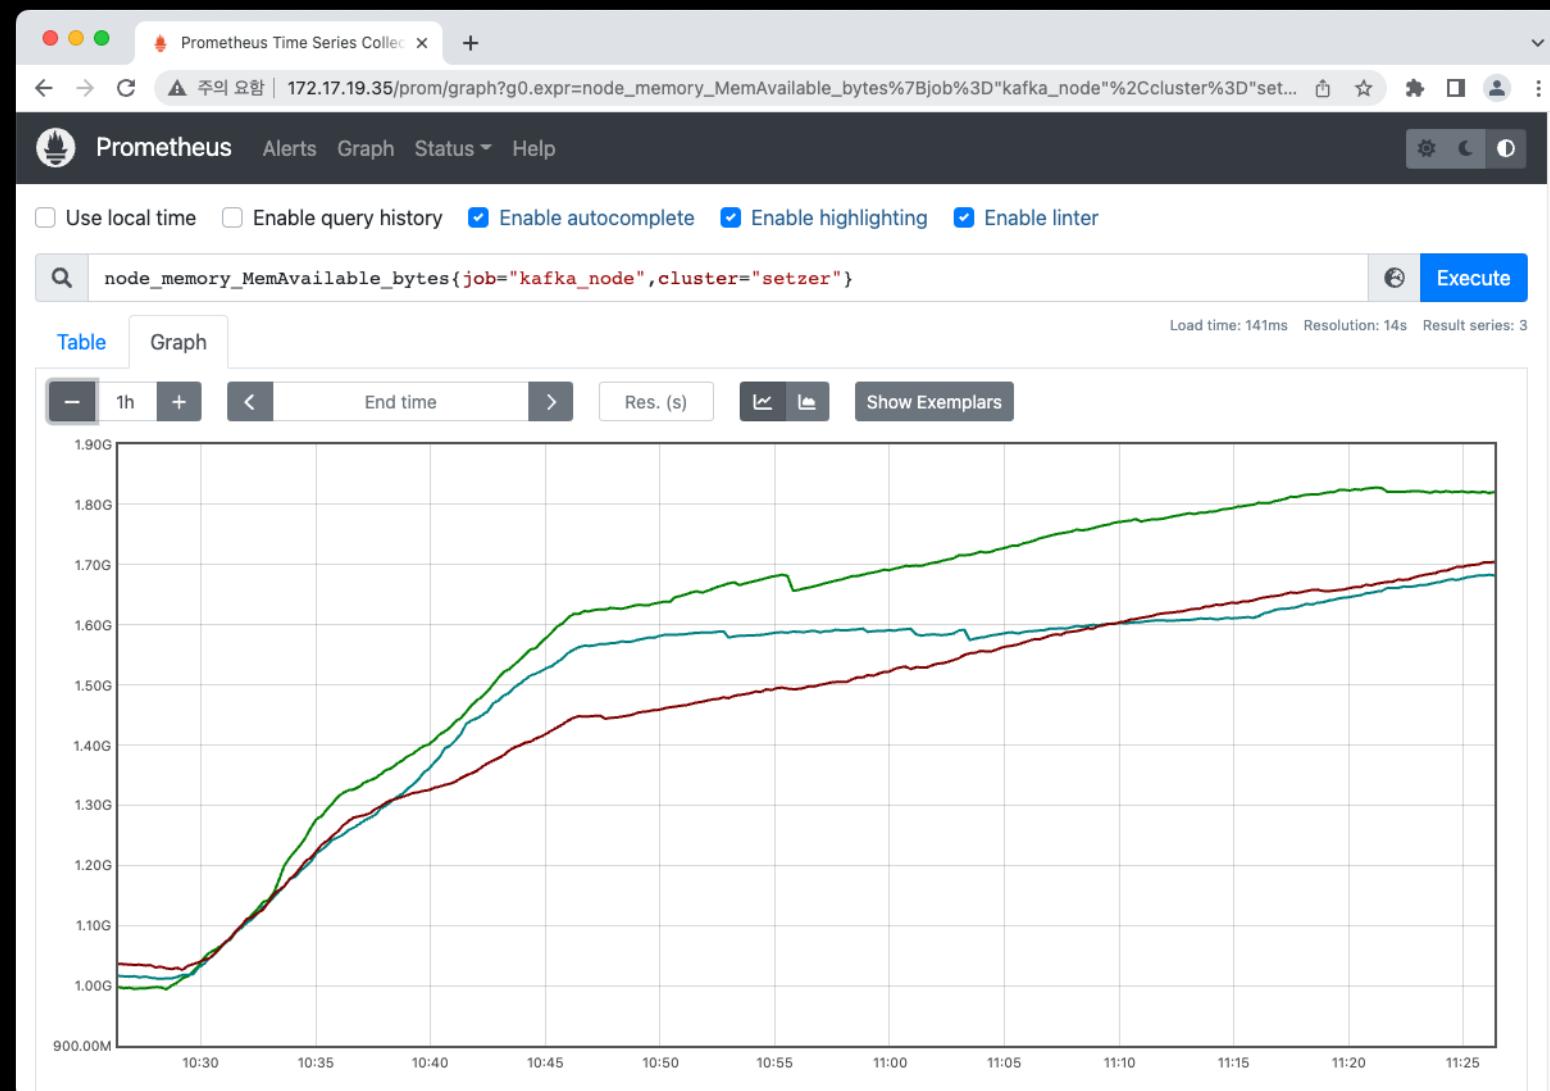
Metrics Example

Average CPU idle



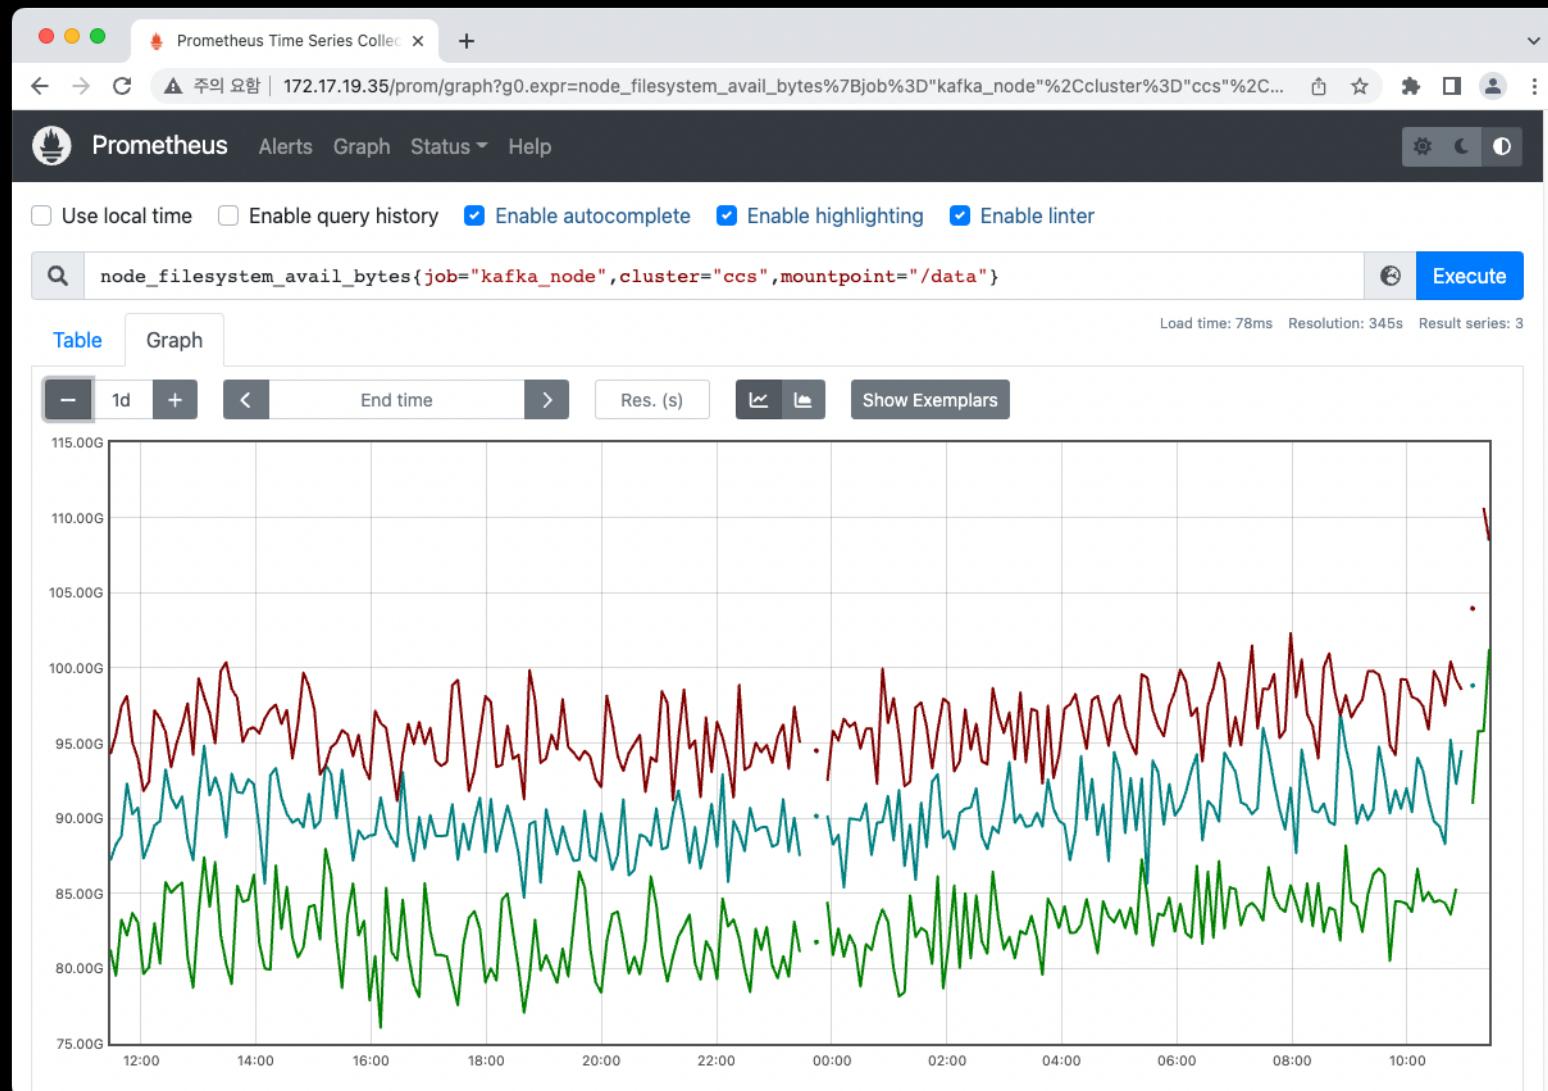
Metrics Example

Memory available bytes



Metrics Example

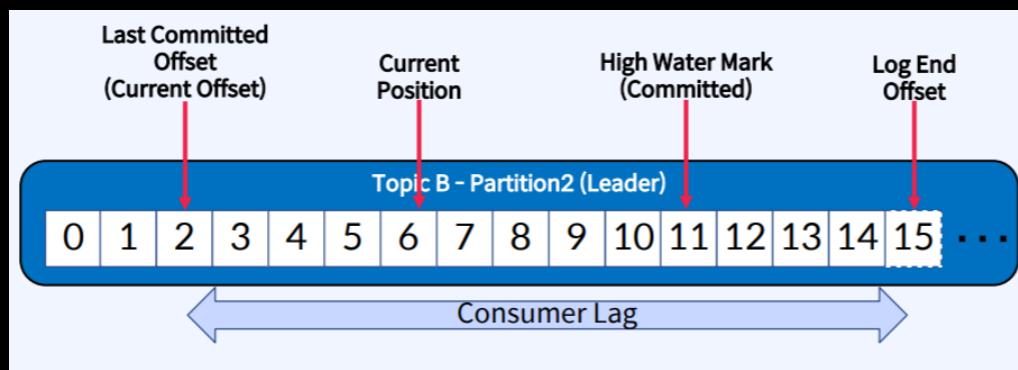
Filesystem available bytes



Offsets

Everything is working fine?

- "컨슈머 랙"?
 - LEO(Log End Offset) - LCO(Last Committed Offset)
 - LEO = 프로듀서가 다음에 보내온 메시지를 쓸 위치
 - LCO = 컨슈머가 마지막으로 커밋한 위치
 - 값이 큰 것 보다 전체적인 경향이 더 중요함
 - 값이 계속 커지는 것은 메시지 유입에 비해 처리가 늦다는 것을 의미



Consumer Healthy

Evaluate consumer lag /w burrow



The Burrow (Franz Kafka Stop-Motion Short Film)

- burrow
 - <https://github.com/linkedin/Burrow>
 - "컨슈머 랙"을 지속해서 모니터링하고 평가할 수 있는 오픈소스 도구
 - 평가된 값을 HTTP로 노출하여 Prometheus를 통해 수집할 수 있음
 - 일정 시간 간격으로 수집한 메트릭 시리즈를 윈도로 설정하여 평가
 - <https://github.com/linkedin/Burrow/wiki/http-request-consumer-group-status>
 - <https://github.com/linkedin/Burrow/blob/master/core/protocol/evaluator.go>

Code	Name	Description
0	NOTFOUND	컨슈머 그룹 없음
1	OK	컨슈머 그룹 / 파티션 상태 GOOD
2	WARN	그룹 혹은 파티션 상태 WARNING. (예) 오프셋은 이동하지만 랙도 증가
3	ERR	그룹 상태 ERROR. (예) 오프셋은 멈췄지만, 랙은 0이 아님
4	STOP	파티션 상태 STOPPED. (예) 장시간 오프셋 커밋 없음
5	STALL	파티션 상태 STALLED. (예) 오프셋은 커밋되지만 랙에 변화 없고 0이 아님
6	REWIND	파티션 상태 REWIND. (예) 오프셋이 이전 값보다 작은 값으로 커밋 됨

Example Config

Scraping consumer healthy from the burrow

```
1 ## prometheus.yml ##
2
3 - job_name: 'consumer-lag'
4   scrape_interval: 5s
5   static_configs:
6     - targets: [ 'burrow:8000' ]
```

```
1 ## burrow.yaml ##
2
3 logging:
4   level: "info"
5 httpserver:
6   default:
7     address: ":8000"
8 storage:
9   felice:
10    intervals: 5
11    class-name: "inmemory"
12    expire-group: 86400
13 sasl:
14   ccs:
15     mechanism: "SCRAM-SHA-512"
16     username: "admin"
17     password: "admin-secret"
18 cluster:
19   ccs:
20     servers:
21       - "172.17.25.20:9093"
22       - "172.17.25.21:9093"
23       - "172.17.25.22:9093"
24     class-name: "kafka"
25     client-profile: "ccs"
26     topic-refresh: 30
27     offset-refresh: 5
28     groups-reaper-refresh: 300
29 setzer:
30   servers:
31     - "172.17.25.40:9092"
32     - "172.17.25.41:9092"
33     - "172.17.25.42:9092"
34   class-name: "kafka"
35   client-profile: "setzer"
36 client-profile:
37   ccs:
38     sasl: "ccs"
39     read-timeout: 5
40     dial-timeout: 10
41     kafka-version: "3.0.0"
42 setzer:
43   read-timeout: 5
44   dial-timeout: 10
45   kafka-version: "3.0.0"
```

Metrics Example

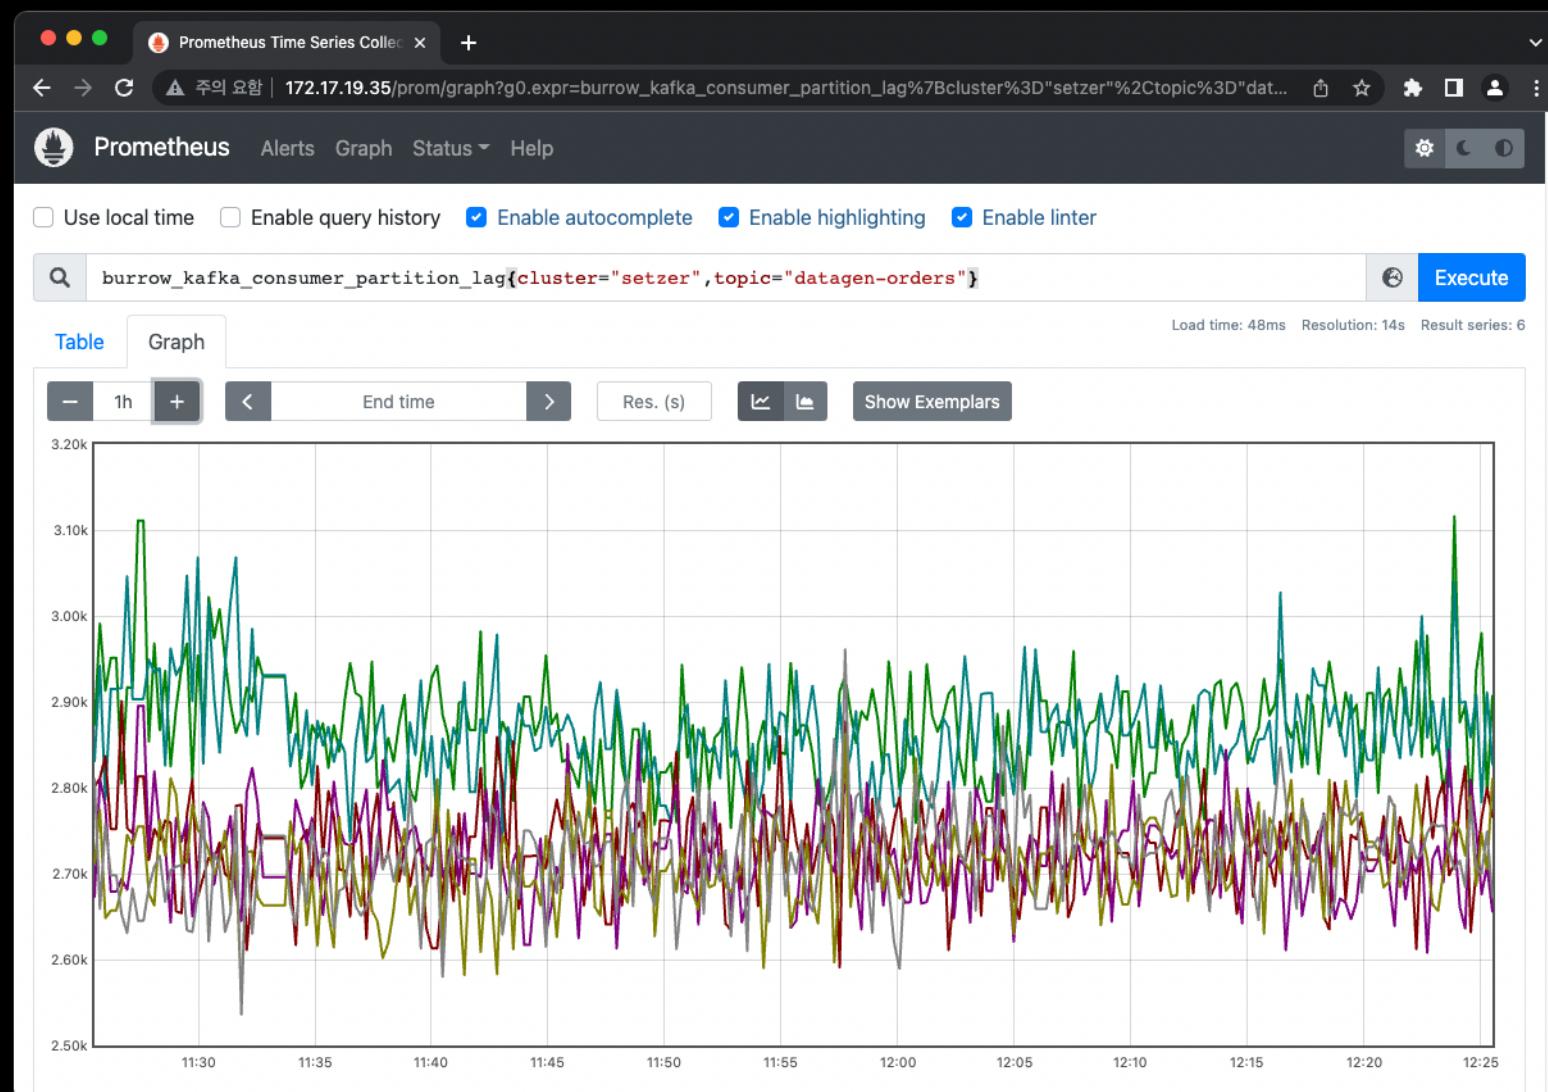
Consumer status

The screenshot shows the Prometheus web interface with the URL `172.17.19.35/prom/graph?g0.expr=burrow_kafka_consumer_status%7Bcluster%3D"newgen"%7D&g0.tab=1&g0.stac...`. The search bar contains the query `burrow_kafka_consumer_status{cluster="newgen"}`. The results table displays 31 series, all with a value of 1. A red box highlights the first three rows of the table.

Series	Value
<code>burrow_kafka_consumer_status{cluster="newgen", consumer_group="burrow-spark-newgen-552c3d1c-aaba-4c4f-990c-3da8faba16b4", instance="felice-lag:8000", job="consumer-lag"}</code>	1
<code>burrow_kafka_consumer_status{cluster="newgen", consumer_group="datagen-consumer-1", instance="felice-lag:8000", job="consumer-lag"}</code>	1
<code>burrow_kafka_consumer_status{cluster="newgen", consumer_group="felice-streams-alarm-action-develop1", instance="felice-lag:8000", job="consumer-lag"}</code>	1
<code>burrow_kafka_consumer_status{cluster="newgen", consumer_group="felice-streams-alarm-action-setzer", instance="felice-lag:8000", job="consumer-lag"}</code>	1
<code>burrow_kafka_consumer_status{cluster="newgen", consumer_group="felice-streams-alarm-alert-dev-develop1", instance="felice-lag:8000", job="consumer-lag"}</code>	1
<code>burrow_kafka_consumer_status{cluster="newgen", consumer_group="felice-streams-alarm-alert-dev-setzer", instance="felice-lag:8000", job="consumer-lag"}</code>	1
<code>burrow_kafka_consumer_status{cluster="newgen", consumer_group="felice-streams-alarm-alert-log-dev-develop1", instance="felice-lag:8000", job="consumer-lag"}</code>	1
<code>burrow_kafka_consumer_status{cluster="newgen", consumer_group="felice-streams-alarm-alert-log-dev-setzer", instance="felice-lag:8000", job="consumer-lag"}</code>	1
<code>burrow_kafka_consumer_status{cluster="newgen", consumer_group="felice-streams-alarm-group-setzer", instance="felice-lag:8000", job="consumer-lag"}</code>	1
<code>burrow_kafka_consumer_status{cluster="newgen", consumer_group="felice-streams-alarm-target-setzer", instance="felice-lag:8000", job="consumer-lag"}</code>	1
<code>burrow_kafka_consumer_status{cluster="newgen", consumer_group="felice-streams-audit-log-bak-develop1", instance="felice-lag:8000", job="consumer-lag"}</code>	1
<code>burrow_kafka_consumer_status{cluster="newgen", consumer_group="felice-streams-audit-log-bak-setzer", instance="felice-lag:8000", job="consumer-lag"}</code>	3
<code>burrow_kafka_consumer_status{cluster="newgen", consumer_group="felice-streams-audit-log-dev-develop1", instance="felice-lag:8000", job="consumer-lag"}</code>	1
<code>burrow_kafka_consumer_status{cluster="newgen", consumer_group="felice-streams-audit-log-dev-setzer", instance="felice-lag:8000", job="consumer-lag"}</code>	3
<code>burrow_kafka_consumer_status{cluster="newgen", consumer_group="felice-streams-felice-token-usage-dev-develop1", instance="felice-lag:8000", job="consumer-lag"}</code>	1
<code>burrow_kafka_consumer_status{cluster="newgen", consumer_group="felice-streams-felice-token-usage-dev-setzer", instance="felice-lag:8000", job="consumer-lag"}</code>	3
<code>burrow_kafka_consumer_status{cluster="newgen", consumer_group="felice-streams-felice-user-dev-develop1", instance="felice-lag:8000", job="consumer-lag"}</code>	1

Metrics Example

Partition lag



Constraint

Weak point of the burrow

- 카프카 인증 지원 부족
 - 카프카 SASL은 다양한 메커니즘 지원
 - PLAIN, SCRAM-SHA-256/SCRAM-SHA-512, GSSAPI, OAUTHBEARER
 - Burrow는?
- 개발이 느린 편
 - sarama는 업데이트되었지만...
 - 코드는 업데이트되었지만 매뉴얼 업데이트는...
 - 쌓여 가는 "Pull requests"
 - SASL/PLAIN, Config Auto Reload, SASL/GSSAPI, ...

Can't watch 24/7

Need a way to know events

- Prometheus
 - targets = 메트릭 스크랩 대상
 - rules = 스크랩된 메트릭 평가 규칙
 - 평가 규칙에 부합하는 경우 Alertmanager로 alert 전달
- Alertmanager
 - alert를 email, slack, pageduty 등으로 전송
- Burrow
 - notifier 모듈 = 컨슈머 그룹 상태를 외부 HTTP 서버로 전송
 - 프로메테우스 rule을 통해 더 상세하게 평가 가능

Rule Example

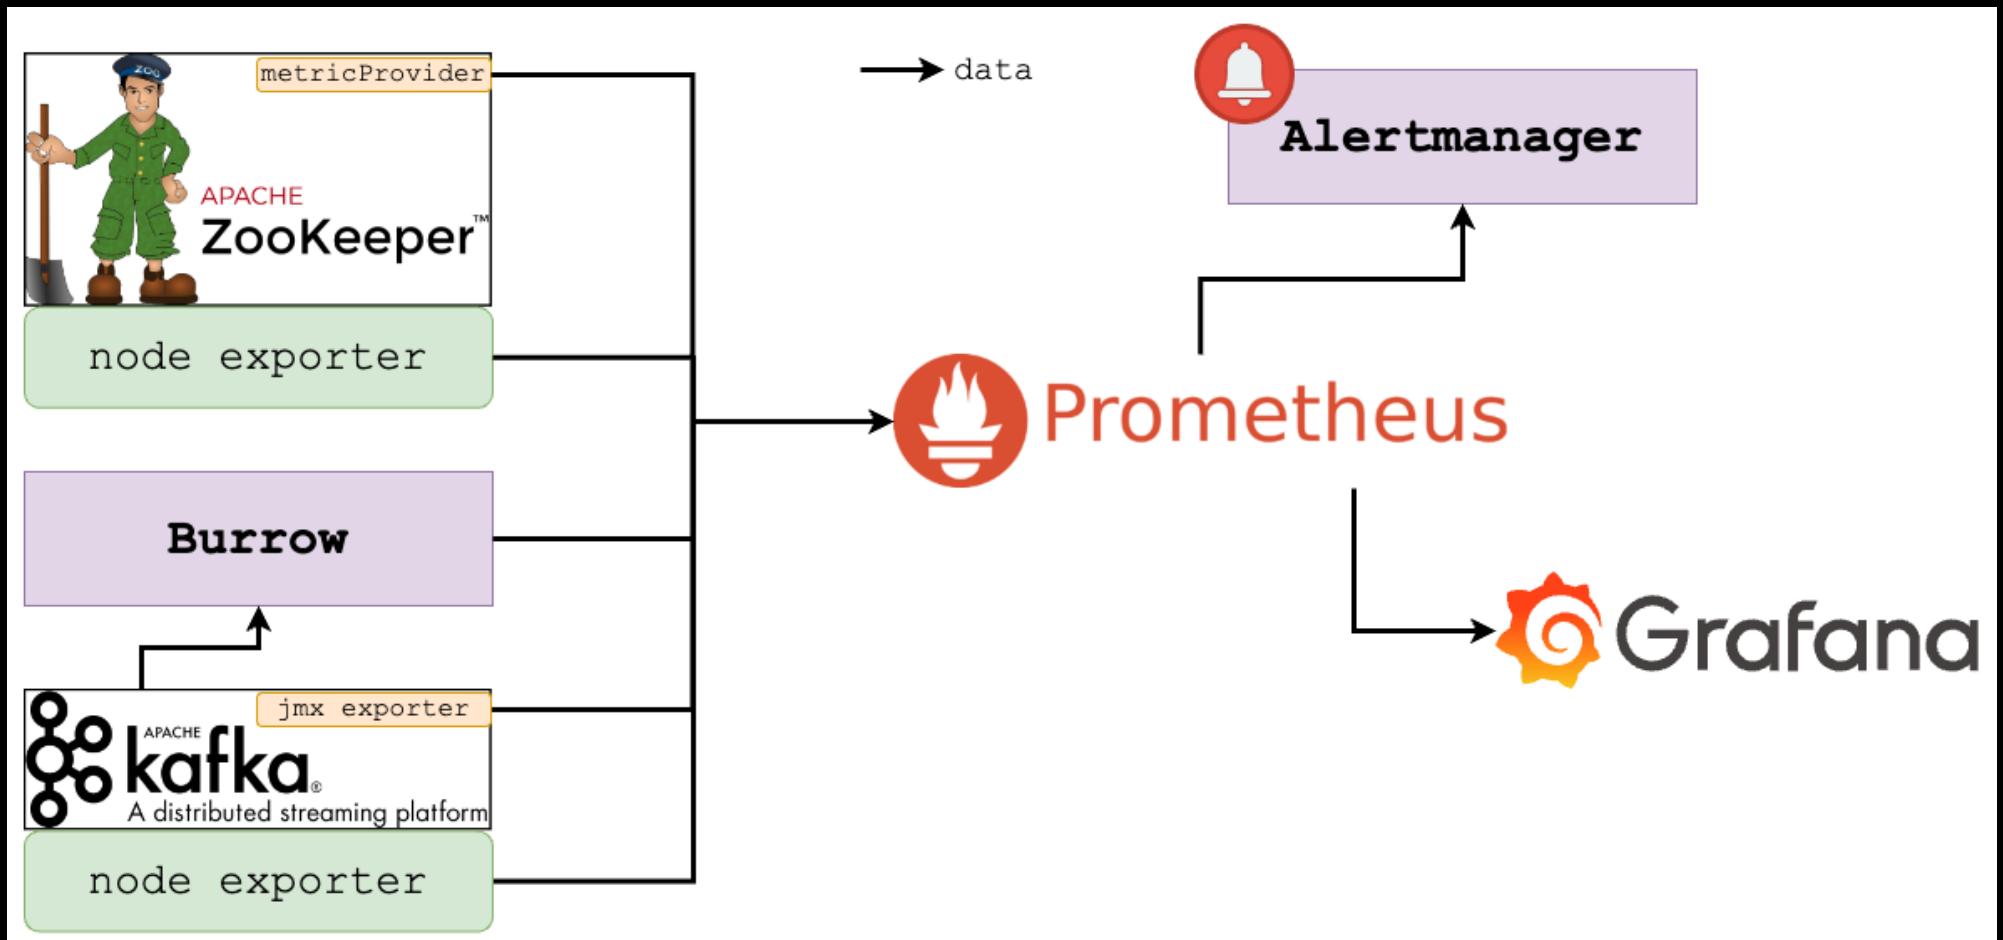
Alerting issues

The screenshot shows the Prometheus UI interface with three alert rules listed in a table.

Rule	State	Error	Last Evaluation	Evaluation Time
<pre>alert: [WARNING] Number of the under-replicated partitions expr: kafka_server_replicamanager_underreplicatedpartitions{cluster="ccs"} >= 1 for: 5m labels: action: ccs_system_0443d0aaacdb4a2a8f1a690e124b9ee8 severity: 5 annotations: summary: [{{\$labels.cluster}}-{{\$labels.job}}-{{\$labels.brokerId}}, under-replicated_partition: {{\$Value}}] Under-replicated partition have found for more than a minute</pre>	OK		6.200s ago	0.460ms
<pre>alert: [CRITICAL] ZooKeeper instance DOWN expr: up{cluster="ccs",job="zookeeper"} == 0 for: 5m labels: action: ccs_system_2aa63985ba954dbfb651b90e6343018c severity: 3 annotations: summary: [{{\$labels.cluster}}-{{\$labels.job}}-{{\$labels.brokerId}}] ZooKeeper instance DOWN</pre>	OK		6.199s ago	0.251ms
<pre>alert: [CRITICAL] CPU usage expr: ((count by(cluster, job, brokerId) (count by(cluster, job, brokerId, cpu) (node_cpu_seconds_total{cluster="ccs",job="kafka_node"}) - avg by(cluster, job, brokerId) (sum by(cluster, job, brokerId) (rate(node_cpu_seconds_total{cluster="ccs",job="kafka_node",mode="idle" [5m]}))) * 100 / count by(cluster, job, brokerId) (count by(cluster, job, brokerId, cpu) (node_cpu_seconds_total{cluster="ccs",job="kafka_node"})))) >= 70 for: 5m</pre>	OK		6.199s ago	3.559ms

Avengers Assemble

Integrate all parts into one system



Client metrics

Scraping metrics via pushgateway

1. client(producer, consumer)에서 메트릭을 추출

- org.apache.kafka.clients.producer.KafkaProducer<K,V>

```
Map<MetricName, ? extends Metric> metrics()
Get the full set of internal metrics maintained by the producer.
```

- org.apache.kafka.clients.consumer.KafkaConsumer<K,V>

```
Map<MetricName, ? extends Metric> metrics()
Get the metrics kept by the consumer
```

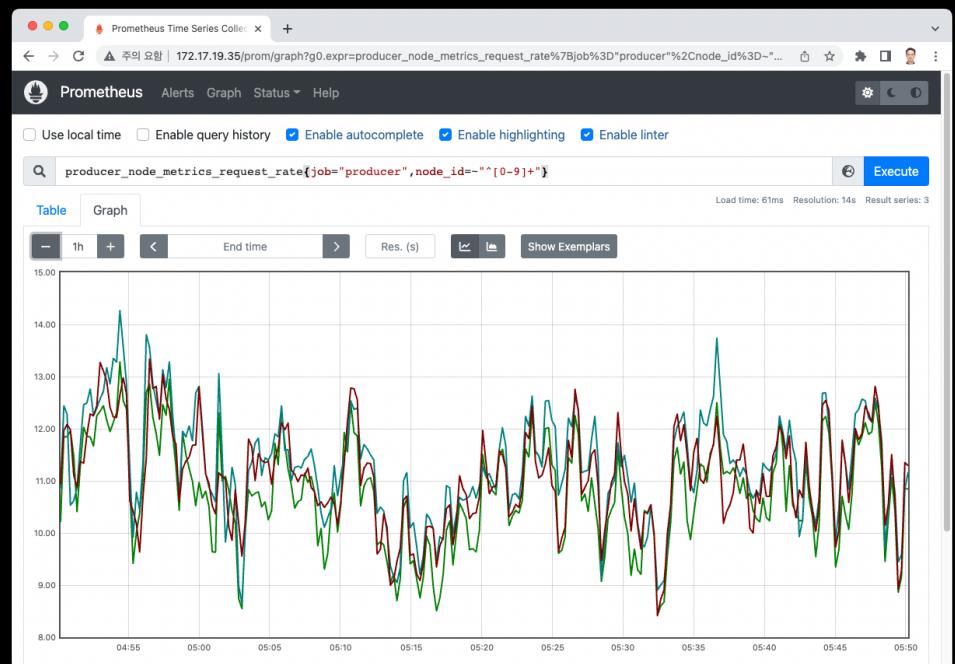
2. pushgateway에 전송

- <https://github.com/prometheus/pushgateway>
- io.prometheus.client.exporter.PushGateway

```
void executeBatchJob() throws Exception {
    CollectorRegistry registry = new CollectorRegistry();
    Gauge duration = Gauge.build()
        .name("my_batch_job_duration_seconds")
        .help("Duration of my batch job in seconds.")
        .register(registry);
    Gauge.Timer durationTimer = duration.startTimer();
    try {
        // Your code here.

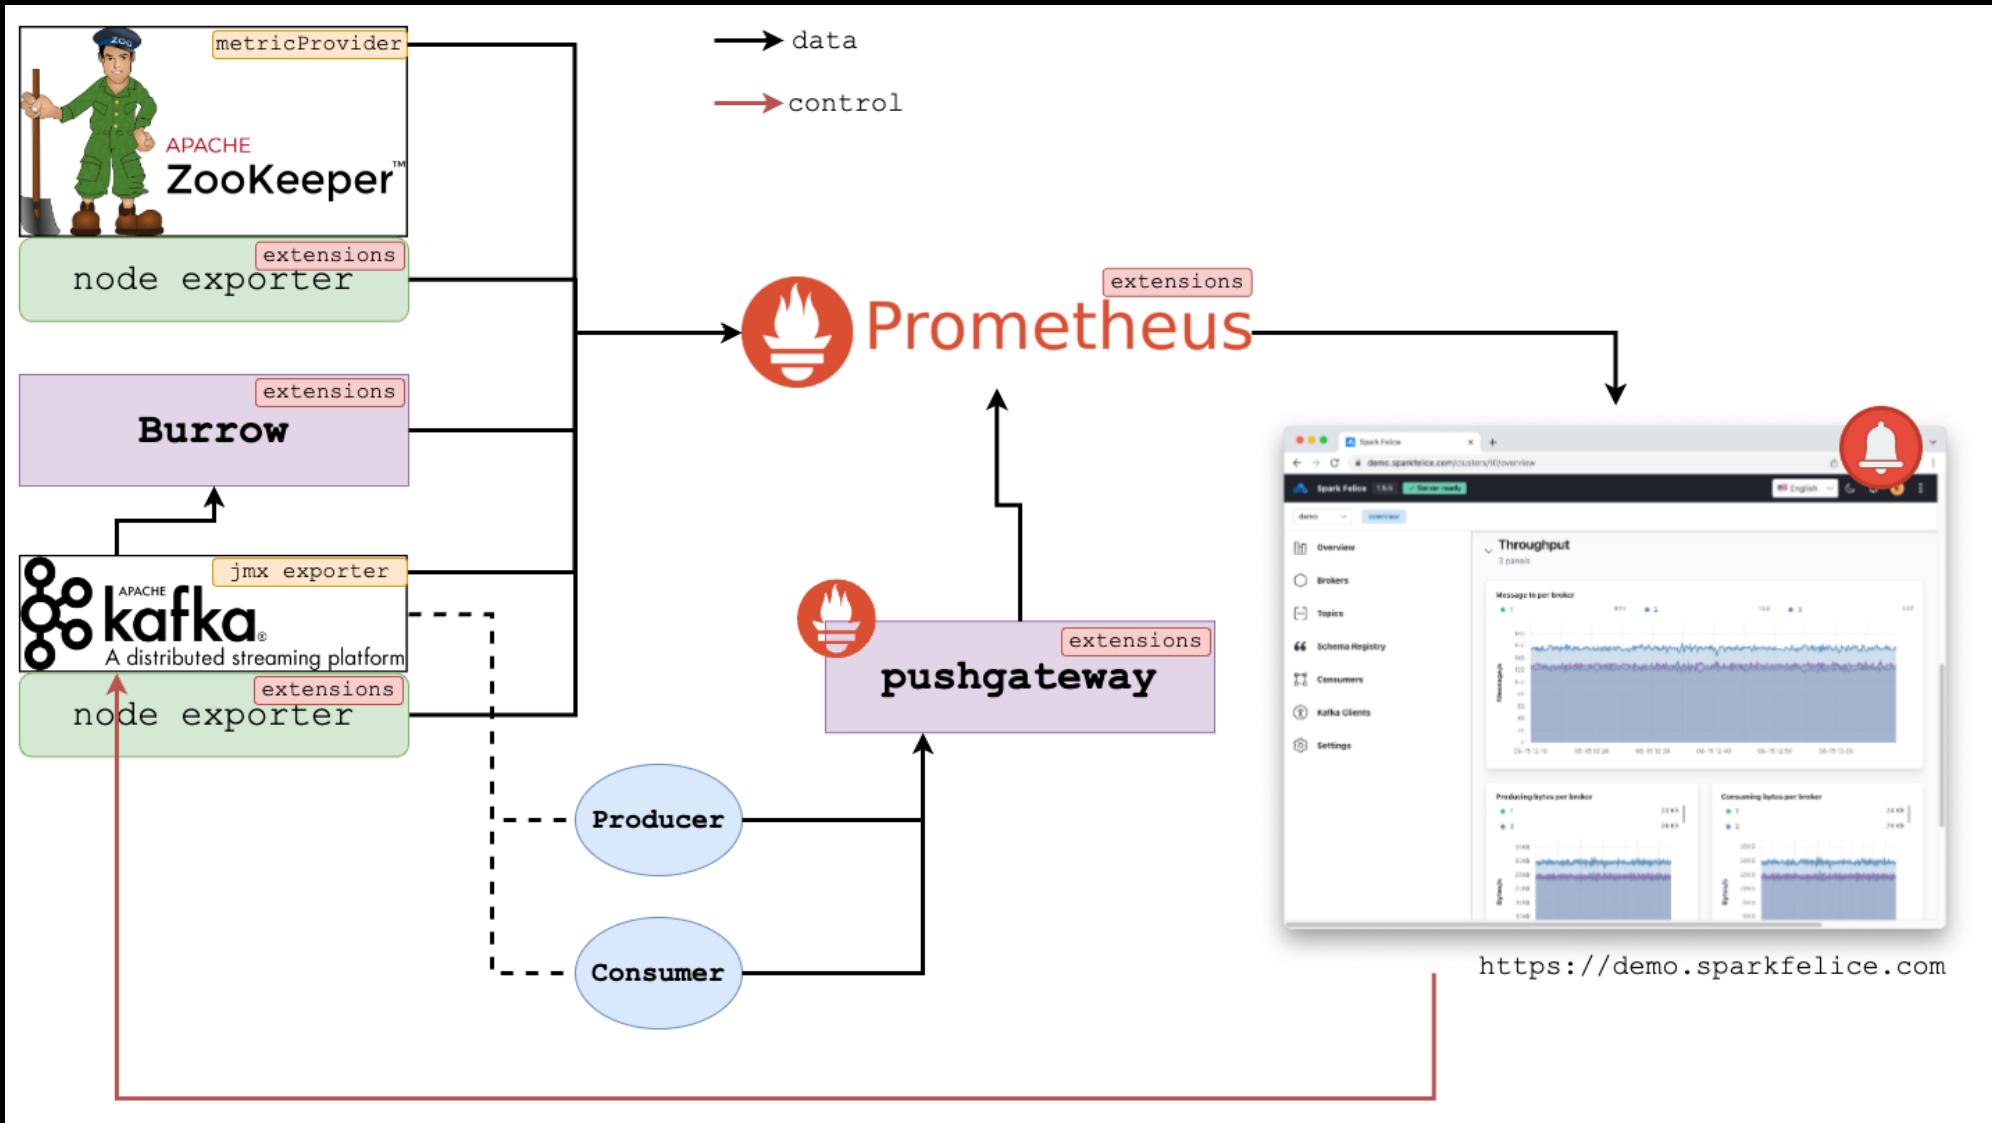
        // This is only added to the registry after success,
        // so that a previous success in the Pushgateway isn't overwritten on failure.
        Gauge lastSuccess = Gauge.build()
            .name("my_batch_job_last_success")
            .help("Last time my batch job succeeded, in unixtime.")
            .register(registry);
        lastSuccess.setToCurrentTime();
    } finally {
        durationTimer.setDuration();
        PushGateway pg = new PushGateway("127.0.0.1:9091");
        pg.pushAdd(registry, "my_batch_job");
    }
}
```

3. prometheus에서 scrap



One more thing...

Our implementations



What's different?

Our custom modifications

- Burrow
 - Notifier, ZooKeeper 모듈 비활성화
 - 설정 파일 변경 시 자동 재시작
 - SASL/PLAIN, GSSAPI, OAUTHBEARER 인증 메커니즘
- Node Exporter
 - 네트워크 상세 연결 정보 추가
- Prometheus
 - 설정 파일 변경 시 자동 재시작
- Pushgateway
 - 메트릭 TTL 적용

Kafka's Fiance, Felice

Our Monitoring and Managing Tool for AK

- 카프카 클러스터 모니터링
 - 주키퍼, 브로커, 서버, 메시지 입출력, 컨슈머, etc.
- 카프카 클러스터 관리
 - 토픽, 파티션, 카프카 계정, ACL, 스키마, etc.
- Alerting
 - 알림 룰, 알림 타깃, 알림 조회, etc.
- Enterprise Features
 - Audit Log, RBAC, etc.
- Community Edition Coming Soon~

Questions?

Hope this helps

